



OKTOFLOW PLATFORM HANDBOOK

Version **0.8.1-SNAPSHOT**

Holger Eichelberger, Ahmad Alamoush, Monika Staciwa

Disclaimer

The contents of this document has been prepared with great carefulness. Although the information has been prepared with the greatest possible care, there is no claim to factual correctness, completeness and/or timeliness of data; in particular, this publication cannot take into account the specific circumstances of individual cases.

Any use is therefore the reader's own responsibility. Any liability is excluded. This document contains material that is subject to the copyright of individual or multiple IIP-Ecosphere or ReGaP consortium parties. All rights, including reproduction of parts, are held by the authors.

This document reflects only the views of the authors at the time of publication. The Federal Ministry for Economic Affairs and Energy (BMBF), the Federal Ministry for Research, Technology and Space (BM-FTR, previously BMBF) or the responsible project agency are not liable for the use of the information contained herein.

Publication: **September, 2026**

DOI: 10.5281/zenodo.19704503

As oktoflow is now used in ReGaP as technical innovation platform and this handbook has been revised for the use in ReGaP, we decided to remove authors that were not active for a longer time or are not active anymore. Thus, the author list of this handbook is dynamic and may be extended again in the future.

Gefördert durch:



Executive Summary

The oktoflow platform (initially developed in the BMWK project IIP-Ecosphere) is a novel platform for Industry 4.0, e.g., based on asset administration shells as interfaces for software components and resources, unified edge deployment, an AI toolkit or seamless configuration of a platform from network settings via services up to applications running on the platform. This platform handbook provides insights into the rationales, ideas and concepts that make up the design and the realization of the platform, ranging from an overall layered architecture over a detailed discussion of the design and realization state of each layer up to cross-cutting mechanisms such as the configuration model or the related code/artifact generation.

This platform handbook addresses the technical side of the platform and builds on the intensive prior work on requirements (usage view and functional/quality requirements of the platform). This handbook shall provide means for deeper technical discussions with partners, stakeholders and interested parties, but also allow for a technical understanding to contribute to the platform, e.g., in terms of protocols, platform connectors, services or demonstration applications.

This version of the handbook focuses on the platform release as of **September 2026** (version **0.8.1**) and supersedes older versions of this handbook/the platform.

Acknowledgements: We would like to thank Christian Nikolajew for his work on the MODBUS/TCP and REST connectors, Jobst Hillebrandt for his work on the ADS connector as well as Thomas Lepper and Aleks Arzer from PZH/IFW of the Leibniz University Hannover for their testing support and input.

oktoflow was partially supported by the BMWK project IIP-Ecosphere (grant 01MK20006C) and DatiPilot ReGaP-PgE (grant 03DPC1511B) and by the BMFTR DatiPilot Innovationcommunity ReGaP, sub-project ReGaP-PgE, (grant 03DPC1511B).

Contents

1	Introduction.....	6
1.1	Motivation and Goals	6
1.2	Structure of the document.....	7
2	Tooling and Basic Technical Decisions.....	9
3	Architecture.....	12
3.1	Overview.....	12
3.1.1	Relation to Reference Architectures	17
3.1.2	Stream (Data) Processing	17
3.1.3	Asset Administration Shells.....	18
3.1.4	Component Interaction Overview.....	20
3.1.5	Virtual Character of the Platform.....	22
3.2	Overall Requirements.....	23
3.3	Support Layer	24
3.3.1	Component Structure of the Support Layer.....	25
3.3.2	The support.boot Component.....	25
3.3.3	The support Component	27
3.3.4	The support.aas Component.....	29
3.3.5	The support.iip-aas Component.....	33
3.3.6	AAS Creation and Usage Pattern	35
3.3.7	Plugins	36
3.4	Transport and Connection Layer	38
3.4.1	Transport Component	38
3.4.2	Connectors Component	44
3.5	Services Layer	50
3.5.1	Terminology and Background.....	50
3.5.2	Service Environments	51
3.5.3	Service Control and Management.....	57
3.5.4	Design and (Plugin-)Interfaces	58
3.5.5	Spring-based Service Control and Management.....	60
3.5.6	Validation.....	62
3.6	Resources and Monitoring Layer.....	63
3.6.1	ECS-runtime.....	63
3.6.2	Device/Resource Management.....	67
3.6.3	Monitoring.....	69
3.7	Storage, Security and Data Protection Layer	70
3.7.1	KODEX platform service.....	71

3.7.2	Influx DB connector	71
3.8	Reusable Intelligent Services Layer	72
3.8.1	Data Processing Function Library	72
3.8.2	RapidMiner RTSA service.....	73
3.8.3	Flower-based Federated Learning.....	73
3.9	Configuration Layer	74
3.10	Application Layer	76
3.11	Platform Server(s).....	76
3.12	Platform Management User Interface	77
3.13	Test support.....	82
4	Architectural Decisions and Constraints	86
5	Asset Administration Shells.....	89
6	Platform Configuration.....	91
6.1	Modeling Patterns	97
6.2	Configuration Model Structure	101
6.3	Special configuration topics	103
6.3.1	Deployment plans.....	103
6.3.2	App Templates.....	103
6.3.3	Identity store	104
6.3.4	State machines	104
6.3.5	Dashboard support.....	104
6.4	Support for Standardized Connectors/Protocols	105
6.5	Platform Instantiation Process	105
6.6	Container Instantiation	108
7	Implementation Aspects	112
7.1	New components	112
7.2	Compiling the Platform	112
7.3	Installing and Using the Platform	115
8	Summary & Conclusions.....	116
9	References.....	117

1 Introduction

1.1 Motivation and Goals

The digitalization of the industry increases the effectiveness of technical systems and related processes, but also affects the complexity of the realizing (software) systems. Currently, several approaches are developed in the fields of Internet-of-Things (IoT), Industrial Internet-of-Things (IIoT) or „Industrie 4.0“ (I4.0)¹. To support the industrial transformation towards IoT, IIoT and I4.0, several software platforms were developed that provide different capabilities [ESA+25].

We understand the term platform as a coherently integrated set of software frameworks or libraries to allow for and enable the execution of user-defined apps, here, in the domain of Industry 4.0 [EN23]. A platform may support distributed execution of the apps, may be installed in a cloud, locally on-premise or in a hybrid form exploiting the edge-cloud continuum.

The oktoflow platform was created in the context of the BMWi-funded² project IIP-Ecosphere, which pursued the vision of enabling innovations in the area of industrial production based on connected, intelligent and autonomous systems in order to increase productivity, flexibility, robustness and efficiency of IIoT and I4.0. IIP-Ecosphere aimed at creating a novel ecosystem for the “next level” of intelligent industrial production, not only for software-based systems, but also for the people involved in this kind of systems, e.g., automation engineers, software developers, AI experts, startups, venture capitalists, etc. On the software side, one core activity in IIP-Ecosphere was to research and to realize a virtual platform that connects factory installations across companies in a vendor-independent manner. In particular, the platform shall provide easy-to-use access to Artificial Intelligence (AI) in secure and flexible manner. This platform, oktoflow, became in BMFTR DATIpilot ReGaP³ in 2025 the technological innovation core for energy applications in the industrial context.

Towards the design of such a platform, we analyzed in [SEA+20, ESA+25] more than 40 research IIoT platforms and 21 industrial IIoT platforms with specific relevance to IIP-Ecosphere. In [SSE21, ESA+21], we discussed the requirements for the oktoflow platform from two different perspectives, namely the usage view and the functional/quality requirements view [ESS22]. The resulting platform design shall be open, extensible, vendor-neutral, secure, flexible, configurable, self-adaptive and based on relevant standards as well as on existing Open Source components. In particular, we aim at developing a **virtual platform**, i.e., a platform that utilizes existing, already installed solutions by integrating with them, using accessible output and resources, enhancing them with AI and, if desired, feeding back AI-enhanced information into utilized systems. Thus, we do not aim at replacing existing platforms as those mentioned in [SEA+20] rather than enhancing them. To a certain extend, this also covers the need of openly integrating various solutions into the platform using different mechanisms. Moreover, we aim at demonstrating how research results, e.g., from systematic variability management, security or data protection, can lead to platform concepts that are currently rarely used in IIoT/I4.0 platforms [ESA+25]. Besides the desirable abilities mentioned above, following the initial decisions made in [SSE21, ESA+21], the platform shall be service-based and virtualized through containers. One relevant Industry 4.0 standard to use and to integrate the parts and pieces of the platform is the **Asset Administration Shell** (AAS) that we aim to apply as self-description and interface to software components across all platform layers. The IIP-Ecosphere consortium discussions regarding a vision of the platform also emphasized the need to directly communicate with production machines, in particular, to utilize edge devices as compute resources and, if feasible, cloud technology. This re-shaped the character of the envisioned platform from a purely virtual to a mixed-virtual platform with

¹ Translates to some degree to IIoT in German-speaking areas in Europe, partly based on own standards.

² <https://www.bmw.de/Redaktion/DE/Publikationen/Technologie/ki-innovationswettbewerb.html>

³ <https://regap.de>

stronger aspects of a usual IIoT/II.0 platform, in particular providing **uniform deployment of services** to heterogeneous execution resources such as edge devices, on-premise servers or clouds.

In this handbook, we aim at discussing and documenting the architecture and the implementation of the platform. This happens in an incremental⁴ fashion as, we intentionally mix requirements, architecture and implementation activities in an agile manner. With this approach, we aim at synchronizing the requirements with the architecture and ensuring that the underlying implementation realizes and fits the architecture. Thus, this document reflects the current state at hands, while we aim at updating this document as part of improving the platform. In other words, in this document, we document and discuss the current state of the platform on a feasible level of detail, the underlying implementation, decisions we made and the tradeoffs that we faced. However, depending on the state of the implementation, this document is not meant to be complete but rather to be a “living document” that is updated incrementally. This version of the handbook focuses on the platform release as of **September 2026** (version **0.8.1**) and supersedes older versions of this handbook/the platform.

This platform release comprises multiple management UI updates, build process shortcuts for updating artifacts (e.g., for Python) and runtime service updates (initially only for Python).

1.2 Structure of the document

A typical first section of a platform handbook could be a summary of the requirements to be realized. As stated in Section 1, the IIP-Ecosphere team summarized the results of the requirements collection for the platform in two other whitepapers, namely the usage view [SSE21] and the functional/quality requirements view [ESA+21]. For pragmatic reasons, these two documents have been prepared partially before and partially while designing the platform architecture, so that they are synchronized with the work described here. In order to avoid inconsistencies, we are not repeating the requirements in this document rather than referring to [SSE21, ESA+21] through requirements identifiers defined there.

In Section 2 we introduce the tooling that is used for developing the architecture model and the implementation. A brief discussion of the tooling and the rationales for certain decisions is relevant at that point as the decisions significantly interact with the modeling concepts, i.e., affect the set of concepts that we practically can use for specifying, describing or realizing the architecture.

In Section 3 we introduce and discuss the architecture of the platform, ranging from the lower transport up to user-defined applications. This section is not only intended to present the architecture as it was designed rather than also the tradeoffs that we faced and the decisions that we made towards the actual architecture. In Section 4, we summarize architectural constraints that must be obeyed by the implementation. In Section 5, we discuss the representation of the platform components in terms of Asset Administration Shells, which are used as a uniform way to represent interfaces and communication among components.

One aim of our work is to research concepts on systematically and consistently configuring such a platform, ranging from network settings over available resources or services up to the wiring of reusable parts and IIoT-applications. In Section 6, we elaborate the structure of and the concepts of the model to specify decisions that must be made to turn alternative or generic components into an installable platform with user-defined applications. We will also discuss, how to utilize such a model, not only to validate configuration decisions, but, in particular, to automatically generate platform instances, artifacts or glue code as one means of supporting platform users to create IIoT-applications.

⁴ Along the realization state, i.e. the releases of the platform software. The version number of this white paper reflects the software release version. Thus, at the beginning some sections may be rather empty.

In particular for Sections 3 to 7 it is important to recall that the platform is currently under agile and incremental development, i.e., details and structures may change. Faster access to such information, we started turning modeling- and implementation level details into github documentation, which is easier and more agile to change than the handbook focusing on the more fundamental structures and decisions.

Ultimately, in Section 8 we summarize and conclude this document. In Section 9 we list references to other work that we rely on.

We recommend **different reading flows** for different audience groups:

- Readers with **architectural interests** find most relevant information in Section 3, which explains the layered architecture from bottom to top as well as the architectural decisions and constraints in Section 4. Potentially, also Section 2 is helpful regarding the utilized tooling and the basic technical decisions as helpful background.
- **Users**, who want to **install the platform** and do first steps are advised to start with the github online documentation. For first steps in configuring the platform or for creating applications on model-level, we recommend the online documentation on modeling concepts and properties as well as Section 6 as background. Please consider that at its heart, the platform is a distributed system and requires certain programs running on the nodes that shall perform application tasks. Depending on your installation, different user permissions also for installing programming language libraries may have to be considered.
- **Users**, who want to create **own applications**, are advised to read the online documentation on modeling concepts and properties as well as Section 6 for background on the configuration approach. If you are also responsible for (your own) software installations on your computer(s), please take the reading flow for installing the platform into account.
- **Users**, who want to use the **graphical management user interface**, we recommend the github online documentation for platform installation and modeling concepts as well as Section 7.3 and the modeling basics from Section 6. Further, the overview of the user interface in Section 3.13 is recommended.
- **Developers for services and connectors** shall know the platform architecture basics, i.e., Section 3.3-3.6 as well as how to develop applications, in particular for testing, i.e., Section 6 for configuration and creating own applications. Further, the online documentation, the guidelines and the code documentation are relevant.

As already indicated above, we are about migrate detailed technical information to github, e.g., besides the modeling concepts, the FAQ and the manual installation steps are now part of the online documentation in github⁵.

⁵ <https://github.com/iip-ecosphere/platform/blob/main/platform/documentation/README.md>

2 Tooling and Basic Technical Decisions

Tooling is an important topic when creating an architecture and when implementing it in terms of executable code. In this section, we briefly describe the tooling decisions made by the involved partners, as they affect the available options for modeling the architecture and for realizing it.

The **architecture** is designed using the Unified Modeling Language (UML) [UML]. We will not introduce UML in this document rather than assuming that the reader is sufficiently familiar with UML. In previous versions of this handbook, we used Eclipse Papyrus⁶, which did not scale anymore (display resolution, complex nested dialogs for templates, vector image export, behavioral diagram routing). Thus, we replaced all architectural diagrams by near-UML figures illustrating the most important structures of the respective component.

Along with the architecture and the design of individual components, also **architectural constraints** and, thus, **implementation rules** arise. We will discuss the architectural constraints of the platform in Section 4 as a specific summary of the architecture section. Section 3 may already indicate or mention such constraints.

For **implementing** the architecture, we must integrate existing components and consider that in particular AI services will be realized in different programming languages.

- For the **Java** components constituting the platform core and the connectors, we rely on Eclipse with Maven⁷, Git⁸ and checkstyle⁹ integrations¹⁰. Fundamental technical decisions are documented in the architectural constraints and, more detailed, in code. As we use Maven for the platform installation, a Java Development Kit (JDK) is required rather than a Java Runtime Environment (JRE). We just mention some of the decisions here: The dependency management and the build process are specified in Maven, thus, all dependencies must be available through official or own/local Maven repositories. The platform provides various own Maven plugins, realizing specific functionality like variability installation or end-to-end testing. Templates for code formatting and validation of the formatting are available for checkstyle in the source code repository as part of the most fundamental platform dependencies project and shall be applied prior to any commit. A common logging abstraction was realized and must be used at least in the platform core. Components of the platform are represented as individual projects using Eclipse as main integrated development environment (IDE). With version 0.8 of the platform, we upgraded the code to Java 17 and test against Java 21 (except for some components like RapidMiner Real Time Scoring Agent (RTSA) still requiring an installed JDK 8 for execution). For the continuous integration, the build/deployment process is specified due to technical reasons in ANT, partially setting CI specific variables, ultimately calling Maven.
- While some AI methods may also be realized in Java, nowadays AI methods are typically implemented in **Python**¹⁰. For Python services (as for Java-based services), a service execution environment is provided, which is responsible for the communication with oktoflow (there the Java service counterpart), so that an AI developer does not have to deal with both languages, protocol details or a plethora of different protocols. The service environment and the integrating Python services can operate with virtual Python environments. Python services

⁶ <https://www.eclipse.org/papyrus/> version 4.8

⁷ <https://maven.apache.org/>

⁸ <https://git-scm.com/>

⁹ <https://checkstyle.sourceforge.io/>

¹⁰ For the required versions, please see <https://github.com/iip-ecosphere/platform/platform/documentation/PREREQUISITES.MD>

must explicitly declare their dependencies, e.g., used AI frameworks as well as (if need) target virtual environments in the application model of the platform configuration to enable automated creation of installation artifacts, in particular containers, and execution in the desired Python environment as services and their dependencies may require the installation of different virtual environments or even Python versions.

- In particular, we prioritize **dependency reduction** over alternative, potentially more modern programming approaches as well as isolated loading of classes and their dependencies to cope with incompatible libraries. Thus, we decided not to use frameworks like OSGi or Spring as foundation as they may lead to (future) dependency conflicts, even among different versions of these frameworks needed in the same platform instance (as we experienced for Spring Cloud Stream as well as the AAS implementation Eclipse BaSyx and BaSyx2). Akin, we prefer in some places boilerplate code over annotation-style programming, e.g., in platform parts where a later revision with yet unclear external decisions can be foreseen. Where adequate, we leverage own, abstracting annotations and interfaces (as basic for plugins) and prefer them over technology-dependent annotations or interfaces.
- Some components require **technical settings** for their startup, e.g., certain internet addresses or basic security certificates to announce the own instance, to request or contribute information. The aim is to reduce such explicit setup information to a minimum as it is a source for inconsistencies. For this purpose, such information shall be managed centrally, instantiated into (binary) components or distributed via discovery protocols where feasible. So far, no automated discovery mechanisms (for I4.0) settings was integrated (BaSyx2 discovery could be an option if AAS metamodel version 3 is enabled), which could ease the setup. Further information not required to startup a component shall be made available via the (joint) AAS of the platform. Technical settings that may be subject to modifications by administrators shall be represented in a uniform and human readable manner. For stored setup information we rely on Yaml¹¹, for machine-readable complex data in AAS on JSON¹². Regarding terminology, we distinguish between **Setup** (the technical information, e.g., in Yaml, in practice often also called configuration) and the **Configuration** (the managing part, in terms of a configuration model, used for generating consistent setup information as well as related code and further artifacts). Related source code shall be named accordingly¹³.
- Components shall internally communicate via **interfaces** that encapsulate technical dependencies. Alternative and optional components shall be realized based on interfaces and register themselves into the platform. For Java, we use the Java Service Loader (JSL) mechanism¹⁴, which associates concrete implementations to their respective (descriptor) interfaces through text resource files declaring the actual implementations. We use that mechanism to define, e.g., plugin (setup) descriptors, factory instances, to compose AAS but also to set up the component lifecycle, e.g., to handle the start and shutdown process¹⁵.
- Since version 0.8, third party libraries must be encapsulated for **isolated loading** into **oktoflow plugins** and, as stated above, be used via an interface defined by the platform. The main reason

¹¹ <https://en.wikipedia.org/wiki/YAML>

¹² <https://www.json.org/json-en.html>

¹³ Initially, we aimed for an alignment with Spring, also calling the technical setups a “configuration”. However, this led to some confusion, so we decided for version 0.3.0 to refactor the platform code according to the setup/configuration naming convention introduced above. However, some parameter/variable names and comments may still use `configuration`, `config` or `cfg` where `setup` would now be correct. We will try to clean up these (local) inconsistencies incrementally over time.

¹⁴ <https://docs.oracle.com/javase/8/docs/api/java/util/ServiceLoader.html>

¹⁵ However, due to the oktoflow plugin mechanism, the Java service loader shall be created directly only if a specific classloader shall be applied. Otherwise, the platform classloader shall be used via `ServiceLoaderUtils.load`.

is to actively mitigate dependency conflicts. As a prerequisite, the platform core components must be free of third-party dependencies (except for the Java library). Further, plugins allow for isolated testing of the integration of dependencies and enable the individual evolution of technical dependencies. For flexible integration with the platform, oktoflow plugins also employ JSL, one descriptor for defining a plugin's contribution, a second determining how the plugin is actually loaded.

- All components shall provide sufficient **tests** for their functionality. Tests shall be executed during the **continuous integration** (CI) of the platform and also usual test metrics shall be recorded. **Test artifacts**, e.g., setup files created specifically for testing components or dependencies used only for testing, must be **strictly separated from production code**, e.g., reside only in test resource folders. In particular for Java components this is important as setup files that are accidentally placed in production resource folders may take precedence over generated setup information, i.e., prevent that the configuration decisions made by the user are enacted.

As stated in Section 1, for several reasons one objective of the platform is to use existing Open Source solutions wherever feasible. However, not all **Open Source licenses** are per se permissible in industrial contexts. Therefore, we reviewed Open Source licenses and categorized them into four categories:

- 1) Usable without limitations, e.g., MIT, BSD-2-Clause, BSD-3-Clause, ISC, CDDL1.0, Eclipse-Dist-1.0.
- 2) Permissible, but potentially problematic, e.g., Apache 2.0, LGPL-2.1, Artistic-1.0-Perl, EPL-2.0, MS-PL, MPL-1.1.
- 3) Commercial licenses.
- 4) Problematic and potentially not allowed (as default or core dependencies), in particular due to copy-left implications, e.g., GPL-2.0, GPL-3.0, EPL-1.0, AGPL-3.0. In some cases, the use of binary artifacts of software under such licenses may still be permitted as long as the license information and the origin are stated, the underlying code is not modified or included and the integrating component is optional.

These categories shall be considered already during the design of the platform and may effectively limit potential candidates. Licenses of the first two categories may be used (with care), the remaining shall be avoided. This is in particular true for platform components that constitute mandatory core functionalities of the platform. Commercial licenses may be used depending on the decision of the installing organization. Thus, platform components relying on commercial licenses shall be optional by default. Similarly, also software under not permissible licenses could be used in optional components, but to avoid later license conflicts, licenses of the fourth category shall be avoided wherever possible.

The source code of the platform is made publicly available in the **GitHub** space of IIP-Ecosphere¹⁶. Moreover, to foster transparency, the development of the platform happens in public. In the future, also the underlying architecture model shall be made available to support external and future developments after the project lifetime. As far as possible, components are subject to CI using the Jenkins server of the Software Systems Engineering (SSE) group at the University of Hildesheim. Upon successful builds, artifact snapshots are deployed by the CI processes to the Maven repository¹⁷ of the SSE group. Java parts including additional artifacts (binary, python, configuration model) of stable releases are deployed to Maven central¹⁸.

¹⁶ <https://github.com/iip-ecosphere/platform/>

¹⁷ <https://projects.sse.uni-hildesheim.de/qm/maven/>

¹⁸ E.g., <https://repo1.maven.org/maven2/de/iip-ecosphere/platform/>,
<https://search.maven.org/artifact/de.iip-ecosphere.platform/transport>

3 Architecture

The architecture of the oktoflow platform aims at realizing the requirements collected in [SSE21, ESA+21] as well as further requirements that are collected or detailed in further work or projects like ReGaP. In this section, we discuss the design of the individual parts and components of the platform. Please note that as mentioned in Section 1, we follow a pragmatic agile development approach, which involves forward and feedback cycles among requirements, architecture and implementation. Thus, depending on the realization state, not all platform components may be completely described in this version of the document, i.e., we will work out sections incrementally depending on the realization state.

We start in Section 3.1 with an overview of the platform layers and dive then into their details in the remainder of this document. At the end of Section 3.1, we detail some further basic aspects, namely, relation to reference architectures in Section 3.1.1, the concept of data flow processing in Section 3.1.2, a brief introduction into Asset Administration Shells in Section 3.1.3, high-level component interactions in Section 3.1.4, and the virtual character of the platform in Section 3.1.5. Section 3.2 takes up the general requirements from [ESA+21] as context for the platform architecture. As basis for the architecture description, we discuss then the layers of the platform, first as overview and then one section per layer, starting at the bottommost (generic) layer.

3.1 Overview

The overall architecture of the platform follows a layered style (see Figure 1) based on components and services (R4 in [ESA+21]). As far as feasible, we aim for a strict (logical) layering, so that for two adjacent layers l_l and l_u (with as “the lower layer” l_l being located below “the upper layer” l_u), only l_u (and not its transitive upper layers) shall access or call l_l directly. Moreover, there are also aspects that cross-cut visibly or invisibly in this layered structure.




Figure 1: Main platform components as block diagram.

Figure 1 depicts an overview of the high-level building blocks of the platform. The overall goal is to enable service-based AI-enabled applications to be deployed to available resources utilizing standardized protocols. Therefore, the platform encompasses on the two lower layers (middle of Figure 1) the platform-internal data transport, connectors to external data, e.g., to machines, basic service interfaces, resource abstraction and monitoring facilities. These two layers form the basis for units that can be deployed to edges, servers or to a cloud. The remaining layers add services, e.g., on data protection, data integration, AI, runtime adaptation/centralized monitoring, overall configuration

and deployment of services/containers. If permitted, services from those layers can be deployed to available resources (bottom of Figure 1) by the platform.

- **Asset Administration Shells (R7)** are used in the platform in two forms: On the one side, AAS represent assets created by different vendors (e.g., a product, a machine, an edge device, an already installed platform, a service, a storage mechanism or an App composed of services). The respective AAS may be provided by a third party. Also, the platform itself forms such an asset that deserves an own AAS. On the other side, we utilize the mechanisms of AAS also for describing interfaces of interacting components within the platform. These components may be internal or external, i.e., also interfaces provided by external AAS may be used. In particular, AAS (submodels) for internal components may be created for the purpose of internal communication rather than external component realization and, thus, may not follow official standardized formats¹⁹. An integration of AAS (different implementation frameworks in terms of plugins) as well as support for realizing (internal) Asset Administration Shells forms the bottom-most layer of the platform.
- The platform contains an event-based **transport messaging** mechanism, e.g., a Broker, so that components and services can communicate among each other independent of the layering. Although this implies certain degrees of freedom and may be used to bypass the indented layering (R7 [ESA+21]) in exceptional cases, the event-based messaging shall not happen in an ad-hoc or chaotic manner undermining the layer structure. Further, uncontrolled messaging may accidentally overload the broker(s), in particular if the broker is involved in the processing of soft-realtime data streams (one potential manifestation of R10 [ESA+21]). As event-based communication and data streaming are essential to the platform, they occur on one of the fundamental layers (Transport) utilizing the external components Broker and StreamingLibrary through plugins.
- **Variability management** and **consistent configuration** typically do also cross-cut layers, as variability instantiations may affect all components. This is already reflected in the requirements, where *configuration model* occurs in many different functional topics, see e.g., also for implicit information R8, R19f, R20, R28, R30, R31, R34, R40-R43, R62, R64, R73, R77, R80, R86, R89, R93-R101, R104, R107, R112, R119-R122, R131, R134 in [ESA+21], but also in the (variability-based) configuration model that crosses several building blocks/components in Figure 1. Moreover, some layers require access to the configuration, in particular at runtime, e.g., to determine whether migrations of components are needed or how adaptations shall be enacted. However, also here a chaotic use of the configuration can easily lead to unmanageable dependencies. Therefore, we modularize the configuration along the layers, and, if required, provide access to the individual configuration via respective interfaces. Also the configuration technology is encapsulated as a plugin. Similarly, only some few selected mechanisms to instantiate variability shall be utilized, in particular code generation, generation of setup files and artifact selection while packaging.

For short, the layers of the platform from bottom to top:

- **Support Layer:** The support layer (not shown in Figure 1) realizes basic abstractions and helpful functions for the upper layers of the platform. This includes logging, resource loading, plugin management as well as basic data format abstractions for YAML and JSON. Further, a set of utility functions for files, archives, system access or network are provided to reduce repetitions in higher layers. This also involves non-trivial management functions or functions to create

¹⁹ At the point of writing, several forms of AAS are in standardization, but most known to us do not aim at platform components. Wherever possible, we utilize existing standards, e.g., for device nameplates, or try to adopt the style of related standards to express proto-AAS, e.g., for software services.

common AAS structures and to foster internal conventions, e.g., how to represent certain information in AAS. Moreover, this layer contains an abstraction of AAS as well as AAS implementation plugins.

- Transport and Connectors Layer:** This layer is responsible for connecting devices among each other and with platform services using appropriate protocols and formats from the I4.0 domain. However, several protocols and formats impose different tradeoffs in functionality, performance, security and legal/normative impact. This layer integrates such protocols in a flexible manner. The role of the **Transport Component** is to abstract over relevant data transport protocols such as MQTT²⁰, AMQP²¹ or OPC UA pub/sub²² and their wire formats (e.g., JSON, XML, Protobuf²³ for MQTT), to provide implementing transport plugins and to integrate the abstraction with the streaming technology (StreamingLibrary). In contrast to recent platforms [SEA+20], where a single fixed transport protocol is not uncommon, we want to avoid making such basic decisions on behalf of the user already on this layer. Further, for the streaming technology several potential approaches with their tradeoffs are known. The idea is to prepare a flexible integration and to link this decision to the selected transport protocol. Similarly, connections to production machines and already installed platforms are abstracted by the **Connectors Component**. Such a Connector may utilize similar protocols as the Transport Component, but also protocols at higher semantic levels such as OPC UA or AAS relying on an own information model. In contrast to the Transport Component, projections and transformations of the original input data of a connector may be ingested into user-defined apps and information/commands originating from the apps may be transported back, e.g., to reconfigure an underlying machine.
- Services Layer:** Openness and extensibility through services of different kinds, in particular AI services, are at the heart of oktoflow. To be useful for an application, services must be parameterized and orchestrated, e.g., their data (streams) must be connected to other services or connectors. While the interconnections will be handled by the Transport and Connectors Layer, the Services Layer defines the basic service interfaces (Services) as well as the **services execution environments**, e.g., for Java and Python. On this layer, Connectors are wrapped into services so that they can be used seamlessly in user apps. Services may be realized in different programming languages and, thus, demand different integration capabilities, ranging from direct calls (Java services) to communicating operating system processes (Python services, GO, or even standalone Java programs). Services are wrapped by code generation into service units for a certain **service execution**. A specific service execution, e.g., Spring Cloud streams, is an implementation plugin of the services layer.
- Resources and Monitoring Layer:** To become effective, services must be deployed to resources/devices (in terms of a Deployment Unit) and monitored at runtime. In the platform, deployment targets such as edge devices shall describe themselves in terms of AAS and perform a registration with the device registry (Devices), which reflects its data into the runtime structures of the platform and the platform AAS. For deployment, the Deployment Unit (called “ECS runtime” in [SSE21]) receives commands via its AAS from the platform and, in its containerized form, downloads a container including installed components and starts the container (ECS implementations plugins for Docker and LXC are part of oktoflow). In this containerized environment, apps are then started through the constituting service

²⁰ <https://mqtt.org/>

²¹ <https://www.amqp.org/>

²² <https://opcfoundation.org/news/press-releases/opc-foundation-announces-opc-ua-pubsub-release-important-extension-opc-ua-communication-platform/>

²³ <https://developers.google.com/protocol-buffers>

implementations²⁴ by the service execution. Also the execution of the services in the container must be monitored, which may involve reusable monitoring probes provided by the platform as well as application-specific probes. The reusable mechanisms are provided by the Monitoring component, which (in terms of probes and signaling) is part of the service environment while the aggregation of the monitoring data happens on central IT level (a default realization in terms of Prometheus²⁵ is a monitoring implementation plugin of oktoflow). The Monitoring component also uses the capabilities of the support layer (monitoring in terms of AAS) and the Transport and Connection layer (fast track signaling, alarms) and may issue alerts in generic as well as application-specific manner to further layers.

- **Storage, Security and Data Protection Layer:** Security and data protection in the platform encompass of two parts, 1) cross-cutting mechanisms that can be used to implement security and data protection in any component, e.g., authentication, and 2) centralized or distributable mechanisms to support security and data protection, e.g., services supporting data protection or data storage. While the cross-cutting mechanisms occur in all layers (directly or indirectly controlled through the platform configuration), this layer primarily focuses on the second part. Thus, this layer realizes in particular components (optionally) enhancing the security and data protection as platform-supplied app services, e.g., for Anonymization and Pseudonymization.
- **Reusable Intelligent Services Layer:** The components described so far (as well as not mentioned administrative services provided by the platform) can be used to develop applications similar to existing platforms [SEA+20]. This layer shall pave the way for open, extensible and reusable intelligent services. In particular, the AI - Toolbox contains re-usable AI services that can be parameterized and orchestrated to form a running application, e.g., Federated Learning with flower based on generated Python services, an optional integration of the RapidMiner RTSA as generic, re-usable AI service or a re-usable basic data processing library.
- **Configuration Layer:** The configuration layer contains components to manage the platform configuration. The Configuration and Instantiation component is responsible for composing reusable and application-specific services and representing the information in terms of the application parts of the platform configuration. The Deployment component is responsible for deciding which services shall be executed by which device (e.g., edge, server or cloud) depending on runtime information available in the platform configuration. Based on these decisions and device-specific information provided by a device AAS, containers are created automatically and made available. In particular, this involves code generation of various artifact types, from app/service code templates over integrating service wrappers for the service execution to build specifications. Further configuration operations target the re-configuration of services or the runtime-selection of alternative services.
- **Applications Layer:** Applications are described in the app part of the configuration model and may ship with application-specific components, e.g., AI services. Although not visible here, glue or transport code generated for services implicitly belongs to the apps. The execution of the apps shall be visualized by (as far as feasible) generic Dashboard components.
- **Management User Interface:** Ultimately, a Web user management interface (UI)²⁶ relying in particular on components of the Configuration layer and the AAS of the platform allows for configuring apps, supports the implementation of apps as well as their distributed execution. It is important to emphasize, that although the management interface is realized as

²⁴ Assembling the containers is managed by the Configuration Layer as described below.

²⁵ <https://prometheus.io/>

²⁶ As discussed in [ESA+21], user interface and dashboards are formally out of scope of our funding contract. However, if feasible, we plan to realize at least a simple (Web) user interface in one of the next releases.

a Web UI, the platform must not necessarily be installed/deployed in a Web/Cloud setting, i.e., on-premise installation and use of the Web UI via a browser is one important installation alternative for oktoflow.

The platform may or may not interface with Clouds or dataspace as desired by the user, e.g., to not include/remove respective connectors and components completely from the individual platform instance/apps upon platform instantiation.

The full stack shown in Figure 1 is not required for all kinds of installations. E.g., on a resource such as an edge device, a cloud or a server, a **specialized runtime** is needed (ECS-runtime from [SSE21]) to take control over containers and services, while monitoring, device management and platform AAS shall be running on central IT. For example, the service manager can be composed from a subset of the layers as indicated in Figure 2, in particular support, transport and connectors and services (using the respective oktoflow plugins indicated in light blue/italics). Similarly, the ECS-runtime, in particular its variant including the service manager can be composed from lower layers and the respective components from devices and monitoring. For managing containers, at least the deployment unit (implemented as ECS-runtime plugin) from the Resources and Monitoring Layer is needed. Service manager and ECS-runtime can run in the same container/on the same device, as individual processes or combined. However, ECS-runtime and service manager may also run as individual containers, the one for the service manager then also containing all dependencies that apps do require, e.g., respective Python installations. The platform monitoring component can be instantiated as individual service, in Figure 2 intentionally without plugin, i.e., not relevant. On top, the “platform service” hosts the AAS with device management, excluding ECS-runtime and service manager.




Figure 2: Layers, components and plugins required to build a service manager, the ECS-runtime and the platform service.

3.1.1 Relation to Reference Architectures

The platform aims at interrelating and adhering to reference architectures such as RAMI 4.0 [RAMI]. Although we use an own naming of the platform layers, they map nonetheless to layers defined by RAMI 4.0 as summarized in Table 1. However, it is important to recall that the platform was initially planned to be a virtual platform, i.e., it shall be able to build on existing installations without implementing a complete IIoT platform stack. Thus, it is for us not relevant to meticulously adhere to all RAMI levels, in particular not to the lower levels targeting field devices (as already scoped out in [SSE21, ESA+21]). In addition, our architecture includes some (crosscutting) layers that do not directly fit into the picture of RAMI²⁷.

Table 1: Mapping RAMI 4.0 and the platform architecture

RAMI 4.0 Axis	RAMI 4.0 Level	Oktoflow Layer/Component
Layers	Asset	Not in scope [SSE21, ESA+21], represented through edge AAS
	Integration	Support Layer, Transport and Connectors Layer
	Communication	Services Layer
Hierarchy Levels	Information	Reusable Intelligent Services Layer
	Functional	Application Layer
	Business	On top of Application Layer via Applications AAS
	Product	Not in scope, represented by data
	Field Device	Not in scope [SSE21, ESA+21], represented through edge AAS
	Control Device	ECS-runtime [SSE21] with deployed services, in particular Resources and Monitoring Layer with contributions from upper layers
	Station	ECS-runtime [SSE21], possibly with access to more powerful resources or UI capabilities for executing or controlling deployed services. Includes Resources and Monitoring Layer with contributions from upper layers.
	Work Centers	Reusable Intelligent Services Layer, in particular Data Integration component
	Enterprise	Application Layer
	Connected World	On top of Application Layer via Applications AAS, including connected platforms.
Life Cycle Value Stream	Type	Component and AAS types prescribing structures
	Instance	Deployed component and AAS instances

In terms of the Industrial Internet Reference Architecture [IIRA], this document can further be understood as a continuation of the usage view(point) [SSE21], the functional view [ESA+21] In terms of a platform architecture as well as its implementation.

3.1.2 Stream (Data) Processing

In an IIoT/Industry 4.0 setting, often the processing of data is viewed in terms of streams of data items (or tuples), e.g., produced in regular fashion by a machine, taken up by edge devices for pre-processing, protocol transformation or retro-fitting, handled further by other devices and (partially) stored in some data stores, e.g., time series data bases. In contrast to other forms of data processing, e.g., batch processing, data stream processing can fulfill (soft) realtime requirements, of course, depending on

²⁷ Crosscutting aspects are better covered by IIRA [IIRA].

the data ingestion frequency (overload, backpressure) and the (relative) speed of the individual data processors.




Figure 3: Viewing IIoT and Industry 4.0 as data streams.

Figure 3 illustrates the basic components of such a stream processing approach, considering “the machine” on the left side as constant (conceptually endless) data source. The data tuples/items produced by the machine is taken up by a data transformer (e.g., preprocessing, anonymization), passed to a second transformer (e.g., AI service) and finally to a sink (e.g., data store, dashboard). From a different point of view, the data flows in forward manner from source to sink. The edges in such a graph indicate the data flow and the nodes the data processors. There could be more processors, different kinds of processors or more complicated forward flows that we do not touch in this brief introduction. Please note, that there is no need for synchronous processing in the nodes, in particular in the transformers. With synchronous processing, we mean that a transformer operates like a mathematical function, i.e., for an input tuple it produces in the same step an output tuple. In contrast, asynchronous means that the processor receives data item(s) and at some future point in time it may emit any number of tuples (including none at all).

In Figure 3, there are also two horizontal lines, indicating borders of physical devices, e.g., the first two streaming components could be running on an edge device, the second transformer on a further device, and the sink on a third device, e.g., a central server. The distribution of components is not fixed, e.g., depending on resource usage, the second transformer could also be executed on the first device or the first transformer on the second device.

Several approaches to stream processing rely on untyped data, i.e., the transformer implementation decides based on the available data fields, what to process. Such an approach can easily fail at runtime, when processing nodes are combined that cannot work together, with negative outcomes ranging from loss of data to runtime errors or exceptions. In contrast, we rely on typed data flows, i.e., for each forwarding edge the type of data item(s) is known during design and built into the respective app. As the design of data processors and data flows will be captured in the configuration model of the app, checking for type and streaming compliance before realizing or instantiating the system becomes possible. In Figure 3, the forward flow indicates three data types, t_1 , t_2 and t_3 . Please note that depending on the requirements and the design of the data processing, the types may be the same or they could differ, e.g., indicating that a processor adds or removes data fields.

While in many applications, a forward flow is sufficient, in particular in IIoT/Industry 4.0 settings it could be desirable, that an upstream processor shall send back data to a downstream processor, e.g., a decision node after one or multiple artificial intelligence nodes shall inform the machine at the data source that some processing parameters must be changed. Akin to the forward flow, we allow types for backward flows. It is just a matter of modelling convenience that we define the forward flow in terms of nodes and connecting edges, while we consider the backward flow as typed notification data channel (t_4) of one or multiple senders and potentially multiple receivers.

3.1.3 Asset Administration Shells

The platform aims at complying with, integrating of and extending existing standards and technologies in I4.0 (R7, R14). This applies to protocols, formats but also model standards such as the Asset

Administration Shells (AAS). For short and without aiming for a complete description, an AAS is an information model, which targets a physical or virtual asset in terms of nested, detailing sub-models. Sub-models may consist among other kinds of elements of typed properties, operations and heterogeneous collections/lists of sub-model elements. AAS and sub-models can be classified as static (all information is determined when creating the AAS), dynamic (some information may change at runtime) or active (callable operations are provided). Similarly, properties and operations can be static or dynamic, whereby in the dynamic case both element types can be linked to an implementation, e.g., provided by a remote implementation server, and, thus, change value (access) or implementation over time. In particular, AAS for different assets of different vendors can be provided, related and integrated, e.g., to link the AAS of a device utilized by the platform into the platform AAS to provide, e.g., a digital nameplate for industrial equipment [BBB+20, ZVEI-N] or the documentation of the device at hands. Moreover, composite AAS can be created, representing, e.g., a complex machine consisting of AAS of the utilized components.

According to the requirements (here R7), the platform shall describe all (distributable) components, interfaces, functions and deployment targets in terms of AAS. Thus, each of the components of the platform that forms an individual asset (of a certain vendor) shall receive an own AAS (as indicated in Figure 1). Moreover, the platform itself shall provide an own AAS and each of the discussed layers shall provide one or more sub-models to link the layers against each other (whereby the sub-models may and shall link to the vendor AAS of the individual assets, e.g., edge devices). As far as feasible, the platform will utilize existing approaches and standards to define the AAS, but also define own ones where needed, e.g., to characterize the capabilities of deployment targets such as edge, server or cloud devices [SSE21]. We will detail the platform AAS and its structure in Section 5.




Figure 4: AAS deployment options (D1 remote deployment, D2 local deployment)

As typically several distributed compute resources are involved in a platform installation and each compute resource shall be used for running oktoflow apps, it shall be described/registered with an own AAS (model, sub-model or as part of joint model/sub-model). Therefore, it is helpful to introduce two basic AAS and component deployment patterns. Figure 4 illustrates the central IT side (the “Platform AAS server”) and two distributed resources *D1* and *D2*, e.g., edge devices. An AAS can be served locally and only be registered in a central registry or it can be deployed remotely to a central server. Serving an AAS locally requires a related web server process (“Resource AAS server” in *D2*), i.e., a further process to be executed on a resource. Deploying an AAS centrally avoids such local server processes, but may lead to increased communication with the central server and, in the case of dynamic or active AAS that allow for operation calls, also to redirections of requests via the central server to the resource (which may anyway be the default behaviour of an implementation, e.g., BaSyX2 operation delegation). To handle requests of dynamic or active AAS, the resource must run a (further)

server instance, the “Resource AAS command server”. A similar server process must exist on the central IT side of the Platform AAS server to offer dynamic properties or operations. In the resource case, this “Resource AAS command server” may forward operations to further processes, or, if the processes are already known when the resource AAS is constructed, also specific server processes, e.g., for the service manager running in an own container, can be linked to the AAS and directly contacted to serve AAS requests.

3.1.4 Component Interaction Overview

In the previous sections, we introduced the layers and the high-level components of the platform as well as the basic concepts of AAS. In this section, we provide a brief overview on the component interactions for a basic walk-through of platform operations. The individual sections on the components in Sections 3.3-3.13 will provide more detail on the interactions. In addition, Section 3.14 will address the cross-cutting topic of testing support for services and applications.

The aim of this walk-through is to bring up the ECS-runtime, the service manager (in terms of a container), some services, to let the services run and to stop all parts in reverse order. Services are modeled as a service mesh forming individual applications (we will detail how to define such a mesh in Section 6). The required high-level interactions are illustrated in the sequence diagram in Figure 5. We will go through them now from top/left to bottom/right.




Figure 5: High-level component interaction for basic platform interactions.

1. At the beginning, the platform AAS-Server is running. An ECS-runtime is started for a certain resource, e.g., an edge device. The ECS-runtime instance then deploys its own sub-model characterizing the device with container operations and a collection of available containers (initially empty) into the platform AAS. A scheduled background process of the ECS-runtime is started to inform the platform AAS about the actual resource state (resource monitoring, not shown in Figure 5). Depending on the device, the ECS-runtime may provide information about

an existing device AAS or create a device AAS on its own (one particular point of openness as the device vendor may or may not provide an AAS). This information is linked from the platform AAS.

2. Via the user interface (UI), the user requests a list of available resources. The UI reads out the AAS submodel for resources including the ECS-runtime instance started in step 1 and displays device information including the actual resource usage. In a similar manner, further information can be obtained, e.g., the available services, the defined applications, the packaged service artifacts or the available containers.
3. The UI requests adding a container via the ECS operations known to the platform AAS, leading to a remote method call to the ECS-runtime (AAS implementation server). For this walk-through, we assume that the container contains the service manager and provides the technical dependencies for services to be executed on the respective device. Starting a container may lead to a download of the pre-built container from a central platform server (indicated by #) or from a file system of the device. Information about the container instance is made available to the platform AAS by creating a structure in the containers submodel of the platform AAS.
4. The user requests starting the container added in step 3, i.e., the UI calls the respective platform AAS operation, leading to a remote call to the ECS-runtime (AAS implementation server), respective operations in the container management implementation, e.g., Docker, and, ultimately, when the container is running, to an automated start of the service manager. In turn, the service manager deploys information about itself, e.g., service operations, into the platform AAS, more precisely into the device entry created by the ECS-runtime so that services on the underlying device can be managed.
5. So far, no app/service is known. The user requests to add an app via an operation of the platform AAS, leading to a remote method invocation to the Service Manager (AAS implementation server). In turn, as for the container, the service manager may download an implementation artifact containing the app, the services, and the related service execution environments for the actual device. The service manager adds entries for the artifact and all contained services to the respective sub-models of the platform AAS.
6. The user requests the start of the app, i.e., all services for the device addressed in the steps above through a deployment plan. The involved Service Managers start the service environment and creates the service instances in the sequence of dependencies, i.e., starting with the service having no data dependencies or for which all prerequisite services are already running. During this step, several network ports may be acquired for internal communication, relations to a global or a local protocol server/broker may be established and individual operating system processes for the services may be started. Further, connector settings may be adjusted to device-provided services as stated in the device AAS, e.g., specific IP address, actual port or even the startup-time selection of the right protocol implementation depending on the device-specified protocol version, e.g., MQTT v3 vs. MQTT v5. These detailed technical procedures are not shown in Figure 5. During service startup, the Service Manager checks the service relations in the platform AAS (services sub-model) for service availability and, as soon as the service is up, creates a relation entry linking two subsequent services in the service mesh of an IIoT application running on the platform.
7. The services are running now, receiving data via the machine/platform connectors, executing functionality specified for the actual application, e.g., AI-based inference. During the execution, background processes collect data for the device and the individual services and inform the platform AAS about actual runtime states, e.g., resource consumption. Here we also indicate in Figure 5 the resource monitoring of the resource mentioned in step 1.

8. The user requests to stop the running app via a respective operation of the UI/platform AAS, which causes a remote method invocation to the Service Manager(s). In turn, the Service Manager removes the service relations in the platform AAS and stops the service environment and the services.
9. The user (directly or through the deployment plan) indicates that the artifact will not be used any longer, i.e., a platform AAS operation is called and causes a remote method call to the Service Manager, which removes service and artifact entries from the platform AAS.
10. As also the service management container shall not be used anymore, a command from the UI to the respective AAS operation leads to a remote method call to the ECS-runtime, which commands a stop of the container through the underlying container implementation.
11. Ultimately, the container shall also be removed from the management realm of the device, leading to a further remote method call to the ECS-runtime, performing a removal of the container information from the platform AAS.

The horizontal dashed, red lines in Figure 5 indicate phases of the operations, i.e., startup (step 1), preparation of containers and services (steps 2-5), service operation (steps 6-7), shutdown (steps 8-11). The vertical yellow dashed lines indicate a potential distribution to different logical or physical devices. Extreme cases are that all components run on the same device, e.g., for testing, or that UI, platform AAS, ECS-runtime and service manager/services are installed on separate devices.

It is important to emphasize that the “user” in this walk-through may be a human, a deployment plan selected in the UI or the platform itself acting on behalf of the user. A deployment plan lists the assignment of containers and services to resources so that the UI can execute the desired deployment automatically.

3.1.5 Virtual Character of the Platform

As stated in Section 1, the platform shall also have the character of a **virtual platform** (R3), i.e., a platform that offers services on top of existing already installed platform functionality. The idea is that the Connectors component in the Transport and Connection Layer maps relevant underlying platform information and functionality into the platform. Where feasible, this mapping shall happen in the form of AAS as it allows for an overarching information model, but also further approaches like OPC UA or MQTT may be used. We see here three alternatives, focusing on AAS as the default approach, potentially using a transport protocol like MQTT:

1. An underlying platform provides its own AAS and manages the access to selected functionality and data. Theoretically, this AAS could be mapped side-by-side into the AAS of the platform. Then, layers such as deployment device management, or monitoring could directly utilize the information. Therefore, a standardized AAS structure for manufacturing platforms would be desirable, but, as far as we know, such a standard currently does not exist.
2. The AAS connector of the platform can map the AAS of an underlying platform into oktoflow. Of course, this may add additional overhead and in some cases a mapping may not be possible at all.
3. One of the other connector types provides a protocol that allows mapping the underlying platform and its operations into the platform AAS. This approach may require manual programming, while the second approach might be realized easier through mapping and code generation.

Besides having access to the AAS of an underlying platform, relevant components of the platform, in particular the resource management and monitoring component are required to operate with multiple AAS instances (for now based on the platform AAS structure).

3.2 Overall Requirements

In general, all platform layers and components discussed below must take the following general requirements from [ESA+21] into account:

Table 2: General platform requirements in [ESA+21]

Requirement	Summary
R1	Vendor and technology neutral platform
R2	Use of standards
R3	Design as a virtual platform
R4	Design based on components and services
R5	Use of Open Source, with respect to the licensing rules of the platform
R6	Open for optional/commercial components
R7	Use of AAS for interfaces
R8	Use of systematic variant management techniques
R9	Means for availability
R10	Soft realtime processing (<100 ms) for production-critical functions
R11	Documentation (also in terms of this handbook)
R12	Documentation of processing steps (of applications, supporting data privacy)

As already indicated in Table 2, [ESA+21] also specifies quality requirements such as R10. Besides security and data protection requirements, there are also data frequency and volume requirements that are not so obvious, in particular as they are assigned to specific topics/components of the architecture in [ESA+21]. To provide an overview, we discuss them here on a global level for the entire platform.

In Table 3, we summarize the cross-cutting quality requirements, i.e., in particular those that may require specific considerations regarding time-critical functionality such as the (stream) processing or data transport. Although the platform aims at the deployment of components to edge devices, both, the services as well as the platform operations belong to the IT realm so that OT requirements such as R35 or the OT sensor sampling frequency mentioned in R28 do not directly apply. However, a machine pulse of 8 ms (R28), an hourly throughput of 7 GByte as well as an expected size of data items with 50 values (R19a) are highly relevant for judging the performance of the platform. As also mentioned in [ESA+21], not all data volume and frequency requirements were indicated while collecting the requirements from the partners, i.e., the platform shall aim for even higher speed (such as a 50 ms cycle time) or a throughput of 600 GByte per day.

It is also important to recall from [ESA+21], that the platform is primarily responsible for its mechanisms and included services, i.e., providers for services to be packaged with the platform will have to obey the quality requirements in [ESA+21]. Further, as also discussed in [ESA+21], the platform is not responsible for the quality of external services, e.g., application-specific or user-specific services (while measures may apply to report or terminate services that potentially taint given runtime requirements).

Table 3: Overview of (global) quality requirements on data frequency and volume

Requirement	Summary
R10	Soft realtime, response time < 100 ms for production critical functionality
R19a	Sample data set of 50 values of different types all 20-30 s
R19e	Output data shall be provided all 5 s
R21	Low impact on data throughput
R22	Overall platform throughput of 500 GByte per year

Requirement	Summary
R28	OT sensor sampling frequency 0.2 ms, machine pulse 8 ms, step pulse 5 s, process pulse 25 s (mentioned in the explanation of the cloud requirement R28)
R35	OT sampling frequency of 2 ms
R91	7 GByte per hour as input for data integration, which may be aggregated to 2 Gbyte per hour.

As an illustration, we discuss the quality requirements now in terms of hypothetical numbers. From the data transport perspective, the requirements command that each machine can ingest a data item with around 50 values each 8 ms, i.e., 125 messages per second. This leads to at least 450.000 messages per hour (per machine/edge device). If we assume a size of 654 Byte payload (actual size of a simple JSON serialization of such as message), a data source produces around 280 Mbyte per hour (just focusing on the raw data payload, i.e., not on additional information, e.g., for routing or meta-information as stated in R79). On a platform-level (R91, R22), aggregating components of the platform will have to cope with multiple parallel streams of this kind, which requires 26 such streams to reach the requested 7 Gbyte (in a real setting with payload and overhead). Of course, the distribution may be different, i.e., more streams at lower ingestion frequency or less streams at maximum frequency, potentially with image payloads, to reach several hundreds of GBytes per hour.

In the discussion of the individual layers/components, we will refer to these general requirements and re-iterate the argumentation only for affected layers or layers that already have been (initially) evaluated.




Figure 6: Structure of the Support Layer, core components left and plugins right (not all plugins are connected)

3.3 Tools Layer

The Tools Layer realizes specialized Maven build plugins to realize particular aspects of the build process of oktoflow components, plugins and apps. These build plugins focus on:

- Resources, in particular for building and installing oktoflow plugins.
- Python, for validating Python source code and for executing Python unit tests.

- Inovation, to steer complex app build processes, in particular to execute the platform and app instantiation process (realized as Maven build plugin on configuration layer, see Section 3.10) specified in terms of Maven profiles.

Besides these three fundamental build process plugins, this layer contains also a small common code library, the `tools.lib` that is used by the build processes as well as the platform, e.g., containing the `PythonUtils` to determine which Python executable to invoke.

3.4 Support Layer

The Support Layer aims at providing useful common functions and abstractions to ease the realization of the platform. Thus, it is more a support library than a full layer, i.e., it does not provide an own AAS representing the interface of the layer. Below, we first discuss the structure of the whole layer, than it's four main components and finally, in Section 3.4.6 the recommended approach to implement platform AAS as well as in Section 3.4.7 the plugins realized for this layer.

3.4.1 Component Structure of the Support Layer

As illustrated in Figure 6, the most abstract component in the Support Layer is `support.boot`²⁸, which introduces the plugin mechanism and the resource loading as well as the fundamental plugin interfaces for logging, common operations, JSON and YAML.

The `support` component adds plugins that (partially) depend on the plugins introduced in `support.boot` as well as further common mechanisms. `support.aas` defines the AAS abstraction, i.e., the plugin interface for Asset Administration Shells. Further, `support.iip-aas` are specific AAS support functions including the AAS-based component lifecycle support as they are used in oktoflow (already introduced in IIP-Ecosphere, thus, “iip”).

3.4.2 The `support.boot` Component

The `support.boot` component introduces the most basic mechanisms including some common functionality classes for collections, file/zip access, JSL, basic network functions, exception-enabled functional interfaces, etc. Moreover, this component defines the plugin interfaces for basic technical dependencies, such as common functionality, logging, YAML and JSON.

Figure 7 depicts a coarse-grained summary of the structure of `support.boot`. Below we focus on the plugin manager, the resource loader and the task tracking realized in `support.boot`.

²⁸ `support.boot` was introduced in version 0.8 while `support` already existed; `support.boot` contains fundamental parts without using plugins while `support` uses the plugins defined in `support.boot`. As used in many build specifications, changing the name of `support` was considered too dangerous. As all `support` components share the same Java package namespace, future versions may move functionalities that are still in components based on past decisions.




Figure 7: Simplified Structure of `support.boot`: core implementing classes, plugin interface (lighter background), delegating frontend utility classes like `IOUtils`. Plugin descriptors are not shown.

3.4.2.1 Plugin Manager

While most of the alternative oktoflow components can be combined without dependency or classpath conflicts, some (versions of the same) components would introduce conflicts, e.g., different versions of the AAS implementation BaSyx. To prevent that the development of oktoflow and oktoflow apps is forced by external configurations to certain dependency versions (the opposite direction is not realistic), we introduced a simple plugin management mechanism in version 0.8. Although proven implementations of such capabilities do exist, e.g., OSGi, and to prevent unpredictable conflicts with actually used and future dependencies, we decided to rely on a rather, simple classloader-based mechanism based on two JSL descriptors.

These are the `PluginSetupDescriptor`, which creates the plugin classloader and the `PluginDescriptor`, which creates specific instances of the plugin. The classpath of a plugin may rely on common components of underlying architectural layers, but not of higher layers, alternative components or other plugins (see also architectural rule C1). The `PluginManager` loads these descriptors and makes instances available through unique plugin identifier names declared by the plugins. Plugins may have a single identifier as well as multiple alternative identifiers, which eases migration from classname based instance creation of the previous platform version to plugins in this version. Different forms of `PluginSetupDescriptor` do exist, most are based on a persisted list of dependencies created during the build process of the plugin. Some example descriptors are: Loading of unpacked plugins from the file system (`FolderClasspathPluginSetupDescriptor`), from already loaded classpath resources or from FAT²⁹ plugin assemblies containing an extended classpath file (`ResourceClasspathPluginSetupDescriptor`). For installation, the platform instantiation shall obtain and unpack the plugins as determined by the configuration model so that the setup descriptors (as determined in the respective extended classpath file) can take them up.

This combination of JSL descriptors allows for:

²⁹ FAT Java archive files (JARs) are specialized ZIP files in which dependencies are included, either as contained JAR files or dissolved into individual files or folders.

1. Separate, priority-based class loading for isolating plugins that require potentially conflicting dependencies.
2. Limited class loading while running the plugin as an own JVM process, e.g., in case of server instances with heavily conflicting dependencies.
3. Proxy plugins using the same classloader to enable a unified plugin architecture, e.g., if similar alternative components are loaded through (and require) priority classloading while others use plugins internally or are free of conflicts.

Java class loaders are organized hierarchically and typically the class loader of the actual class also loaded the application, i.e., may be suitable as parent for isolated class loading. However, this is not always correct as, e.g., in server environments sometimes the so-called context class loader represents the application class loader. To select the correct class loader, the class `PluginSetup` defines the class loader for oktoflow. Dependent on the startup, e.g., of apps, this class loader may be re-defined adequately and shall be used for any dynamic loading operations.

3.4.2.2 JSL support

Initially to cope with different versions of JSL in different JDK implementations, we created a set of utility methods in the class `ServiceLoaderUtils`. Since the introduction of the plugin mechanism, the class loader to be used for JSL is also taken from `PluginSetup`, i.e., all service loaders in the platform shall be created via `ServiceLoaderUtils`.

3.4.2.3 Resource Loader

In many cases, platform components rely on (file) resources that must be resolved and loaded at runtime. In Java, this usually happens via the class loader (in oktoflow the one in `PluginSetup`), i.e., Java archive files (JARs) contain such resource files and the Java class loading mechanism provides transparent access to them. However, besides the standard path starting at the root of such an archive file, in some cases the packaging of FAT JARs may dictate further paths. In the platform, FAT JARs are an alternative for packing app services into service artifacts. As an unknown number of additional resolution strategies may be required, `support.boot` realizes the `ResourceLoader`, which allows registering additional `ResourceResolver` instances directly or via JSL. All platform components are encouraged to utilize the `ResourceLoader` or to contribute required resolution strategies.

3.4.2.4 Task Tracking

For longer running tasks, such as service deployments, tracking and reporting the state of the execution to the user is required, e.g., on a user interface. However, the platform is a distributed system, i.e., task information must be passed among the executing resources in a manner, that multiple resources can collaborate on the same task. For this purpose, the oktoflow platform provides a thread-based task tracking mechanism.

3.4.3 The support Component

Below, we detail the `support` component as illustrated in Figure 8. Besides further plugin interfaces for REST (server), HTTP (client), websockets, ssh, Java bytecode manipulation and runtime metric probes, this component also implements basic program setup classes, a registration mechanism for already installed programs/dependencies required by the platform or services as well as a runtime data collector for tracing build and testing times. The mechanisms defined in this component are allowed to utilize the plugins from `support.boot`, i.e., logging, commons, YAML, Json, also for testing.




Figure 8: Simplified structure of support: core implementing classes, plugin interfaces (lighter background), utility classes like `PidFile`. Plugin descriptors are not shown.

3.4.3.1 Setup

As stated above, we differentiate between a Configuration (overarching, for the entire platform and its applications, basis for code generation) and Setup (instantiated configuration information to be read by platform components upon startup). Inspired by the Spring framework, we represent the startup information uniformly in YAML. Further a simple mechanism to uniformly read command line arguments (also in Spring style) is implemented here.

3.4.3.2 Installed Dependencies

Services may require individual technical dependencies, e.g., a certain version of Java or Python. However, operating systems like Linux ship with certain versions of these dependencies, which, in turn, affect also the available versions of related libraries. Similarly, Windows makes assumptions and even, during updates, suddenly tends to switch to newer versions, e.g., of Python. To allow the platform and in particular the applications and services to run with expected versions of such underlying technical dependencies, we introduce the “installed dependencies” mechanism. This is a simple YAML file which declares the paths of the executables for certain keys, e.g., `PYTHON39`. As services in the configuration model of the platform can specify also the required system dependencies, the service implementation can request the actual path from the installed dependencies and execute the respective binary. In particular, for generated containerized environment, the platform ensures that the required dependencies are installed and respective paths are declared in the installed dependencies file. For application testing, it is important to know that the installed dependencies YAML file is searched primarily in the Java classpath, the operating system root (intended for containers) and in the location specified by the Java system property `iip.installedDeps`.

3.4.3.3 Other classes

`support.support` also defines abstract file formats, a `Version` class (representing platform and service versions), the `PythonUtils` and the `dependency Updater` for plugin dependencies. While some of these classes may be migrated in the future to `support.boot`, the `Updater` is one class which needs `Json` functionality, i.e., an implemented plugin, and, thus, must reside one “layer” above the required plugin interfaces.




Figure 9: Simplified structure of `support.aas`: AAS abstraction, core implementing classes, plugin interfaces (lighter background) and the platform lifecycle mechanism. Plugin descriptors are not shown.

3.4.4 The `support.aas` Component

Reusing the mechanisms of `support` (and transitively of `support.boot`), the `support.aas` component introduces the interfaces for over abstracting Asset Administration Shell implementations, the identity (authentication) mechanism as well as basic mechanisms to be used as implementations for AAS submodels, e.g., system metrics, the distributed network manager and semantic id resolution.

3.4.4.1 Asset Administration Shell Abstraction

A core aim of this component is to abstract over the used AAS implementation. This allows for flexibility (the AAS implementation can be exchanged), embraces oktoflow’s plugin mechanisms for isolating technical dependencies, but also to mitigate risks of impacts by the currently evolving AAS standard and its implementations. Thus, the abstraction described here aims at supporting the application of AAS for the description of interfaces (R7), the application of standards (R2) and enables openness for different AAS implementations, including potential upcoming commercial implementations (R6).

Currently, we employ BaSyx1 as the default AAS implementation of oktoflow and, besides three different versions of BaSyx1 (AAS metamodel/API v2) realized as plugins, we also provide in the same fashion an integration of BaSyx2 (AAS metamodel/API v3). Please note that not all classes and types defined by the abstraction are depicted in Figure 9, e.g., there are also multi-language properties, reference elements or entity elements.

The `aas` component mainly consists of the instance factory (`AasFactory`) as well as interfaces defining the functionality to be provided by an AAS implementation³⁰. It is important to distinguish here between AAS interfaces (such as `Aas`, `SubModel`, `Property` and `Operation` following the AAS metamodel) and the associated (nested) builder interfaces used to create concrete instances of these interfaces in an abstract manner, i.e., so that the client code needs no knowledge about the actual underlying instance creation approach employed by the AAS implementation. The AAS interfaces provide access to the respective information and, to a certain degree, also allow for modifications of local or deployed AAS elements. Moreover, the builder interfaces allow for a concise coding style and additional consistency checks, e.g., preventing typical usage errors of the underlying AAS implementation.

Instances of the AAS interfaces can only be created through the `AasFactory` and the builders, i.e., the top-most AAS-builder can be obtained from the `AasFactory` and all subsequent builders for nested AAS elements (sub-models, element collections, properties, operations) can transitively be obtained from the actual builder. Specific extensions to the typical AAS interfaces are the deployment support (`DeploymentBuilder`), the remote protocol support (`InvocablesCreator` and `ProtocolServerBuilder`) as well as the `AasVisitor`. The `DeploymentBuilder` aims at realizing and encapsulating typical deployment recipes, such as local or remote AAS deployment. The protocol support encapsulates a specific remote communication protocol to implement the dynamic/active behavior of an AAS (as realized by the underlying AAS implementation). This (related) `InvocablesBuilder` creates function objects delegating the respective operation to the protocol/implementation, while the corresponding `ProtocolServerBuilder` registers these function objects with a matching server implementation. Ultimately, the `AasFactory` is responsible for creating a matching pair of instances for a given protocol.

In addition, the AAS abstraction encompasses an `AASVisitor`. As usual, a `Visitor` allows traversing a data structure in an extensible, polymorphic manner (based on inversion of control) without knowledge about the structure, need for explicit alternatives over types or type casting. Moreover, a visitor instance can be applied to any element in the data structure and, thus, also perform a partial traversal. Further, there is usually not a single `Visitor` implementation rather than many, each one for a specific purpose. Besides the interface, we provide the `PrintVisitor` which emits the structure of the AAS in textual form in particular for testing/debugging. Further, we provide, as usual, an empty basic implementation, the `BaseAasVisitor` to be used by visitor implementations.

In situations, where many AAS elements shall be created, deployed or manipulated in a short time, the AAS abstraction, which, by default, applies implicit caching of the AAS element instances, may be a performance or resource usage obstacle. For such situations, caching can be disabled, e.g., below submodel level, and certain operations allow for a direct non-cached interaction with the underlying AAS implementation, in particular the operations `create` (receiving a parent-level builder instance)

³⁰ We follow a pragmatic and agile approach here, i.e., we follow the metamodel, but we do not aim to be complete from the very beginning. We add interfaces and operations only on usage demand. Ultimately, at latest at the end of the IIP-Ecosphere project, the abstraction shall be complete with respect to the most recent, implemented version of the AAS specification.

and `iterate` (receiving a temporary, non-cached AAS abstraction instance of type-filtered submodel elements).

Along with the further evolution of the AAS concept, more and more standardized AAS structures will be defined. One such structure is the Technical Data Submodel [BBB+20] including manufacturer information, nameplate etc. The AAS abstraction layer takes up relevant submodel specifications (in types) and allows to create and read such structures in terms of an API based on the AAS abstraction. Since version 0.7.0, we provide here generated generic, uniform realizations based on the abstraction for all realized AAS implementations of the

- Generic Frame for Technical Data for Industrial Equipment in Manufacturing [IDTA 02003-1-2]
- Handover Documentation [IDTA 02004-1-2]
- Hierarchical Structures enabling Bills of Material [IDTA 02011-1-0]
- Draft Submodel PCF [IDTA 2023-01-24]
- Time Series Data [IDTA 02008-1-1]
- Submodel for Contact Information [IDTA 02002-1-0]
- Nameplate for Software in Manufacturing [IDTA 02007-1-0]

A concrete implementation of the AAS abstraction provides an `AASFactory` along with required (plugin) descriptors and implementations of the elements. Except for the visitors, which are based on the abstraction rather than a concrete implementation and, thus, can directly be created on purpose by client code, instances of all other concepts can be obtained directly or indirectly from the `AASFactory`. Concrete AAS factories are supposed to be realized as dependency isolating plugin announcing themselves via the `AasFactoryDescriptor`. Multiple AAS implementations may be part of a specific platform installation, one playing the role as default plugin (used for platform operations), while the others can be requested in specific situations, e.g., an AAS connector may state the specific plugin id of the underlying implementation to use.

The current default implementation of the AAS abstraction is based on Eclipse BaSyx. The `aas.basyx` plugin and similarly the `aas.basyx2` plugin implement the interfaces, typically in terms of adapter/wrapper³¹ classes, i.e., classes that delegate the actual operations to the underlying BaSyx implementation. Each of the plugins ships with its own communication protocols, e.g., BaSyx1 with the Virtual Automation Bus (in variants TCP, HTTP and HTTPS) while BaSyx2 relies on operation delegation through REST. As BaSyx ships with a large number of dependencies and not all of these dependencies may be needed on an edge device, e.g., when deploying an AAS remotely to a central server (cf. Section 3.1.2) persistent storage to a database is not needed, we aim for a dependency-reduced `aas.basyx` component and an `aas.basyx.server` component including all dependencies.

3.4.4.2 Network Management

In addition to the AAS abstraction, the support layer also provides basic network management functionality, in particular for TCP port negotiation. The network manager supports two modes, based on registered and dynamic/free ports. Both modes are relying on a self-selected key for the respective port, e.g., representing a service or a channel/topic identifier. Central services can register themselves with a platform-wide known key. Dynamic services are supported by assigning/reserving free (ephemeral) ports. Furthermore, the network management support can record the number of instances accessing a certain service represented by its known key. This is in particular important if services shall be started/stopped dependent on the actual use, i.e., if no further instance is using a service, it can be stopped and the resources can be freed.

³¹ https://en.wikipedia.org/wiki/Adapter_pattern

Network managers can be stacked, i.e., a parent network manager can contain (more) centrally registered addresses (e.g., for overarching communication brokers) while local managers focus on local (ephemeral) ports. The `NetworkManagerAas` realizes the active AAS frontend network manager instances, in particular for a central platform manager instance.

3.4.4.3 Platform Component Lifecycle

A further basic capability is to start up components in a uniform but extensible manner. This is particularly important as individual components may rely on different technology imposing different technological requirements on the startup process. Moreover, it supports the transparent realization of optional and alternative platform components. Therefore, this component defines the `LifecycleDescriptor`, allowing components to do the necessary startup/shutdown operations, declare a startup level (priority) and, if required, stop a component. A `LifecycleDescriptor` defines a priority (akin to startup levels in Linux) and may indicate, whether it desires to terminate the execution of the containing platform instance upon a certain event or condition. A `LifecycleDescriptor` announces itself through JSL and is taken up by the `LifecycleHandler`. The `LifecycleHandler` provides generic startup classes for all components, e.g., with or without the ability to terminate the platform instance, which trigger a respective processing of the lifecycle descriptors. Furthermore, to handle conflicting functionality, the `LifecycleProfile` specifies a set of `LifecycleDescriptor` instances to be executed when the profile is stated as command line parameter of the component startup. These profiles also allow for virtualization of such partial component lifecycles.

3.4.4.4 System-level Monitoring Support

System-level properties such as number of CPUs or GPUs, their actual load or temperature are particularly difficult to access in Java. Moreover, edge devices may have vendor specific interfaces including OPC UA or MQTT to access such information. To enable the generic use of such information, also in the platform AAS, we included the required basic access functionality as an interface and a rather simple default implementation into the support layer. Specific implementations can be added via JSL. One example is the `support.defaultSysMetrics` plugin, which relies on OSHI⁴² and JCuda³². Alternatively, the process information plugin interface could be used/extended.

The platform includes an optional system-level monitoring plugin for Phoenix Contact PLCnext, which accesses some system properties like CPU or board/case temperature via GRPC/protobuf provided by PLCnext (starting with firmware released in 2022). Similarly, oktoflow provides an optional system-level monitoring plugin for the Bitmotec Bitmoteco system.

3.4.4.5 Identity Support

Some mechanisms in the platform require a certain form of authentication, ranging from anonymous over username/password up to X509 tokens³³, keystores with certificates or (public) cryptographic keys as well as SSL key managers. However, storing such information in the configuration model or even in code is not acceptable. Therefore, the platform provides an `IdentityStore` with a pluggable implementation. By default and in particular for demonstration installations or testing, a YAML file with the identities is read either from the classpath, a file from the home directory of the actual process or a file determined by an environment variable. Moreover, advanced and sophisticated implementations for central identity and authentication token management can be plugged in here. Upstream components shall refer to an identity through a logical name, which provides access to the registered authentication token provided (if known) by the identity store. To allow for more

³² <http://www.jcuda.org/jcuda/JCuda.html>

³³ Originally, a generic form of identity tokens was provided by the connectors component, mainly for OPC UA. This now became a more general mechanism of the platform.

flexibility and to ease identity management, several default names, e.g., starting with a specific device name, if not found, the name of a device group, e.g., edges or servers, etc. can be used.

An example for a named YAML identity store is shown below:

```
name: HM'22
identities:
  "amqp":
    type: USERNAME
    userName: user
    tokenData: p**
    tokenEncryptionAlgorithm: UTF-8
```

The name of the store is used for logging when the identity store is loaded, i.e., which identity store is actually taken up in case that a differentiation among alternatives is needed. The `identities` are described as token objects, here the identity token with key `amqp` as a username token for user `user`, password `p**` and token “encryption” algorithm `UTF-8`. If a file entry is specified, e.g., pointing to a relative keystore, the token data is used to open the keystore and, depending on the keystore type, may then omit the user name.

3.4.4.6 Semantic Id Resolution Support

One specific ability of AAS is to mark used elements with a so-called semantic identifier, i.e., a reference to a dictionary detailing what is contained in a certain AAS element. With increasing use of semantic identifiers in the platform AAS, also a resolution of these identifiers becomes important, e.g., on the user interface to display associated value units and descriptions. Besides ECLASS³⁴ IRDI identifiers, also URL-like IRI identifiers are used, e.g., in the specifications of AAS submodel formats. A semantic id resolution mechanism must take care of such identifiers, potentially considering mechanisms implemented by the AAS framework as well as potentially commercially licensed access to catalogs and web services as they apply for ECLASS.

For this purpose, oktoflow provides a flexible semantic id resolution support. The `SemanticIdResolver` interface provides access to the resolution mechanism. The result of a successful resolution (inspired by the ECLASS dictionary) returns the version, the revision, and, in multiple languages, the name, structure name and a free text description of the referenced value unit or concept. The actual resolution shall be realized in terms of oktoflow plugins and, as fallback, through fallback catalogues provided by the platform (to be able to at least resolve the semantic ids required on the user interface even without internet access). One example plugin performs online resolution using the ECLASS web service relying on the identity management (Section 3.4.4.5) to access the required authentication certificate.

3.4.5 The support.iip-aas Component

The `iip-aas` component on top specializes the AAS abstraction for use within oktoflow³⁵, e.g., further functionality that eases the realization of the platform, e.g., mechanisms how to dynamically link alternative and optional AAS sub-models of different components into the platform AAS as illustrated in Figure 10.

³⁴ <https://eclass.eu/> we are grateful for the support of Eclass and the ability to use the Eclass catalogue in the context of a research license.

³⁵ As the initial name was IIP-Ecosphere platform, there is still the “iip” in some names.




Figure 10: Simplified structure of `support.iip-aas`: Dynamic AAS realization mechanisms, basic AAS structures relying on `support.aas`, device identity providers and fallback semantic id resolution.

One basic ability is that AAS (sub-models) for the different platform layers can be collected and deployed as a single representation depending on a given deployment mode. Therefore, the `iip-aas` component defines the `AasContributor` interface and the `AasPartRegistry`. The `AasContributor` is a (plugin) interface supposed to be implemented by upper platform layers to create the respective AAS (sub-models) and to register the implementing function objects with the protocol builders of the AAS abstraction in `support.aas`. An `AasContributor` can indicate whether prerequisites are met so that its AAS can be created. Instances of `AasContributor` are supposed to be announced/registered via JSL. The `AasPartRegistry` provides access to those plugin instances and, e.g., triggers the creation and the deployment of an entire AAS for an installation. These contributor instances are also used to increase availability (R9) and resilience of the platforms with respect to downtimes of the AAS. If the AAS server disappears, as a basic mechanism the deployment of contributed AAS and sub-models is executed again to re-populate the AAS server and to allow for continued (management) operations. The `AasPartRegistry` also maintains the AAS setup information (supported by the classes package `config` – to be renamed to `setup` in the future).

To ease the development of platform components that supply an own AAS, `support.iip-aas` defines a basic lifecycle descriptor for AAS-contributing components. Further, the `ActiveAasBase` realizes supporting utility methods for typical AAS runtime modifications, in particular for executing them asynchronously (production code) or synchronously (test code).

The next set of classes shown in Figure 10 focuses on the realization of specific AAS as well as access to the contained structures from platform code. `SubmodelClient` is a base class for all platform parts that need to access AAS submodels, for reading individual values, writing values or, in particular, for calling operations. Subclasses of `SubmodelClient` shall add specific operations and, thus, to provide a helpful code-level API to the AAS submodels. This is illustrated by the platform (nameplate) AAS

For semantic id resolution, the for oktoflow relevant fallback cataloges (based on the generic `YamlSemanticCatalog`) are realized, i.e., a simplifying small excerpt of `ECLASS` (with associated constant definitions in `Ec1ass`) and `AAS IRI` definitions from [ZVEI-N].

For using the provided capabilities in an upstream component or layer, we suggest the creation and usage/design pattern illustrated in Figure 11. As stated in the previous sections, we are using the AAS abstraction (`support.aas`) as a frontend, i.e., through the `AasFactory` without direct dependencies to the underlying AAS implementation. The platform-specific AAS library (`support.aas.iip-aas`) provides helpful classes and mechanisms that we use in this pattern.



However, to obtain a single, central AAS server and to hook the individual parts into that server with the right setup information, we need a lifecycle descriptor. A basic form, that creates also the AAS platform server instances if needed, is provided in terms of the `AbstractAasLifecycleDescriptor`, which utilizes the `AasPartRegistry` to build the AAS. To become active, the `CLifecycleDescriptor` must be specified as JSL service. In this combination, the AAS of component *C* is created at the right point in time in the life cycle (of the containing platform

componend) and automatically deployed to or registered with the platform AAS. During this creation process, also further AAS may be created, e.g., to represent a device AAS including vendor information [ZVEI-N, BBB+20].

For using the information in the AAS during the execution of other platform components, one could now request the platform AAS instance from the `AasPartRegistry` and operate on it through the abstraction interfaces provided by `support.aas`, e.g., to find a certain operation and to call it. However, if all platform parts do that directly, evolving the structure of individual sub-models becomes nearly impossible (or simply a mess). Thus, each component defining a part of the platform AAS shall also provide a (submodel) client implementation. For this purpose, `support.aas.iip-aas` provides basic client implementations, e.g., the `AbstractSubmodelClient` (for properties and operations defined on sub-model level) or the `AbstractSubmodelElementsCollectionClient` (for an element located in a submodel elements collection in a certain sub-model). The component providing the client shall now define an interface for the respective operations (`CClient`) and implement that interface in terms of either a specialized basic client, in Figure 11 shown as `CAasClient`. Upstream components that want to access the AAS, shall use the client interface and the concrete client implementation. While the `CClient` interface does not seem to be required here, it helps testing against mocked instances, e.g., in the command interface of the platform.

For modeling AAS operations, we follow the convention, that usual AAS operations behave like synchronous calls and, thus, must not be tracked via the `TaskTracker` in `support.boot`. Top-level AAS operations that shall be tracked shall be equipped with name suffix “Async”, return their task identification immediately and continue running in parallel. Lower-level operations that can be tracked are marked with the name suffix “ByTask”, offer an additional parameter “taskId” and use the task id for reporting their status. The actual distributed status reporting is realized in the transport layer.

As we started our work on oktoflow using very early implementations of AAS frameworks, e.g., when no user-defined types were available for properties or operation parameters/return types, we still simplify the modeling. For AAS properties, we usually rely on primitive types. Where possible, we avoid complex types in operation parameters and, if required, use JSON strings to transport complex or multiple values, e.g., objects, arrays or maps. Thus, to simplify later code revisions of the platform and to avoid conflicts with, e.g., annotation-based JSON libraries, we decided to provide some support for JSON marshalling in `iip-aas`, e.g., to handle return values and alternative exceptions that may occur during operation execution. Similarly, as there were no mechanisms to programmatically resolve AAS references, we decided to represent references as Strings carrying the name of an element in a submodel element collection denoted by dependencies or associations or as URLs.

3.4.7 Plugins

The support layer also implements most of the oktoflow plugins. Table 4 summarizes the core plugins, most defined/used by the support layer. Plugins can be utilized through the `PluginManager` or, in particular for testing, as classical dependency via JSL. However, for testing it is important to consider that the root classspace must remain free of dependencies as otherwise the loading of further plugins may be affected. Thus, we recommend even for testing to load plugins through the `PluginManager`, so that in most cases dependency separation through isolated classloading can be achieved and components are tested as they will be used later. While the oktoflow core is deemed to be free of direct dependencies, implementation components such as connectors may use these plugins or rely on own dependencies. In contrast, using plugins as dependencies does not lead to isolated loading and, thus, must be handled with care, i.e., cannot be applied in all situations.

Table 4: Summary of core plugins.

Plugin	Purpose	Based on	Default-Impl.	As Test-Dependency
support.log-slf4j-simple	Logging	slf4j ³⁶ including slf4j-simple	x	exclude slf4j
support.yaml-snakeyaml	YAML reading/writing	snakeyaml ³⁷	-	
support.json-jackson	JSON and XML reading/writing	FasterXML/Jackson ³⁸ , glassfish ³⁹ , jsoniter ⁴⁰	-	
support.websocket-websocket	Websocket client/server	Java-websocket ⁴¹	-	
support.processinfo-oshi	Native process information	OSHI ⁴²	-	
support.rest-spark	HTTP/REST server	spark ⁴³	-	
support.http-apache	HTTP/REST client	Apache HttpComponents ⁴⁴	-	
support.commons-apache	Common utility functions	Apache commons ⁴⁵ , jodatetime ⁴⁶	-	
support.ssh-sshd	SSH client/server	Apache Mina SSHD ⁴⁷	-	
support.metrics-micrometer	Monitoring probes	micrometer ⁴⁸	-	
support.bytecode-bytebuddy	Java bytecode manipulation	bytebuddy ⁴⁹	-	
test.amqp.qpid	AMQP broker for testing	Apache QPID	-	

One special case is the logging plugin which, for which an implementation is already required before and when the `PluginManager` is started, thus, we provide a default/fallback implementation.

In special situations, in particular for the logging and the metrics plugin, it may make sense to rely for a more consistent integration or to reuse existing setup/instances on the version provided by the respective platform component (e.g., the Spring Cloud Stream plugin for service execution). Then the dependency to the underlying implementation used in the plugin can be excluded in the component's POM and, as it is implicitly replaced by the provided dependencies of the component at hands. In such cases, the tests of the "customized" plugins shall be executed as part of the component tests to ensure future compatibility. Akin to the discussed plugin, all implementations of upstream platform

³⁶ <https://www.slf4j.org/>

³⁷ <https://github.com/snakeyaml/snakeyaml>

³⁸ <https://github.com/FasterXML/jackson>

³⁹ <https://mvnrepository.com/artifact/org.glassfish/javax.json>

⁴⁰ <https://jsoniter.com/>

⁴¹ <https://github.com/TooTallNate/Java-WebSocket>

⁴² <https://github.com/oshi/oshi>

⁴³ <https://github.com/perwendel/spark>

⁴⁴ <https://hc.apache.org/>

⁴⁵ <https://commons.apache.org/>

⁴⁶ <https://www.joda.org/joda-time/>

⁴⁷ <https://mina.apache.org/sshd-project/>

⁴⁸ <https://micrometer.io/>

⁴⁹ <https://bytebuddy.net/#/>

components have been turned into plugins for isolated loading. The platform instantiation may decide whether plugins or usual (JSL) dependencies shall be used.

oktoflow plugins can usually be relocated, i.e., paths in classpath specifications are relative. Depending on the plugin's resolution mode, the base directory for turning relative classpaths into absolute ones may be the plugin directory, i.e., the folder where the plugin was installed to, or the local maven repository (either specified in the environment variable `M2_REPO` or, as default value, the folder `.m2/repository` in the actual user's home directory).

Besides the plugins mentioned in Table 4, further upstream components, e.g., the service execution (for Spring Cloud Stream) or the configuration modeling/code generation (EASy-Producer) constitute own plugins, primarily for dependency isolation, but also to allow exchanging the respective technology if desired.

3.5 Transport and Connection Layer

The Transport and Connection Layer is responsible for connecting devices, services and resources among each other. We will discuss the two interrelated components in this layer, the Transport Component (Section 3.5.1) for the low-level platform-internal data transport and the Connectors Component (Section 3.5.2) for external data input/output.

3.5.1 Transport Component

The Transport Component is responsible for turning objects into a specified wire format and to transport the data using that wire format from a sender to a receiver, e.g., among (distributed) services. Wire format and transport protocol shall be exchangeable and extensible. The Transport Component is in particular responsible for fast (soft-realtime) communication while, in contrast, AAS is more for storing stable data of low frequency changes and for representing (distributed) operations/component interfaces. This decision was made based on early experiments [Sta20], where AAS operation calls showed a round-trip time of 23 ms and property accesses of about 4-10 ms. For comparison, plain Java Remote Method Invocations operate in this setup at 2-4 ms, which may impact the required 8 ms machine pulse in R28 if multiple sources/sinks are involved.

Please refer to older versions of this handbook for a discussion of potential data transport and data streaming technologies and how we made our decision for the technologies integrated into oktoflow.

3.5.1.1 Design

Figure 12 depicts an overview of the packages and (top-level) classes in the Transport component. The Transport component is intended to be deployable as re-usable component rather than to act as a standalone communication container. The main concepts in this layer are:

- The `TransportConnector` allowing to bind transport protocols into the infrastructure. A transport connector allows sending/receiving of data on (virtual) channels. As receiving usually happens in asynchronous manner, implementations that rely on a `TransportConnector` are informed via the `ReceptionCallback` about received data.
- To avoid creating transport connectors again and again, Transport holds a global (inter-device) and a local (intra-device) transport connector.




Figure 12: Transport Component overview (comments partially cropped)

- The actual wire format to be used for transport may differ from protocol to protocol. For example, low level transport protocols such as MQTT or AMQP support arbitrary binary payloads (might be with individual size restrictions) while higher level protocols such as OPC UA may define their own payload format. However, to be open and flexible with respect to the wire format and to utilize a minimum of data formats within the platform (R19), we foresee a mechanism for data transcoding. Specifically, for binary wire formats, the `Serializer` transcodes programming language objects into a binary representation and back. More generically, a `Serializer` is a `TypeTranslator` that can be applied also in other situations, e.g., data processing. In turn, `TypeTranslator` is a combination of `InputTypeTranslator` and `OutputTypeTranslator` with cross-over template bindings⁵⁰. Intentionally, we leave the actual technical approaches for transcoding open here (some candidates are JSON, OPC-JSON, XML or protobuf). The actual instances depend on the data types used in the application and are supposed to be generated from the configuration model. While instances of `TypeTranslator` are supposed to be attached where needed (and may be combined with `Serializer` instances), `Serializer` instances shall be usable dynamically on-demand, e.g., for a certain `TransportConnector` implementation. For this purpose, we provide a `SerializerRegistry`. Certain default type translators for primitive types are defined in `TypeTranslators`.
- The `TransportConnector` instances shall be available to other components of the platform where an internal data protocol is needed. To obtain `TransportConnector` instances, we define a `TransportFactory` (basis for transport plugins) and exhibit the actual protocol, the wire format and the broker data connector(s) from the platform configuration in the Transport AAS.
- Three default protocol plugins are shipped with the platform, namely MQTT v3 (based on Eclipse Paho), MQTT v5 (also Eclipse Paho) as well as AMQP (based on the RabbitMQ AMQP

⁵⁰ At a glance, `TypeTranslator` shall be sufficient, but in some situations, it is convenient that only the required direction must be implemented rather than both. This is in particular true for the machine/platform connectors, which require either direction for different types but usually not both directions. As `TypeTranslator` inherits from the input/output type translators, it is also possible to use a fully-fledged `TypeTranslator` in these situations.

client). Each protocol plugin is an own alternative component, the installed ones determine the `TransportFactory` behavior through a JLS descriptor. The default protocol plugins support optional Transport Level Security (TLS) and, thus, contribute to the realization of R40.

- The streaming approach currently located in the Transport Layer as transport protocols and wire formats must be provided accordingly. However, as discussed above, the streaming approach shall also remain exchangeable through glue code generation. Thus, the platform provides also transport plugins for the default streaming approach (Spring Cloud Stream), the so-called Binders, which are realized in turn through the Transport Component. A basic spring component implements convenient mechanisms for applying Spring Cloud Stream, e.g., to add serializers to the `SerializerFactory` through the component setup or to bind the `SeralizerFactory` to the data conversion mechanism of Spring Cloud Stream (`SerializerMessageConverter`). In addition, Spring Cloud Stream ships with generic serialization approaches, e.g., for JSON or XML that may be used out-of-the-box. By default, the platform ships with six alternative Spring Cloud Stream protocol binders, a generic one just using Transport, one for MQTT v3 (based on Eclipse Paho and HiveMQ-client), MQTT v5 (based on Eclipse Paho and HiveMQ-client) and AMQP (based on the RabbitMQ AMQP client). These binders support optional Transport Level Security (TLS).
- The transport component defines several global platform streams (`StreamNames`), e.g., for status (`StatusMessage`), alert (`Alert`) and trace (`TraceRecord`) messages or, as forward declarations, for upstream components, all accessible via Transport. The status notification mechanism informs interested parties when containers or services are dynamically added or removed. The notifications consist of a message data structure, which is sent on a pre-defined transport channel. Further, status notifications may report on their progress and may be task-related, i.e., carry a task id (cf. Section 3.4.2.4). As task-based execution does not have a result, status messages may carry result or failure information in this case. Alerts are created by monitoring components to signal abnormal or undesired situations. Traces make the operations of the platform visible. Moreover, the transport component defines a global instance of the default `TransportConnector` and send methods, that may queue messages until the transport connector can be utilized.

As several transport protocols rely on a central server instance, often called a Broker, it is important to mention that we do not prescribe the amount or deployment strategy for communication servers (Brokers for the mentioned concrete protocols) within a platform installation. If needed, the platform shall create a matching (test) broker implementation during platform instantiation. Further, the platform configuration provides opportunities to define multiple brokers (to be reflected in the Transport AAS) while the broker(s) to be used shall be instantiated through the platform configuration or the network managers into the respective deployment units. Moreover, based on the provided mechanisms of the protocol implementations and the streaming library, different levels of resilience or recovery can be realized, while failover to alternative broker servers may require additional implementation work.

3.5.1.2 Validation and Evaluation

We discuss now briefly the validation of the design and the implementation of the Transport Component as it has a major impact on the performance of the entire platform. We start off with a discussion of the regression testing approach and turn then to an initial performance evaluation.

The implementation of the Transport Component is subject to regression testing and continuous integration. Testing protocol integrations requires some form of server or broker instance. Therefore, further Open Source components are utilized so that the tests are self-contained, e.g., embeddable protocol brokers to simulate the platform side in the respective tests. The required dependencies are only active in testing, i.e., they are not part of a platform installation and, thus, here relaxed license or

Java version rules may apply if needed. In the regression tests, we use protobuf and a simple JSON implementation for serialization as well as Apache HiveMq or Moquette as MQTT broker and the Apache Qpid broker as AMQP broker.

For the Spring Cloud Stream binders we realized a simple setup validating the discussed streaming capabilities. This is reflected in the communication setup shown in Figure 13. Ingested by a Source (the regression test), a mocked stream component (Transformer) modifies the data (synchronously) and passes the data to the broker (representing the platform/server). The communication between these instances is handled by the Protocol Binder under test as well as the Serializer selected by the test. The Protocol Binder is based on the respective protocol implementation and in the test bound against a corresponding embedded test server/broker. To test also the flow back, a shortcut client based on a corresponding TransportConnector receives the data and ingests modified data asynchronously, which now flows through the Broker, the Serializer and the Protocol Binder back to the Source acting also as Receiver. Combining Source and Receiver is a relevant setup, as a machine/platform connector (to be discussed in Section 3.5.2) also ingests data and may receive information, e.g., to reconfigure an edge device or a machine. The regression test has access to the sent/received information and, thus, can validate the entire flow.




Figure 13: Regression testing data flow for the Transport Component.

In addition, it is also important to understand the (early) fulfillment of quality requirements. We determine the respective properties in terms of a performance experiment. Figure 14 details the setup of this experiment, which in fact is a variant of the regression test setup. Here, the Source produces a stream of data items at a certain ingestion frequency. Each data item consists of at least 50 values with repeatable characteristics (R19a). We concentrate on the payload and scope out meta-information (R79) for now. A simple Anonymizer takes a produced data item and turns one property (a name String) into simple pseudonyms. An “AI-Service” inspects the data and sends for 5 received data items one “command” back to the Source. Again, on the forward flow, the processors operate synchronously, while the backward “command” flow is ingested asynchronously. The number of received data items is recorded in all processors by simple monitoring probes and written in parallel once per second to a log file. An additional stream is used to asynchronously send experiment control commands to all involved processors, e.g., to terminate the experiment and to close the monitoring log. Items on the experiment control stream are not recorded by the probes.

The processors in Figure 14 can be executed locally (in one process, in multiple processes) or distributed on separated hosts as indicated in Figure 14. For the distributed execution, two brokers are used, one in the local realm and a remote broker in the platform realm. In the local realm, we currently use the same transport protocol/mechanism as in the platform realm, i.e., we focus at the moment on an Inter-Process Communication (IPC) setup rather than an edge setup where at least one stream goes to a different resource or the platform. Replacing the transport protocol, using different brokers or exchanging the wire format for serialization may be subject to future experiments. In this experiment we focus on the basic transport characteristics of the utilized approach.

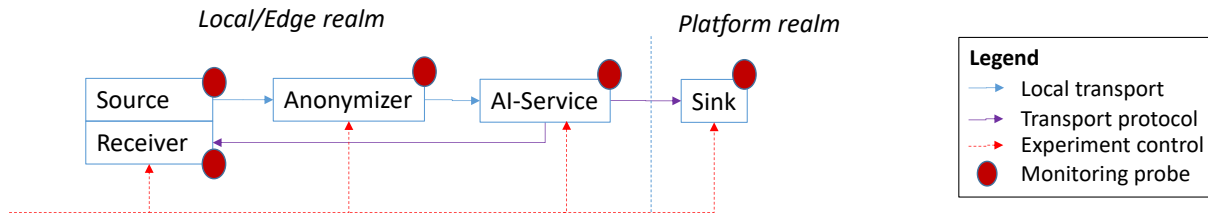


Figure 14: Performance testing data flow for the Transport Component.

For executing the experiment, we use a selection of the binders available in the platform (HiveMq v3, v5 with QoS AT_LEAST_ONCE, AMQP) with the setup as shown in Figure 14 and a respective (local, embedded) broker (Apache HiveMQ 2020.4, Apache Qpid 8.0.2). As baseline, we realized a plain network communication binder/distributed broker based on Netty⁵¹, an asynchronous networking library, and the network port management of the platform. For the source, we use a message ingestion rate⁵² per experiment and vary it from slow pace (R28) up to congestion. As wire format, we use a simple JSON serialization (leading to 650 Bytes of payload). We run the experiment for 1 minute and exclude by default the first three seconds as well as the last second where fluctuations due to network, just-in-time compilation and broker startup activities may occur. Further, some time may elapse until the average throughput is established, which we consider in this experiment as part of the stable measurements although it may significantly cause variations.

The measurements for this initial experiment have been taken on an Intel Core 7-8750U @ 1.90GHz with 32 Gbyte running Windows 10 and OpenJDK 13+33. As we aim at the moment for initial measures, we do not pay specific attention to a clean setup, e.g., getting rid of potentially other process influences such as a virus scanner or system updates.

Figure 15 illustrates the average ingestion rate at the source on the horizontal axis and the average arrival rate at the sink on the vertical axis. Until an ingestion rate of around 1000 messages per second, all binders scale similarly. Over 1000 messages per second, the behavior of the four binders differ significantly. The arrival rate of the MQTT v3 binder starts dissociating from the ingestion rate at around 1500 messages per second. For MQTT v5 this happens at around 2100 messages per second and for AMQP at a rate of roughly 2300 messages per second. While the MQTT v3 binder tries to cope with the ingestion rate until 6500 messages per second (dropping at the sink to 1400 messages per second), the MQTT v3 and the AMQP binders stop operating around 2700 messages per second. In contrast, the experimental Netty binder scales well until 7200 messages per second. Then the sink rate starts dissociating from the ingestion rate and above 9300 messages per second the simple experimental broker implementation stops operating as indicated by the trendline in Figure 15. Moreover, there are noticeable differences in settling time for the average throughput (not shown in Figure 15): All binders require more than 10 seconds to reach the respective average throughput, while Netty requires higher settling times for lower ingestion rates and AMQP leads faster to a stable throughput than both MQTT versions.

⁵¹ <https://netty.io/>

⁵² The ingestion is based on the Spring Default Poller, which is controlled by a fixed delay between message ingestion time slots (translates to a minimum ingestion rate) and a maximum number of messages ingested within a slot (determines a maximum ingestion rate). The effective ingestion rate is within the minimum and maximum ingestion, but subject to an internal congestion control of Spring Cloud Stream.




Figure 15: Average stream throughput measures for the four utilized alternative binders with trend lines.

As Figure 15 relates source and sink throughput rates, it does not reflect the total number of translated messages. Due to the streaming setup, the messages among source, processors and sink and also messages on the “command” channel (one item per five input messages) are communicated. Thus, the absolute number of transmitted messages per second is higher (least around factor 3.2). Table 5 details these numbers for the measured protocol-client-server combinations. In particular, our HiveMq readings amount to similar ranges as reported in [KGR20], where two server machines with up to 16 CPU cores but no stream processing approach were used.

Table 5: Total number of translated messages per second in best source/sink transmission situation.

Total number of translated messages per second	
MQTT v3: HiveMq, HiveMq embedded server	6172
MQTT v5: HiveMq, HiveMq embedded server	8908
AMQP: Rabbit MQ client, Qpid embedded server	9531
NETTY	30298

In summary, the required rate of 125 messages at 8 ms machine pace (R28) is supported by all brokers and works in combination with the Spring Cloud streaming approach. At around 50 values per message (650 Bytes of payload in a JSON serialization), a stable ingestion of 1000 messages per second leads to (calculated) 2.1 GByte of data transmission per hour. Moreover, the Netty binder can cope with (calculated) 15.6 GByte of data, which even qualifies for R91⁵³. It is important to emphasize that we focus here on pure IPC transport characteristics without significant data processing load. Moreover, we use a single stream, i.e., multiple (moderate) input streams from different edge devices may easily aggregate to even higher frequencies and volumes. In a realistic setting, we expect a multi-server setup as platform installation and potentially also a redundant cluster-based message handling for individual tasks, e.g., in the data integration, so that the envisioned approach qualifies for the given data (transport) quality requirements, in particular frequency and volume.

Further experiments indicate that the discussed behavior is similar when running the data processing within a single JVM, i.e., as threads, or in separate processes. Measurements on real edge devices with inter-device (cross-realm) network communications are subject to future work. As soon as further

⁵³ Based on the transferred messages in Table 5, this leads to 13.5 GBytes up to 66 GBytes per hour.

parts of the platform are available that potentially impact the data size or the performance (meta-information, security, etc.), further experiments shall be performed.

3.5.2 Connectors Component

The Connectors Component is responsible for the communication with external data sources and sinks, e.g., machines (potentially connected via some form of edge devices) or already installed platforms (the virtual platform aspect). The aim here is to allow for a bi-directional, typed communication.

Relying on the design of the Transport Component, it is desirable that the machine/platform connectors utilize type translators or serializers for the inbound communication, i.e., to translate received information (if feasible already filtered in application-specific manner) into application-specific datatypes that can further be processed in the platform. For the outbound direction, (AI-)services or humans may send instruction or configuration changes to the connected machines/platforms. These decisions are represented as information, e.g., commands, and are translated/sent through the connector to the machine or platform. Here, type translators turn the application-specific data types received from the platform side into information suitable for the external side. As stated in Section 3.5.1, application-specific type translators are realized by code generation to ease the development of applications.

We differentiate between:

- **Connector type:** Generic implementation of the platform connector interfaces for a certain protocol or information model, usually by wrapping an existing implementation (usually an external library or a framework). Connector types are implemented manually based on the connector component of the platform as an optional platform plugin.
- **Connector instance:** A connector type is created at runtime by an application after it was adapted for the specific application context, e.g., by adding type transformers for data types used in the application. These adaptations are realized through code generation, while some specific adaptations may be hooked in manually. To be considered, new connector types must be added appropriately to the configuration metamodel.

Usually, when we talk about an implementation of a connector, we implicitly refer to the respective connector type. Similarly, when we talk about the use of a connector in an application or its generation, we mean a respective connector instance.

Regarding related approaches, please refer to one of the previous versions of this handbook.

3.5.2.1 Design

For the design of this component, it is important to recall that in contrast to the Transport Component, the Connectors Component already deals with processing and translating application-specific data. For example, it is not performant to just ingest, e.g., an entire OPC UA namespace upon each data modification or, if polling/sampling shall be applied, in each poll cycle. It is more important to select the required data in an application-specific manner and to focus on the information that is required by an application running on the platform. We call the step of translating an outbound protocol into an internal protocol (and back) “protocol adaptation”, i.e., a (generated) plug-in `ProtocolAdapter` will be responsible for this task. One form of implementing the protocol adaptation is in terms of existing `TypeTranslator` and `Serializer` instances from the Transport Component, either as the realizations are part of the platform and can be re-used or because they are defined as part of the application and can be generated or are provided as hand-crafted components. However, also other forms of type translation may occur. This applies to connectors that handle generic payload (where the payload format must be translated to application-specific instances and can optionally be filtered/translated). Further, it applies to connectors that are based on a specific information model,

such as OPC UA or AAS. In the latter case, we aim for specific `TypeTranslator` instances that are linked to a generic model interface abstracting over the underlying information model. However, not all approaches support the same range of concepts and types, e.g., OPC UA allows different kinds of custom datatypes while AAS does not. Thus, connectors will differ in the offered functionality of such an interface and it may be helpful to provide meta-information stating the connector capabilities in order to dynamically guideline the code generation for a certain connector.




Figure 16: Event-based connector and push-based protocol-adaptation.

Moreover, connectors may differ in their data provisioning style. For performance reasons it is desirable to utilize **event-based ingestion**, i.e., the underlying protocol or information model informs the connector about new or changed data. Message passing approaches like MQTT or information-model based approaches like OPC UA provide such events. In this case, as illustrated in Figure 16, the “Protocol” notifies the Connector about new data. In turn, the Connector consults the `ProtocolAdapter` to translate the external data into an application-specific type, which, dependent on the “Protocol” capabilities, can be done in terms of payload translation or by querying the abstracted model of the “Protocol” (not shown in Figure 16). When the data is translated, the respective instance is passed on to a registered streaming Service in asynchronous manner. For the outbound direction (not shown in Figure 16), the Service ultimately receives the data as a stream and passes that on to the Connector upon a received data item, which then consults the `ProtocolAdapter` in the backward direction ultimately leading to a send/write command on `Protocol`.




Figure 17: Poll-based connector and subsequent protocol adaptation.

Protocol implementations not offering such notifications are subject to **polling**. One example here is the current OPC UA implementation that we use, where OPC change events just indicate a change of the model but not of which element. In its basic form, polling repeatedly obtains information from a source and ingests the information regardless whether the information changed in the meantime. This may be intended, e.g., to realize equidistant input. However, if not desired, it can also unnecessarily allocate transport resources. To avoid this, a caching mode can be defined for a connector instance. The mode indicates whether there shall be **no** caching, i.e., ingestion of all received data, ingestion only if **hash** codes are different or ingestion if the contents of the data is not **equal**. While a comparison based on hash codes is typically faster than a comparison for equality, it may also fail if hash codes accidentally are the same for equal data items (hashing collision).

As illustrated in Figure 17, the Connector then actively (based on connector settings) polls information from the “Protocol”. As before, the Connector consults the ProtocolAdapter and notifies the registered Service about the data to be ingested. The outbound direction works as discussed for event-based ingestion. Realizing the polling cycle in the Connector rather than the Service allows for connector-specific polling strategies as well as for a uniform interface towards the stream-based data processing in the platform.




Figure 18: Overview of the Connectors Component.

While event-based injection and polling may appear to be an alternative choice, a Connector may, if feasible, implement both alternatives and let the user (via the setup/platform configuration) decide about the desired approach. In particular, connectors for protocols based on information models may support both forms (such as OPC UA). Figure 18 presents an overview of the main classes in the Connectors Component of the platform. The component consists of:

- The Connector interface represents an external data source/sink. Connectors based on an information model shall exhibit a ModelAccess instance to interact with the information model rather than the payload. The Connector interface defines four template parameters, consisting of the data types accessible from the platform, i.e., CI for input into the connector and CO for output produced by the connector, and the data types for the handshake with the underlying protocol implementation, i.e., I for input into the data sink and O for output issued by the data source. A Connector can be connected to the data source/sink as specified in the ConnectorParameters and security settings like IdentityToken or certificates. When connected, received data of type O is passed through a ProtocolAdapter and an interested party is informed through a ReceptionCallback (from the Transport Component) in terms of a data object of type CO. Via the write method, data of type CI can be passed in, is translated by the ProtocolAdapter and handed as an instance of I to the underlying protocol. Finally, a Connector can be disconnected or, ultimately, disposed. All types used in CI, CO must be types used for data transfer among the services, preferably their interfaces. These types are generated by oktflow based on the application model. CI,

CO must not be any internal types used by the connector implementations. In contrast, I, O may be connector-specific types as they just represent the external/machine side and are not used further in the hosting application.

- The `TranslatingProtocolAdapter` is a default implementation of the `ProtocolAdapter` and relies on type translators, i.e., `InputTypeTranslator` and `OutputTypeTranslator` defined in the Transport Component. The `ProtocolAdapter` and its related classes will be detailed below. In particular protocol adapters to information models have a relation to a `ModelAccess` instance, which allows the type translation to interact with the underlying data model. The input type translator is responsible for turning generic values or data transport types to connector internal types, e.g., needed for implementing certain protocols and vice versa the output type translator is responsible for the opposite direction. In most cases, generic type translators for objects can be used⁵⁴.
- The `AbstractConnector` provides a basic implementation, e.g., for handling the `ReceptionCallback`, for utilizing the `ProtocolAdapter`, etc. leaving just methods open that are protocol specific. The `AbstractChannelConnector` specializes the `AbstractConnector` for channel-based protocols such as MQTT and, in turn, requires a specialized protocol adapter (as we will detail below).
- The `ConnectorRegistry` collects information about installed and used connectors. Installed connectors are registered through an instance of `ConnectorDescriptor` upon infrastructure startup (in Java through JSL) with the `ConnectorRegistry`.
- The information provided by the `ConnectorRegistry` is also the basic information to be presented in the AAS of the Connectors Component. Further, selected capabilities of the connectors are made available through the `installedConnectors` sub-model of the platform AAS. Created connector instances register themselves upon connect/disconnect with the `ConnectorRegistry`, which in turn leads to an update of the `activeConnectors` sub-model, i.e., connected connectors appear as sub-model elements and disconnected connectors are flagged as inactive. Further, connector instances provide access to their input/output data types. Ultimately, connector instances link to their descriptors in the `installedConnectors` sub-model to indicate their origin and capabilities.
- The `ConnectorFactory` is a proxy to dynamically create the most appropriate connector instance if there are alternatives, e.g., MQTT v3 and MQTT v5. Such a `ConnectorFactory` takes the `ConnectorParameters` and may decide on the supplied device service information (if present) on the actual connector type. An implementing connector may rely on existing connector implementations (as we do for MQTT).

As we will see in the next section, a connector instance can be wrapped into a service to be executed by the service management of the platform. Thereby, multiple connector instances of the same connector type (handling different types of input/output, specified in the configuration model as quadruples of CI, CO, I and O), can be wrapped into a single service so that the four template parameters of a connector type are actually no limitation for data flow modeling. If, e.g., for resource consumption reasons, underlying instances of the implementation shall be shared, a connector can create a `SharedSpace` which is passed among the connector instances in the same wrapping service.

Currently, nine specific connector types are realized in terms of individual oktoflow plugins. These are the

- generic **AasConnector** for integrating external AAS into the platform (based on the `AasFactory` from the Support Layer),

⁵⁴ In progress: Integration of direct reading/writing typed access bypassing the type translators for performance reasons.

- **OpcUaConnector** for OPC UA 1.04 (based on Eclipse Milo)
- protocol-specific **MQTT connectors**, one for MQTT v3 and one for MQTT v5, also based on Eclipse Paho akin to the Transport Component.
- **generic MQTT connector** which selects dynamically from the MQTT v3/v5 connectors based on device information.
- **serial** connector, e.g., for connecting to EAN or QR code scanners. The actual format is provided through serializers, may be based on predefined `InputParser` and `OutputFormatter` classes.
- **MODBUS/TCP** connector for connecting, e.g., to energy meters. Supports reading/writing to MODBUS/TCP devices, translates usual 1-4 byte types and allows for configuring the device id as well as the vendor-defined byte order (little/big endian).
- **REST** connector for reading from and writing to REST resources. The rest connector implementation is abstract and must be complemented with implementation-specific class representations of the data to be handled. These class representations are created by the oktflow code generation when the connector is used in an app.
- **InfluxDB** connector for writing to and streaming from Influx databases with Influx v2 authentication support via issued tokens and Influx v1 support for username/password authentication. Result streams are requested by simple timeseries or string queries (both requiring monotonic ascending timestamps), multiple entries per datapoint are joined into the data transport format of the platform (optional fields may be helpful) and ingested based on the timestamps of the data points in the database or, if given, overridden by a fixed data point delay given by the query.
- **file-based** connector for streaming from/writing to files in given formats. Multiple files can be read in sequence, a single file can be written. Files can be stored in the file system or may be app resources. The format is defined by the serializers to be attached, which may, in case of generated app integrations, be based on the `InputParser` and `OutputFormatter` classes of the Connectors component, i.e., a file-based connector could be used to stream CSV data while writing back JSON data. Akin to the InfluxDB connector, data read from files is streamed into apps either based on a) polling using a fixed data time difference or, through a `ConnectorInputHandler` or a `DataTimeDiffProvider` plugin b) triggering using arbitrary connector trigger queries.

The internal behavior of the connector types as well as the used interfaces differ depending on whether the connector is based on a typed (often hierarchical) information model or on payload transport, where usually the structure of the payload is not known to the transport protocol. This affects the role of the `ModelAccess`, the (payload) formatter/parser and, in turn, the `ProtocolAdapter`.

- Some protocol implementation libraries like OPC UA or Asset Administration Shells (AAS) are based on a **structured information model**. Accessing this model in a uniform manner is key for the uniform generation of application-specific connector code. We represent the access to the information by the `ModelAccess` interface, which allows to read/write properties (based on a hierarchical naming scheme to be interpreted in the context of the underlying protocol), to call operations, and to register (the implementation counterpart of) custom types. Typically, the generated connector code passes an initial path name to the `ModelAccess` instance and then, due to performance reasons, incrementally, indicates substructures to be accessed (`stepIn/stepOut` operations of `ModelAccess`) forming an incremental context within the underlying information model. For accessing a property or for calling an operation, the connector code passes a qualified name or, usually, a relative name within the actual context and requests the value of a property, writes the value of a property or calls an operation

defined on the model with respective parameters. The specific `ModelAccess` instance of a `Connector` can perform translations between value instances of the model and the (generated) types used in the platform/application. `ModelAccess` provides also opportunities to establish monitors on the underlying information model, i.e., to be notified on specific changes, as well as to register programming language counterparts of custom types defined in the model. For payload-based protocols such as MQTT, implementing the `ModelAccess` interface is not needed.

- For **payload-based transport protocols**, the payload is often binary, i.e., the structure of the (wire) format is not defined by the transport protocol. To handle binary input or output of channel connectors in a generic and open manner, introduce (payload) parser and formatter, for which wire-format specific implementations are provided. A parser allows generic access to a binary payload. In opposite direction, the formatter emits a given data type instance for a specific wire format. Currently, oktoflow provides formatters and parsers for JSON and separator-formatted data, e.g., by tabulators or commas as we found in some application cases. Parsed data is supposed to be mapped to types defined by the user in the respective app configuration. After parsing, data is either accessible via (hierarchical) names, i.e., in terms of nested data types, or in positional manner along a depth-first traversal of the defining data type. In the opposite direction, data is handed in depth-first traversal sequence to the data formatter, which produces the respective format in terms of binary data to be passed to a channel connector. Individual data values can be read in typed fashion from an input parser via an associated input converter. In the opposite direction, the output converter takes typed data and converts it into the output format.
- The role of the `ProtocolAdapter` and the involved type translators varies depending on whether the underlying protocol is payload-based or based on an underlying information model. For payload transport, the used communication channels may be relevant to the data parser-based type translation. This is taken into account by the `ChannelProtocolAdapter` and its default implementation, an extension of the `TranslatingProtocolAdapter`. For an information model-based protocol, the `ModelAccess` instance must be made available to the type translators as well as further initialization such as defining the polling mode must be carried out. This is enabled by two refined type translator interfaces, namely `ConnectorInputTypeTranslator` and `ConnectorOutputTypeTranslator`, both with a corresponding basic implementation.

3.5.2.2 Validation

The functional validation of the Connectors Component and the specific connector types happens through regression tests. Therefore, we follow the same basic idea as explained for the Transport Component in Section 3.5.1.2, i.e., we set up a corresponding protocol server/broker in a way that data sent to the server is (potentially after modification) echoed back to the connector. The test code produces protocol output data (of type O) either by modifying the underlying information model (event-based ingestion, polling) or by sending respective payload. The connector under test translates the data and issues an instance of type CO to a `ReceptionCallback` in the test code, which turns the information into an instance of CI and writes it back into the connector. The respective information must occur on the protocol side and can be analyzed and asserted by the test code.

Further functional tests have been performed, e.g., in the context of VDW/UMATI based on externally implemented OPC UA structures. In these validations, typically first a manually instantiated connector based on handcrafted serialization or type translation is created to test the expected intake of specific formats (e.g., OPC UA or JSON via MQTT). In addition, connectors generated via the configuration model and the platform instantiation are employed (cf. Section 6). Further, we conducted performance analyses of the serialization mechanisms and the generated connectors [EW25, NE25]. The comparison

of serialization mechanisms indicated performance peaks up to factor 10 for different JSON libraries in the context of the use case studies. The generated connectors are (nearly) as fast as the handcrafted ones, sometimes sacrificing a bit time for the generic, open approach as well as a more generated schematic model-based integration.

3.6 Services Layer

The Services Layer introduces the basis for deployable services, i.e., their interfaces, data flows, monitoring support, management and AAS representation. We separate this layer into two major components, one component to control/manage service instances and a second providing a unified execution environment for services. We start with a discussion of the terminology and background in Section 3.6.1. In two further sub-sections, we turn then to the two major components of this layer.

While the service management component is generic and can be realized in the same way for all services, service environments are typically specific for the programming language used for realizing individual services. We support this in terms of a language-specific service execution environments (due to R113, at least for Python and Java) supporting a unified integration and easing the development of services. In Section 3.6.2, we discuss the Service Execution Environments.

The Control and Management component (Section 3.6.3) is closely related to the ECS-runtime and acts as a mechanism to take control over services running on distributed devices. Control operations are, e.g., starting, stopping, reconfiguring or updating services. These operations are offered through an AAS, which also provides access to runtime monitoring information for individual services. Specific operations involve multiple services, such as switching among equivalent services or migrating services among resources, where the control and management component is responsible for the orchestration of such operations.

3.6.1 Terminology and Background

In this section, we briefly introduce our notion of the term service and discuss the bigger picture.

Several notions for services are used, ranging from web services to microservices. In oktoflow, a **service** may receive input data, transform data, or produce output data in continuous stream-based manner. Typically, source services produce data, transformer services receive data and emit modified data, and sink services receive data. Technically, a service is (a thread in) a process implemented in any programming language, while oktoflow starts/stops/manages these threads/processes, i.e., keeps them alive and runs them continuously. Input and out data are defined in terms of types and their correct composition is controlled by the configuration model (cf. Section 6). In a service, data processing can happen synchronously, i.e., an input data item is turned directly into zero or multiple output items, or asynchronously, i.e., the service receives data and produces output data at any time later if at all. A service indicates its state (R4c), meta-information (R4b), name, identification, version, kind/category (source, transformer, sink) as well as the typed input- and output relations/data paths (R4a, [SSE21]). Moreover, it allows for certain runtime operations such as passivation, migration, runtime switch to an equivalent service, reconfiguration or (re-)activation.

We distinguish between **platform-provided services** and **application-specific services**. Application-specific services are designed and implemented for a specific application at hands. The application configuration determines the input/output types as well as relevant service properties and the platform/application instantiation generates interface, integration and template code for that service. The user implements the template code and, finally, the application instantiation assembles a complete application automatically integrating the user-supplied application-specific services. In contrast, a platform-provided service (for short platform service) is shipped with the platform and, thus, shall be applicable for more than one specific application, i.e., it shall be generic and parameterizable. Often, the service implementation is even more generic, e.g., an executable tool that

shall be integrated as oktoflow service. Such services are customized into an application, i.e., the application instantiation generates a specific service stub for the application including sufficient glue code to turn the reusable generic service into an application-specific service, e.g., by using specific input/output types or by nailing down data transport from/to the underlying service implementation as well as the passing of parameters. Moreover, **hybrid services** may occur, typically generic platform services that can operate without application specific code, but which can be customized through add-ons or plugins that turn an instance of the hybrid service into an application-specific service. Thus, hybrid services can ease the realization of application-specific services and still act as platform-provided services.

Further, it is important to answer the question “**Where do services come from?**”. Details of the mechanisms will be introduced later, in particular in Sections 3.10 and 6. However, a coarse picture may already be helpful here. Services are specified in the app/platform configuration, in particular through their technical information, their meta-information and the input/output datatypes. Also, the relations among the services in terms of application-specific service meshes are defined in the app configuration. The platform instantiation/code generation either integrates platform-provided services into the realization of such a mesh or, in case of application-specific services, turns this information into service interfaces contained in application code templates to be edited by the user. Moreover, the code generation creates support artifacts such as data classes, data serializers or basic service implementations (for all relevant programming languages, e.g., Java or Python). Further, the instantiation process binds the service (interfaces) with service/glue code to the selected service execution/streaming engine. The binding happens through dynamic class loading. Dependent on the service configuration, data may be handled synchronously or asynchronously. As part of the code generation, also service descriptors required by the Service Control and Management component are created. Further, information for AAS nameplates is taken from the configuration and turned into respective implementation artifacts to build up the related AAS and submodels, e.g., one vendor AAS per service type known to the platform.

The notion of a service is cross-cutting, i.e., it occurs in many topic areas in [ESA+21] and, thus, a summary of all relevant requirements is important for the design and realization. Besides these functional requirements, we must also consider the decisions made so far, i.e., that services may offer a two-folded communication: 1) communication at lower pace for commands, status and quality properties via AAS and 2) soft-realtime communication via streams whereby the stream-integration shall be generated and flexible to allow for an exchange of the streaming approach. This is in particular important for monitoring (R4b, R4c, R4e, R4f, R133) of runtime properties and the runtime stream management, in particular to start, stop, connect (R20), update (R135), configure (R32), adapt (R69 and R31c, see also dynamic service selection in [ESA+21]) or dispose (R134c) services on demand. To be integrated in a flexible manner, monitoring and service management must be realized based on explicit interfaces and the oktoflow plugin approach, so that an exchange of the implementations becomes possible. If feasible, existing interfaces shall be utilized.

Moreover, the Service Layer must set the scenes for the management of heterogeneous service implementations (R113), including platform services that are more likely to be realized in Java or as Java interfaces to service implementations of underlying frameworks or platforms.

3.6.2 Service Environments

In oktoflow, the service environments provide implementation and execution support for services realized in different programming languages. Java services and non-Java services are integrated differently into (a Java-based stream-based) service execution engine. While Java services can be directly called, non-Java services are executed as processes and receive their control commands and data via inter-process communication/network.




Figure 19: Design of the Service Environments.

The service environments provide the basic types for service implementation. The Java service environment also contains a variety of base classes to ease the integration of native services or services realized in different programming languages such as Python. We will first detail the Java service environment, then the related Python service environment.

3.6.2.1 The Java Service Environment

Figure 19 illustrates the concepts and relations of the service environments, in particular the **Java Service environment**. From the Java point of view, the central package (environment) represents both, the basic service environment for Java and Python. This package (on the left side of Figure 19) defines the `Service` interface with all operations discussed for a `Service` in Section 3.6.2, the `ServiceState` enumeration and the main service kinds (source, transformer, sink) in terms of the `ServiceKind` enumeration. The service environment also provides support for service parameters, i.e., the customization of generic services through values that are determined upon service start or may be changed at runtime.

Several abstract classes provide basic mechanisms for realizing services, e.g., through operating system processes or using the REST protocol (on device-local networks). We just mention some examples here. `AbstractProcessService` provides the abilities to create an operating system process through a standardized naming scheme, to manage the process instance via the service status, including activation and passivation, and to customize the console input/output streams. The `AbstractProcessService` defines four template parameters, namely the received input data type (I) from the perspective of the platform/application, the input data type (SI) of the implementing process, the output data type (SO) of the implementing process and the output data type of the service (O) from the perspective of the platform/application. An `AbstractProcessService` requires two `TypeTranslator` instances, so that incoming data can be translated into a format that can be handled by the implementing process and, further, that the output of the implementing process can be re-ingested into the platform data streams. Here, we rely on the `TypeTranslator` interface from the Transport Component (cf. Section 3.5.1). Moreover, the `AbstractProcessService` requires a

`ReceptionCallback` (the interface from the Connectors Component in Section 3.5.2), so that data can be processed asynchronously.

As one path of refinements, the `AbstractStringProcessService` forms the basis for process-based services that communicate through a string format via console streams with the implementing process, e.g., JSON. This fixes two template parameters, namely `SI` and `SO` to `String`, also allowing to provide a basic realization of some inherited abstract methods, e.g., how to receive data from the implementing process. The `AbstractRestProcessService` is a further refinement which expects REST based communication represented as `String`.

Another path of refinements targets Python services, i.e., specialize the `AbstractProcessService` for executing Python and the Python Service Environment. This includes synchronous processing via command line streams (`PythonSyncProcessService`), asynchronous processing via command line streams (`PythonAsyncProcessService`) and asynchronous processing via websockets as alternative forms of integration.

Further, the `DefaultServiceImpl` is a base class that implements all methods not implemented so far empty. As “adapter” classes in Java, this class shall ease defining own services without the need of providing empty methods that are technically required but not seem to be needed by the service.

All service basis classes discussed so far, partially like connectors due to type parameters, handle exactly one input and one output type. All “Multi” classes wrap such instances in a way that the resulting service can handle arbitrary types. For services, this is the `GenericMultiTypeService`, for Connectors, which can be wrapped into a service in uniform manner, the `MultiConnectorServiceWrapper`, and for testing apps, the `MockingMultiConnectorServiceWrapper`, which injects testing data provided in resource files.

Besides the basic classes, the Java service environment offers the following:

- The `Starter` registers all services given in a YAML service descriptor and starts the application AAS command server on a given port.
- The `ServiceMapper`, a helper class that binds a service against a given AAS command server, usually the one created by the `Starter`. Further, the `ServiceMapper` registers the available metrics (see below) in the AAS command server.
- **Generic services:** One generic service that is provided by the Java Service environment is the `TraceToAasService`. This (sink) service provides/contributes to an AAS used as application endpoint. The main purpose is to provide access to trace messages explicitly emitted by the services of an application as well as to act as a uniform frame for operations that steer or reconfigure the respective application, e.g., through backward data flows. However, although potentially desirable, the service currently does not turn received data rather than optional trace messages sent by the individual services in an application. Trace messages can be enabled for debugging or for demonstration. These messages carry their origin, the action as well as an action-specific payload, e.g., the received data. The service collects all trace messages and displays them for a given time frame in its trace submodel⁵⁵ or in a channel of the websocket status server of the platform so that they can easily be taken up by the management UI. As some data tends to be rather large, e.g., the drive oscilloscope or the magnetic identification data of the EMO’23 demonstrator, the underlying `TransportConverter` can be set up to filter or modify/clean certain payload types. Further, the `TraceToAasService` is intended to act as a hybrid service, i.e., it can be used as basis to

⁵⁵ ‘Not recommended for BaSyx1, as the insertion and cleanup changes to the submodel cause memory overflows and increasing CPU load.

implement an application-specific service, which then may display processed data in an application manner or which may provide operations to be called by a device connected to the endpoint AAS. The `TraceToAasService` (as well as mocking connectors) include a generic `DataRecorder`, which is only active in application testing mode (`--Diip.test=true`). The recorder stores all received data in terms of JSON in order to ease the realization of concrete services, e.g., to pass mocked data sent out or received at a certain service as example input for data analysts or AI developers.

- **Runtime service switching (switching):** Additional mechanisms and base classes to enable a runtime switch-over between compatible, alternative services. The application configuration represents the alternative services in terms of a service family.
- **Runtime monitoring (metricsProvider):** Customized mechanisms to ease applying oktoflows metrics plugin to services based on the work of Miguel Gómez Casado [Cas21, CE21]. All information to be monitored is represented in terms of gauges, counters or timers as defined in oktoflow's monitoring plugin interface. This customization also allows to represent and query distributed meters in a uniform manner, in particular to map them into AAS (R7, R14). The `MetricsProvider` defines the unified access to predefined micrometer elements such as the system memory, but also custom meters, e.g., to measure the stream throughput. Services can be explicitly marked as `MonitoredService` to receive an instance of the metrics provider in order to define and measure application-specific metrics. For example, the `AbstractProcessService` discussed above is a `MonitoredService` to provide access to the runtime metrics of the underlying process through oktoflow's process monitoring plugin interface. Moreover, services can be marked as `UpdatingMonitoredService` if regular updates of the measurements are needed.
- **Test support (testing):** Additional support classes for testing services in their environment, in particular the `DataRecorder`.

A service realization is free to fill the service meta-information as desired, e.g., through code generation or by reading the information from a file. As the default Service Management and Control component relies on service deployment descriptors, one obvious approach is to represent the relevant information in terms of that (extended) descriptor. As these descriptors are given in YAML format, the classes `YamlArtifact`, `YamlService`, `YamlProcess` and `YamlServer` are part of the service environment to represent that information. It is important to recall that we need here only a part of the potential information in the deployment descriptor, e.g., the technical information on how to transfer network ports or how to start a Python process is not required. Further, these classes can be used as a basis to realize the parsing of the deployment descriptor of the service management and control operations in Section 3.6.2. For this purpose, parts of the (Java) service environment are imported into the Service Management and Control component and used there.

Besides the generic Java service environment, there is also a one for the current default stream processing approach (Spring Cloud Stream), the `environment.spring` in Figure 19. To provide a better integration with the logging of a Spring environment and also to reuse the already customized metering probes, the Spring service environment integrates the default logging and meter plugin implementations of oktoflow as direct dependencies and replaces the underlying libraries with the libraries that are used in the actual Spring version (also running the plugin test suites to ensure functionality and compliance). Further, the Spring service environment handles Spring-specific integration topics, e.g., the startup code in the refined `Starter` class hooks into the Spring startup process, i.e., it re-configures the Spring Rest server port and attaches that port to the `MetricsExtractorRestClient` used by the upcoming services. Further, this `Starter` version fires up the AAS command server of the parent class at a point in time when this is permissible for Spring. The `Starter` class is then executed by a specialized Spring loader (cf. Section 3.6.3).

3.6.2.2 The Python Service Environment

So far, we exclusively discussed the Java side of the service environment. Except for the generic service classes, the helper/support classes, the monitoring and the Spring-specific implementation, the **Python service environment** is a mirror of the Java service environment. Differences are:

- The Python environment is accessed through the Java representation of Python services in the streaming engine, i.e., the Python environment realizes some form of command server as well as the intra-device soft-realtime data transport, possibly with a fixed (local) wire format.
- We apply a reduced monitoring to the non-Java service environments, because stream measures can be taken on the Java side. In contrast, resource measures such as execution time or memory consumption can be combined with the related Java process, i.e., the non-Java service environment delivers the information for the own process, which is then combined to a unified measurement on the Java side.

As illustrated at the bottom of Figure 19, the Python service environment implements similar service concepts as the Java environment. However, the Python service environment does not need to be a complete mirror as only the parts to execute services and to enable the communication between the Java side and the Python service implementation are required. Therefore, the Python service environment provides the `ServiceEnvironment` class, which imports service implementations dynamically from four modules named `datatypes`, `serializers`, `interfaces` and `services`. The first three modules are generated from the configuration model and provide datatype implementations, related serializer/type translator implementations and service interfaces specified in the configuration model (implemented based on `Service/AbstractService` from the Python service environment). The `services` module contains the (manual) implementations of the services.

For the data transport between the Java side and a service environment in Python, we can imagine the following alternatives:

1. Use of the platform Transport Component using a local transport server/broker. However, as the Transport Component is open regarding the transport protocol and the serialization, this would imply that all service environments (except for the Java service environment) must implement all transport protocol variants and all serialization mechanisms. If there is not just “the Python environment” but further language-specific environments, this leads to a plethora of required protocols. In particular, if the user organization decides to implement own protocols, these protocols must be mirrored into each environment. We do not think that this is a feasible solution and opted for restricting the transport layer to Java code. Thus, the communication with Python may be based on a (local) protocol here, e.g., HTTP/REST as well as the serialization mechanism may even be fixed (we decided for JSON). While this approach is feasible as only a few variants are needed, it requires server processes for the communication on Java and Python side.
2. Extend the AAS communication protocol to transfer data (in BaSyx1 VAB, in BaSyx2 HTTP/REST). This is a specific decision of fixing protocol and serialization as mentioned in alternative 1. Also here, the backwards channel would require further server processes on the Java side as well as compliance with potentially changing BaSyx protocols. However, extending an external protocol also imposes compatibility and sustainability risks if decisions for the underlying implementation are made, that conflict with our decisions.
3. As the non-Java service implementations are executed as local processes, also command line input- and output streams provided by the operating system may be a low-risk option (as some service candidates and many Unix command line programs do). Here, in particular the Java side must carefully parse the output of the executed services/service environment not to

confuse “normal” or logging output with data output. However, command line streams are said to be a performance issue on Windows-based systems.

4. Apply a local inter-process communication approach as alternative to the command line integration suggested in the last alternative. An obvious option that also shares synergies with other platform-provided service integrations is REST⁵⁶ (representational state transfer). However, REST is designed for synchronous communication and does not allow for asynchronous service execution. While this may be adequate for service implementations that offer a REST API, it is an unnecessary limitation for the Python service environment. As a full-duplex enabled alternative that allows for synchronous and asynchronous service execution, we apply WebSockets⁵⁷ for local communication between Java and Python. Another alternative that could be integrated similarly is some form of RPC⁵⁸ (Remote Procedure Call), e.g., gRPC⁵⁹ with Protobuf.

Currently, the Python ServiceEnvironment implements both, the third (command line streams) and, as alternative, the fourth alternative (WebSockets) using generated object-to-string serializers with JSON as default transport format. We also use the command line streams for the command protocol. On the Java side, specific classes are bound against the Python service environment and the service deployment descriptor specifies the required service-specific Python artifacts as well as the Python command line parameters. More specifically, as already indicated above, PythonAsyncProcessService (command line streams) and PythonAsyncWasProcessService (WebSockets) are responsible for continuously running the Python ServiceEnvironment and the PythonSyncProcessService is an experimental call-and-return implementation of a Python service integration. While the two continuous classes communicate data and commands with the ServiceEnvironment, the PythonSyncProcessService transfers only data items and calls Python upon each data item. Besides the different communication styles, all Python integration classes support synchronous and asynchronous services and their call styles.

As Python is a system-level program, it is relevant to understand how the platform determines which Python version/installation to use. This is of particular interest, as operating systems tend to provide a more recent version, but sometimes the most recent version may not be compatible with your requirements. Moreover, certain installations may have multiple Python versions installed, so it is important to provide a mechanism, which Python to execute actually per service.

- For a (containerized) application, the platform and the platform supplied (generic) services take the information on installed dependencies into account (cf. Section 3.4.3.1).
- For direct execution of Python, e.g., in tests or in the build process, the platform (and the build processes) considers the system environment variable IIP_PYTHON as location of the python binary. If IIP_PYTHON is not set, the platform tries to utilize the first Python in the system path, as fallbacks /usr/bin/python3 or, ultimately, python.

3.6.2.3 Validation

The service environment is subject to automated regression and integration testing. In particular the monitoring classes are tested extensively [Cas21]. Also, the remaining classes of the Java/Python service environments are executed in regression tests, i.e., the Java based build environment also executes Python unit tests. However, many methods are intended to be used by a stream-based application. As done with the components before, a manual implementation of a test application and execution in particular of the Spring service environment might be helpful here, but may fall short for

⁵⁶ https://de.wikipedia.org/wiki/Representational_State_Transfer

⁵⁷ <https://de.wikipedia.org/wiki/WebSocket>

⁵⁸ https://de.wikipedia.org/wiki/Remote_Procedure_Call

⁵⁹ <https://grpc.io/>

the plain Java environment (test metrics are currently accounted per Java project rather than across projects). While currently the test coverage of the service environment could be increased, the classes defined there are tested in terms of integration tests, e.g., through test artifact for the Spring Service Management and Control implementation.

So far, no performance evaluations of the generated code and the underlying service environment have been conducted. Therefore, the manually implemented service chains from the experiments discussed in Section 3.5.2.2 could be used as baseline.

Besides service-level tests, performance experiments for BaSyx1 have been performed in [Cas21]. Retrieving a meter via an AAS on a Lenovo Z50-70 laptop requires 4-5 ms after a settling period of 200 repetitions, whereby most of the time is attributed to the AAS communication. In contrast, initial requests are comparatively slow (8-10 ms), probably an effect of JVM settling periods. Moreover, some meters can schedule own update operations, which doubles the round-trip time. In the current implementation, the `MetricsProvider` performs such updates only on request, thus, saving roughly factor 2 response time in average. The internal operation of the meters, in particular parsing the JSON information requires at maximum 70 μ s, i.e., most of the response time can be attributed to communication and AAS operations. A more extensive performance experiment is presented in [CE21], showing that an integration of the transport layer with a monitoring values cache attached to a remote AAS can be as fast as a local AAS on the monitored device. In other words, the AAS that is updated in parallel is not impacting the data paths to other components, which are informed via the transport layer (i.e., pub-sub with fixed JSON wire format). However, AAS properties realized through Java functions are only (partially) supported in BaSyx1; for BaSyx2 we currently write the metrics values directly into the AAS at higher overhead. In particular, the pub-sub approach decouples the startup of the AAS implementation server as only the platform transport broker/server is used, which is started as part of the platform.

3.6.3 Service Control and Management

The Service Control and Management component defines the service-interface of a (compute) resource towards the platform. It must provide means to load a service implementation onto the resource (in terms of service artifacts, e.g., from a central platform server), to identify the descriptive information about services (id, name, description, version, service kind) and to provide access to runtime capabilities, e.g., the state of the service, reconfiguration capabilities, or runtime monitoring values. As the execution of the services happens within their (programming-language) specific environment, the control and management component can be realized in generic manner.




Figure 20: Service states.

Individual services must comply with a lifecycle that can be queried and influenced by the platform. The underlying lifecycle state machine is depicted in Figure 20. Services (in their artifacts) can be downloaded and transitions via UNKNOWN to AVAILABLE on the hosting resource. When triggered

through the ServiceManager, a service is deployed (DEPLOYING) and gradually turns into the RUNNING state. If nothing bad happens at runtime, a service is stopped through the ServiceManager (turns to STOPPING and STOPPED) and if requested, may be started again or removed from the resource (UNDEPLOYING, afterwards UNKNOWN and potentially again AVAILABLE). At runtime, a service may be reconfigured, adapted, migrated (which may need passivation and activation) or updated (if CREATED, PASSIVATED or RUNNING). Further, a service may fail, which can lead to a recovery procedure (in the lower sub-state machine in Figure 20). If the service becomes operational again, it continues in the upper sub-state machine and there into the last “normal operation” state (via the UML H* deep history state).

3.6.4 Design and (Plugin-)Interfaces

Figure 21 illustrates the design of the Service Control and Management component (services as plugin interface for concrete implementations). At the core of this layer is the ServiceManager, which performs operations such as starting and stopping individual services. While the transitions displayed in dark green in Figure 20 are controlled by the ServiceManager, the transitions in red are performed by the service and monitored (via AAS property value polling) by the ServiceManager.




Figure 21: Service interfaces and management

Services are packaged and transported in terms of artifacts, i.e., an artifact may contain multiple services realized in different programming languages. Instead of the actual instances that may be located in a different container, the ServiceManager primarily operates on descriptors, such as the ArtifactDescriptor detailing structural information on contained services. Access to artifacts and services happens through identifiers, whereby several operations and information accesses are delegated by the service manager to the descriptors or through the respective AAS to the (device-local distributed) application AAS implementation server directly approaching the service instance.

Typically, **services can operate as a single process** and do not require further resources, e.g., a central server process to communicate with. However, there are services with this requirement, e.g., the anonymizer and pseudonymizer KODEX (cf. Section 3.8.1), which optionally may cooperate with a server instance, or a federated learning approach, which usually requires a central model management and exchange server. There are settings, where such server processes may be installed and controlled by the user, but we support also application-specific server processes, which are integrated into the lifecycle of an application, and, thus, are started and stopped along with the application by the platform. Moreover, certain applications may require some form of interaction with the server, e.g., to re-configure KODEX at runtime or to request the replay of a model snapshot of a federated learning server. It is desirable that such interactions happen in a uniform manner so that application

components can rely on and reuse the approach. For the platform, it would, thus, be desirable to represent such server processes within the platform AAS and to allow, e.g., for runtime re-configuration of service parameters (which may also trigger a model exchange). Furthermore, it would be desirable that the resource which hosts the server process can, as for services, be decided freely within the available resources of a platform instance. Moreover, a uniform, standard-compliant communication between client and server would be desirable, most preferably via the Transport Component (cf. Section 3.5.1) of the platform to also exploit the flexible exchange of protocols. However, directing an externally determined communication via the Transport Component may require changes to the server/service to be integrated, which may not be possible in all cases, for which we prioritize this requirement as optional, i.e., to be decided based on the server/service at hands.

While we believe that in most cases the open service concept of oktoflow is also applicable to services in a server/service setup, the server side requires specialized capabilities. In general, we consider a server as a realization of an (internal) service, which does not exchange data via the service input/output interfaces. This allows for utilizing AAS publishing, distribution, re-configuration and execution mechanisms of usual services, in particular the service environments for Java and Python, while not requiring an input/output modeling for the internal service/server communication in the configuration model (cf. Section 6). For this purpose, some server-specific capabilities had to be added to the service environments, e.g., that services know the (optional) service/server communication channel they are relying on. It further requires specialized capabilities such as assigning server processes to an application in the configuration model, registering and announcing the actual server network information via the platform network management (cf. Section 3.4.4.2), starting/stopping server processes along with applications (while maintaining the number of service instances using the ports in the platform network management) or provisioning of specialized transport streams for (private) service/server communication. The required capabilities are partially realized in the service manager and partially in the service environment as generic or specific functionality of basic services.

The **AAS for the Service Layer** consists of a services sub-model indicating as sub-model element collections the (locally) installed artifacts and the contained services, the installed services and their properties as well as the data paths/relations among services. When a service is started, its state changes and for each data path a relation instance is created, i.e., a relation represents the instantiated data path between two service instances and points to the actual start and end service. Start and end service occur in the AAS as soon as the respective service is created. In turn, this information is used by the service manager to determine available services, e.g., during startup of dependent services in service chains. Most operations provided by the `ServiceManager` (also via AAS) are parameterized by an artifact or service identifier. However, internally the operations are bound to the resource the respective artifact/service is installed on, so these operations do not occur at the services in the services collection rather than for the resource in the resource collection. We will detail the resources in Section 3.7.1 as part of the design of the ECS-runtime. As all those operations may fail, the implementation must not only return a result but also carry information about thrown exceptions when calling an AAS operations.

The service manager AAS is primarily intended as service-level control and monitoring interface. Services are supposed to register themselves with the respective local AAS command server (see also Figure 4 in Section 3.1) to react on command requests. Similarly, when monitoring information is requested, the (central or locally deployed) AAS communicates with the respective AAS command server. In case of services not implemented in Java, the respective service environment must provide an AAS command server and pass the information on to the service instances.

The `ServicesAasClient` provides access to the properties and operations of the AAS of the service layer. Actually, both implement the same interface called `ServiceOperations`, which defines the

basic operations of `ServiceManager` not requiring the (repeated, potentially inconsistent instantiation of) service descriptors. The `ServicesAasClient` can be used by upstream layers to conveniently access the services AAS.

Applications consisting of service meshes can conveniently be managed through **deployment plans** as opposed to manually starting and stopping individual services. One application can have multiple deployment plans, i.e., different deployments. Moreover, a deployment plan may be enabled to be started multiple times. Upon each start, a new application instance with individual data paths/streams and individual service instances is created. During startup, a deployment plan can directly re-define service parameters so that application instances may, e.g., target different devices. We utilized this capability in one of our public demonstrators, where the same application in two distinct instances with different service parameter settings was used to execute a federated learning setting between two individual cobots. Starting multiple instances and re-configuring them at startup prevents specifying the same application multiple times in the configuration, which may easily lead to inconsistencies. Similar to startup, application instances can individually be shut down using their application instance id (more conveniently on the platform management user interface where you just select the application instance to be stopped instead of recording the instance id).

As discussed above, soft-realtime data streams shall not be transmitted through AAS rather than through the streaming engine (for our default engine using one of the protocols of the transport layer). If the service implementation is done in Java, the streaming engine will directly communicate with the service (potentially involving glue code generated from the platform configuration). If non-Java service implementations are used, the service representation in the streaming engine must route the data to the respective service environment, which shields the services from the actual communication and passes the data in adequate form to the respective service instances.

The requirements in [ESA+21] do not explicitly define the properties that shall be monitored for services. R29a, R70, R122f just indicate that services may have quality properties, e.g., to support adaptive service selection. Monitoring probes may be generic or bound to the services and, thus, are realized in the service environment (in particular the default one for Java, cf. Section 3.6.2). Similarly, the creation of related parts of the AAS are realized there. Further, probe services may be inserted to perform application-specific monitoring. However, probe services are currently not realized.

3.6.5 Spring-based Service Control and Management

Different technologies can be used to realize and execute service meshes, i.e., to efficiently pass data along pre-defined data paths between the services, to transform data where needed etc. As part of such a service chain, data is turned into some form that can be transported by the utilized protocols. This serialization as well as the transformation of data to fit the input/output requirements of a service is part of the mechanisms of the Transport Layer. As discussed in Section 3.5.1, we rely on Spring Cloud Stream as default (stream-based) service execution engine. An integration of the transport level protocols and serialization mechanisms for Spring Cloud Stream was introduced in Section 3.5.1. As also stated there, we foresee that the platform shall support also other service engines in a flexible manner. Thus, the design of the Service Control and Management Component must allow for the execution of the management binding against alternative service execution technologies. For this purpose, the `ServiceManager` as well as the related descriptors are defined in Figure 21 as interfaces (in the package `services`), while the actual implementation is realized as an oktoflow plugin.

The default implementation of the `ServiceManager` in `services.spring` relies on Spring Cloud Stream, the Spring deployer mechanism and, in turn, on the Spring Boot framework. For Spring-based services, the packaging happens in terms of specifically packed JAR files (as discussed below). The Spring-based `ServiceManager` refines the descriptor classes so that the information can be loaded directly from the Spring application description. Further, it utilizes the Spring deployer mechanism, i.e.,

the local Spring deployer. The deployment specification also allows defining external service implementation processes, e.g., for Python, so that the data communication is managed by Spring-services while the actual implementation of the service operates in an own process. By default, services are executed in their own processes so that services can be restarted in case of failures (R9) without accidentally shutting down healthy services. However, such a single-process deployment may not be desired in some cases so that the deployment descriptor allows for specifying “ensemble” services, i.e., Spring-services that must be executed within the same process.

Figure 22 illustrates the structure of a generated application artifact for execution with Spring Cloud Stream. Here, an artifact consists of combined binaries provided by the service execution framework (the startup code) and binaries that make up the application (below `BOOT-INF`, requiring a specialized `ResourceResolver`, cf. Section 3.4.2.2). Within the application parts, such an artifact contains all application dependencies (in `lib`) including, e.g., the service manager of the platform and its transitive dependencies, but also generated parts such as the application interfaces (folder `iip`, also containing the startup class for the Java service environment `Starter.class`), the Spring Cloud Stream application specification (`application.yml`), the service deployment descriptor (`deployment.yml`) and the logging configuration (`logback.xml`). Depending on the integrated services, more artifacts may be included, e.g., reusable service binaries (here `kodex.zip`), customized service artifacts (`kodex_pseudonymizer.zip`) or the Python code for executing a Python service in the Python service environment (`python_kodexPythonService.zip`). Such specific binaries are referenced in the deployment descriptor and unpacked by the service manager upon start. However, this structure does not serve for isolated loading using oktoflow plugins. For this purpose, the contents of classes, libraries and the `classpath.idx` file (serving for the root classloader then) is separated into application-related files for isolated loading, named with suffix `-app`, i.e., `lib-app`, `classes-app` and `classpath-app.idx`. As the original Spring loader shown in Figure 22 is not able to cope with this separation, we realized the component `services.spring.loader`, which can handle both isolated and non-isolated app loading depending on the actual structure of the artifact replaces the Spring loader. Unfortunately, the base classes of Spring are not designed for reusability so that several classes have to be duplicated.

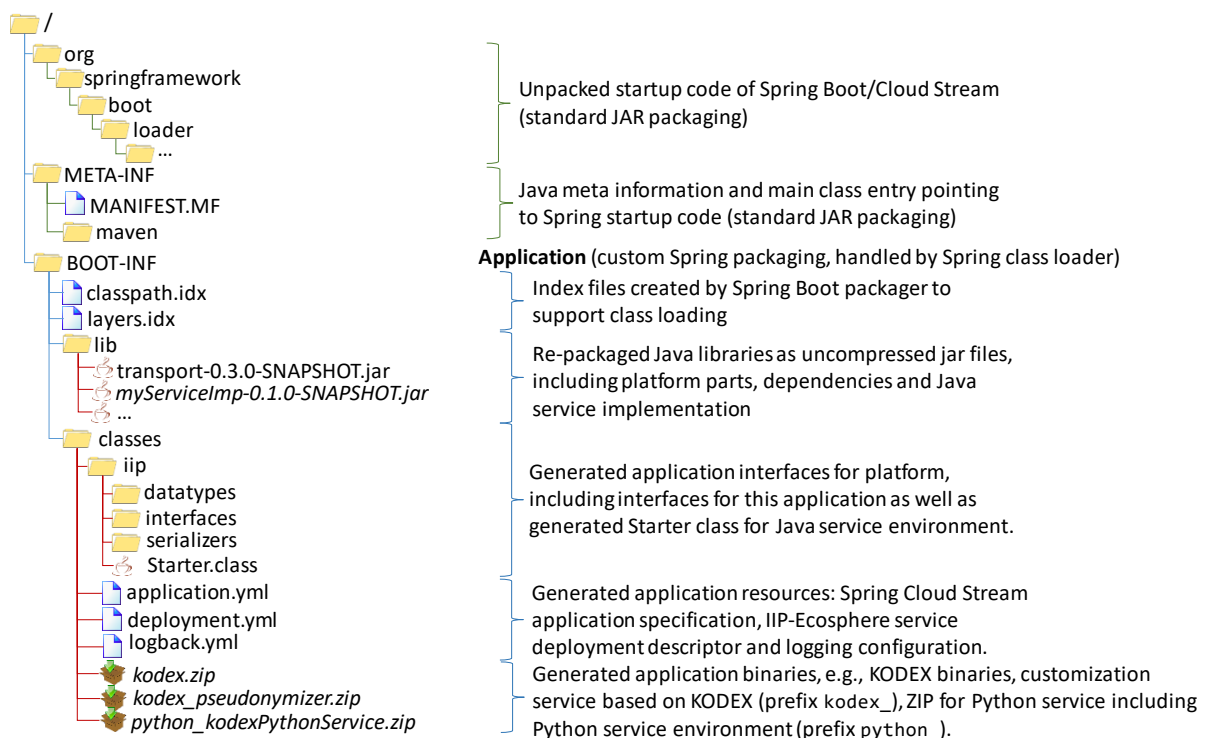


Figure 22: Structure of a JAR application artefact for the Spring Cloud Stream engine (non-isolated loading).

The packaging shown in Figure 22 represents an executable Java ARchive (JAR) and may even be executed without the platform (provided that a respective setup, e.g., communication ports are given). In principle, this is also one major functionality of the service management - besides passing environment settings such as (dynamic) ports to the services stored in the service artifact. However, as an executable JAR, the packaged Spring application prevents shared libraries, i.e., besides plugins, libraries that must not be packaged into the artifact to reduce the footprint of the artifact. To allow for shared artifacts, the service management supports a secondary format, which is independent of the Spring packaging approach (although based on `services.spring.loader`). This packaging structure is shown in Figure 23.

In contrast to Figure 22, the ZIP-based artifact is not an executable JAR and contains packaged rather than unpacked JARs. The application JAR including the generated class `iip.Starter` must be in the top-level directory of the ZIP. There must also be the service deployment descriptor (`deployment.yml`) so that the service manager can read the information about the contained services. Moreover, also the “binary” artifacts, e.g., for KODEX or python must be on top-level, so that the service manager can extract the artifacts for executing them in terms of operating system processes. In contrast, the Spring application definition (`application.yml`) and the logging configuration shall reside in the generated application artifacts as they will be loaded by the respective Java libraries on demand via class loading. The dependencies of the application are located in the `jars` folder (or optionally on top-level). Here the difference to Figure 22 is that any shared jar can easily be removed from that folder (during the packaging process or manually for experiments) and provided through a shared libraries folder⁶⁰ known to the service manager. The ZIP shall may contain in the file `classpath` a listing⁶¹ of Java libraries in their intended class loading sequence. For isolated loading, we also employ the file `classpath-app`, i.e., as for Spring classes and libraries for the root classloader are named in `classpath` and for isolated app class loading in `classpath-app`.




Figure 23: Structure of a ZIP application artefact allowing for shared libraries (as variant of Figure 22).

3.6.6 Validation

The `ServiceManager` and its AAS are validated in terms of regression tests. As the `ServiceManager` and the descriptors are interfaces/abstract classes, the validation must set up a pseudo

⁶⁰ Currently, this folder is specified in the setup information of the service manager. This information could be relocated into the service deployment descriptor in future versions of the platform.

⁶¹ Relative file names in Windows or Linux notation, separated by `:` or `;` depending on the operating system. Such a file can be created by Maven. Before execution, the file is rewritten to comply with the `@` argument file format of Java 9. Thereby, the JARs in the root folder of the archive are added by the service manager to the start of the classpath. If no such file is present, a wildcard classpath is constructed, which may cause accidental class loading conflicts.

implementation for basic testing. The Spring Cloud Specific functionality is tested through a handcrafted service artifact with simple contained services and multiple deployment descriptor targeting different artifacts, e.g., with or without process ensembles. This artifact (`test.simpleStream.spring`) is based on the (Spring) Java service environment (cf. Section 3.6.2). However, just running the tests may not be possible as in particular at the service control and execution varying, potentially conflicting dependencies may be needed, e.g., the Spring service execution together with the Spring-based BaSyx2. In other words, we must set up a dependency-free environment for the tests (`services.spring.pluginTests`), load the spring-based service execution as one of multiple plugins and, through the service manager interfaces, execute the tests that are defined in `services.spring`, which, in turn, start the processes of the test application (`test.simpleStream.spring`). In these tests, the setup of the `ServiceManager` provides a broker, dynamic network settings are handled by a local `NetworkManager` and the service manager is utilized for starting and stopping services. The running services are validated in terms of their data throughput and the actual metric values that the services provide, i.e., that the metrics defined in the service environment (cf. Section 3.6.2) become part of the AAS of the service management. Moreover, also the dynamic aspects of the AAS are validated, in particular during startup in order to figure out whether a service is already running.

Furthermore, the Spring Cloud based Service Manager was validated in a Linux VM-based server setting as well as on a Phoenix Contact AXC 3152⁶² PLC-Edge with 2 GByte RAM and 8 GByte memory card providing additional hard disk space. On the Linux server, the Service Manager was executed directly on the operating system as well as in a Docker Container, on the AXC we focused only on the Docker setup. In both cases, we were able to manage a simple demonstration application (adding the artifact, starting, stopping, removing the artifact) and to verify that the expected input/output behavior of the services can be observed. As starting and stopping individual services involves powering up a JVM, the service manager takes a certain operation timeout (with a default length of 30 seconds) into account. This is sufficient for the Linux server (and the regression tests mentioned above), but does not work on the AXC 3152, where a longer timeout is needed.

3.7 Resources and Monitoring Layer

The Resources and Monitoring layer enables the deployment of services/connectors to (edge, server, cloud) devices, allows for overarching management of the devices and provides aggregated monitoring information about running resources and services. Moreover, the platform components for the overall management of resources, namely the device management and the platform monitoring are located in this layer. We will discuss the ECS-runtime in Section 3.7.1, the device management in Section 3.7.2 and the monitoring in Section 3.7.3.

3.7.1 ECS-runtime

Flexible and heterogeneous deployment to edge, server and cloud devices is a central capability of oktoflow. [ESA+21] defines several requirements for the envisioned deployment approach. R25c and R25d target the (central) management of resources and, thus, are addressed by the device management in Section 3.7.2.

As described in [SSE21], each device shall execute a basic runtime component (`ECSRuntime`) providing the AAS of the device and managing the containers in which individual services are executed. Figure 24 illustrates the design of the ECS-runtime including the device management, which is represented by the `EcsAas`. In turn, the `EcsAas` creates the device AAS (via the `DeviceAasProvider`), the device management operations (via `DeviceManagementOperations`), the representations of the

⁶² <https://www.phoenixcontact.com/online/portal/de?uri=pxc-oc-itemdetail:pid=1069208&library=dede&tab=1>

managed/running containers via the `ContainerDescriptor` and the representations of the `ContainerOperations`. The Service Management and Control component from Section 3.6.2 contributes information to the `EcsAas`, e.g., the running services and their instantiated relations. Through the ECS-runtime, the device can receive and execute commands from the platform, such as downloading or starting a container or opening a remote access channel (via the `DeviceManagementOperations` and the `RemoteAccessServerFactory`).




Figure 24: ECS-runtime for Service Deployment

Different ways to install such an ECS-runtime are possible depending on the capabilities of the underlying device. Figure 2 in Chapter 3 already discussed the context/stack for this component, i.e., on top of the AAS support, network management, transport and connectors and (optionally) service layer, the ECS-runtime is supposed to provide a resource abstraction to manage the containers containing services to be executed on a resource.

- One approach is to provide an automatically created container with the instantiated ECS-runtime as well as one or multiple (dynamic) containers for the services.
- Depending on the capabilities of the device, e.g., whether a suitable version of Java is available, the ECS-runtime could also directly be installed on a device.
- A device could also ship with its own AAS and as part of that the required interfaces for a specifically implemented and installed ECS-runtime is standardized, which also may be available from the app store of a device vendor.

Measures to install, manage and update such installations are subject to the device management (Section 3.7.2). Moreover, different container technologies must be considered and addressed in a uniform manner through the ECS-runtime, two are indicated as oktoflow plugins in Figure 24, i.e., for Docker and LXC (then implementing specific `ContainerManagement` and `ContainerDescriptor` classes). A further alternative could be an integration with a container orchestrator like Kubernetes⁶³. Due to the general requirement R7 that all interfaces in the platform shall be based on AAS enabling an interoperable integration of heterogeneous devices, we view container technologies as a technology to be integrated and used to realized oktoflow management operations rather than the other way round. Ultimately, the platform configuration decides, which of the available technologies shall be installed.

⁶³ <https://kubernetes.io/de/>




Figure 25: Container states

Akin to the service manager, the container manager provides AAS operations, e.g., to download, start or stop a container. The respective container states are depicted in Figure 25.

Partially, device-specific functionality is supposed to be realized via other plugins, e.g., target-system specific implementation of the `SystemMetrics` (cf. Section 3.5.4) or how to provide and access to the device AAS (`DeviceAasProvider`). It may be that the device is already older and does not provide an AAS. For this purpose, the ECS-runtime allows to customize the AAS origin via the `DeviceAasProviderDescriptor`, which determines the component that returns the address of the respective device AAS (the component may also create the AAS if needed). Currently, three implementations of the related `DeviceAasProvider` are shipped with the platform, a pragmatic one reading the AAS from a simple Yaml file (reading manufacturer/product images from the same location), an implementation reading the AAS from an AASX package file via the AAS abstraction/plugin of the platform and a multi-provider that selects the first provider (by default Yaml or AASX) that returns an AAS address. For the Yaml/AASX providers, the underlying information is retrieved as classpath resource, either as `nameplate.yml/deviceId.yml` or `device.aasx/deviceId.aasx`, respectively, whereby `deviceId` is determined by `IdProvider` of the support layer. Besides nameplate information, the device AAS may also specify device-supplied services and their properties, e.g., port or protocol version, which, through an agreed key name, can be taken up, e.g., by connectors, overriding values specified in the application configurat.

As stated above, the platform provides a plain Docker⁶⁴ container manager. As for the service descriptors, the `ContainerDescriptor` of a specific container is manifested in terms of a Yaml file, which, which describes the properties of the packaged container at the same location (accessible via an URI stated in the container AAS). We refrained from adding the descriptor to the packaged container as this may not be permissible for some container formats. In addition, the platform offers an LXC container manager. We selected LXC [Sch23] due to the use in other industrial platforms [SEA+20] and as Docker changes its (commercial) licensing. It is important to mention that LXC is licensed under GPL and, thus, must be an optional component of the platform due to platform licensing rules. LXC is integrated via the Java LXD library JLXD⁶⁵ and requires specific installation steps, which are detailed in [Sch23]. In the current state, it is not supporting all container creation strategies and may be limited regarding a container registry.

The `EcsAasClient` provides access to the properties and operations of the AAS of the resources layer. The `EcsAasClient` implements the same interface `ContainerOperations` as the

⁶⁴ <https://www.docker.com/>

⁶⁵ The official repository is at <https://github.com/digitalspider/jlxd>. However, due to required bugfixes and the need for a deployment to Maven central, which was not provided by the original authors, we rely on a fork of JLXD, which is part of EASy-Producer <https://github.com/SSEHUB/EASyProducer>.

`ContainerManager`. This interface defines the basic container management operations. The `EcsAasClient` can be used by upstream layers to conveniently access the ECS-runtime AAS.

At a glance, the diagrams do not indicate much monitoring support for the ECS-runtime. As the Java service environment (cf. Section 3.6.2) provides a generic and extensible monitoring approach, we re-use it here although the ECS-runtime is not a “service”. Thus, the ECS-runtime defines an internal `MonitoringProvider` as well as regularly monitoring publication cycle, which is started as part of the JSL lifecycle descriptor of the ECS-runtime.

As for the generic platform components, regression tests validate the basic operations of the ECS-runtime, i.e., an artificial test container manager and its AAS. For the Docker-based container management, the regression tests utilize a small Open Source container image and exercise the implemented operations. Akin to services, currently advanced container operations such as update and migration are not implemented.

Experiments with containers and AAS indicate that properties and operations work as expected. Simple operations can be executed in at maximum 5 ms runtime (or significantly less for monitoring properties as discussed in Section 3.6.2). Complex operations, e.g., starting a container depend on the time that is required by underlying operation of the container implementation, e.g., Docker. Direct execution on a non-virtualized operating system was not necessarily better in this regard. However, this experience strongly depends on the AAS and protocol implementation and, thus, is not representative.

We also validated starting Docker containers via the ECS-runtime and the container manager, running the ECS-runtime directly on the underlying operating system as well as running the ECS-runtime in a Docker container. For the latter, a Docker-out-of-Docker (DooD)⁶⁶ setup is required. Moreover, to achieve a certain genericity of the ECS-runtime container, it is advisable to mount the containers via a folder of the host operating system into the ECS-runtime container. The functions of the container management were validated on a Linux virtual machine running on a VMWare ESXi server as well as on the Phoenix Contact ACX 3152 mentioned in Section 3.6. As the container operations require a certain execution time, the minimum overhead created by an AAS-based management operation is not relevant here. The service capabilities are validated in several examples, most under regression testing, as well as some public demonstrators⁶⁷. The service/server functionality is currently subject to regression testing as well as the preparation of upcoming public demonstrators.

It is important to mention, that the sizes of the Docker container depend on the platform and application services that are installed. An ECS-runtime with a DooD setup requires a container of around 1.1 GByte size (packed image of 444 MBytes), a service manager demands 509 MBytes (336 MBytes packed image) and a combined installation of ECS-runtime and service manager into one container 600 MBytes (286 MBytes packed image). All containers can be installed and executed successfully even on an ACX 3152, typically with the platform server and the central broker installed on a server, e.g., the Linux virtual machine mentioned above. The running containers in idle mode allocate roughly 200 MBytes main memory (1.4 GBytes remain free on the ACX 3152), although at least 3 JVMs (ECS-runtime, Service Manager and a local broker for the services) are running. If a simple service chain with 2 services is started, further 400 MBytes are allocated by one JVM per service, i.e., roughly 800 MBytes to 1 GByte memory remains free on the ACX 3152. Here, dependent on the actual load and service demands, we allow for some optimizations, e.g., to combine services with process backends, e.g., Python into the same JVM (ensemble services) or to limit the maximum memory allocation of the involved JVMs adequately. For the latter, the platform configuration model allows

⁶⁶ <http://tdongsi.github.io/blog/2017/04/23/docker-out-of-docker/>

⁶⁷ For details, see <https://github.com/iip-ecosphere/platform/blob/main/platform/documentation/examples/examples.md>

settings for the platform services as well as for individual application services, which are turned into executable artifacts by code generation (cf. Section 6).

3.7.2 Device/Resource Management

The device management supports the administration of devices, i.e., compute resources. As stated above, e.g. along with the ECS-runtime in Section 3.7.1, the notion of devices in oktoflow is rather broad as it involves edge, cloud and (on-premise) server devices. From a practical point of view, the scope includes all devices that potentially can run an ECS-runtime (including the IT infrastructure from [SSE21]) and/or a Service Manager. Also, different forms of installation for an ECS-runtime as discussed in Section 3.7.1 are subject to the device management. It is important to recall that following [SSE21], Industry 4.0 field devices such as machines are out of scope.

From [ESA+21] we know that the main requirements for the device management refer in particular to the "Device Description Store", the "Device Configuration Tool" and the "ECS runtime" introduced in [SSE21]. This includes the abstraction of vendor dependencies (R25.a), on/offboarding (R25a) and device management (R25b). Common management functions which are neither listed in [ESA+21] nor [SSE21], e.g., mechanisms for human interactions (acknowledgements), management techniques such as device templates or import functions for "asset data providers" [SSE21] are desirable, but also well covered by other platforms [SEA+20]. Thus, in [ESA+21, SSE21] it was intentionally left open, whether the platform just focuses on the essential capabilities mentioned in [ESA+21, SSE21] or provides also additional useful capabilities. Please note that quality requirements regarding data processing time limits, e.g., soft realtime, do not apply to management operations of the device management.

Besides this freedom, there are requirements that also prescribe the design of the device management. One important requirement is R7 which requires the use of AAS for the interfaces of all layers/components in the platform. On the one side, the device management must take the information in the platform AAS on available resources into account and use the operations provided there to manage resources, i.e., this component can require its own operations in the resources sub-model elements collections described in Section 3.7.1. On the other side, the device management shall provide relevant own additional operations (such as onboarding, selection of device templates) to upper layers such as the user interface of the platform. If adequate, these operations are parameterized with the resource identifier from the resources sub-model (cf. Section 3.7.1). The functionality of the device management is influenced by given information (through AAS events and polling, R11), but may also directly influence the resource sub-model, e.g., adding/removing devices (potentially requiring subsequent operations, e.g., shutdown/migration of containers or services).

The structure of this component follows the architectural suggestions in [Pid21]. An overview is depicted in Figure 26. The component offers two AAS interfaces, a southbound interface in `DeviceRegistryAas`, and a northbound interface in `DeviceManagementAas`. The southbound interface is intended to enable a self-registration of devices and to notify the platform that they are available (heartbeat). The northbound interface provides device information to higher-level components in the platform.

At the core is the `DeviceManagement` interface, which is composed of operation interfaces covering different aspects indicated by the requirements, such as resource configuration, remote management or firmware operations. The separation into different interfaces allows for a unified handling of implementation, AAS (client) and testing. The `DeviceManagementImpl` class unifies these interfaces (by delegation) and implements a default remote management approach via Secure Shell (SSH) based on temporary sessions created by request on trusted, registered devices. To rely on an existing, mature implementation, the communication is performed here via SSH streams rather than AAS or the transport layer (a pair of channels might be used for this purpose), although SSH may impose issues to

Windows devices. Further parts of this component are the device registry abstraction and the storage abstraction. For both parts, the implementation is left open here and can be realized by alternative oktoflow plugins. A specific AAS client class (EcsAasClient) offers access to an extension of the ECS-runtime from Section 3.8.1 to create a remote SSH endpoint on demand.




Figure 26: Device management

As indicated in Figure 26, for a specific setup, the device management offers the following alternatives:

- **ThingsBoard**⁶⁸ as central management component. ThingsBoard ships with load balancing mechanisms, an own selection of internal protocol and frequently is setup via Docker containers. It must be installed separately in addition to server components.
- **basicRegistry** as a simple, in-memory implementation of the device registry interface. Can be used instead of ThingsBoard, but does not provide a user interface or persistency mechanisms, i.e., can be used for test setups.
- **MinIO**⁶⁹ calls itself as the world's fastest object storage server. MinIO requires adequate setup on the server side. However, MinIO is licensed since April 2021 under AGPL, i.e., only an optional integration is permissible and it might not be the primary choice for installations.
- **S3Mock**⁷⁰ is a Java/Scala-based object storage server for testing S3 implementations. In contrast to MinIO, it can be used without license limitations (MIT license), provides access via the Amazon S3 interface, but accepts any authentication, i.e., can be used for test setups. In contrast to MinIO, the S3Mock integration contains a lifecycle descriptor and, if the setup includes a storage server section, starts also a local storage installation (on the central IT side).

It shall be noted that due to usage of the well-known S3 protocol/interface for the object storage, the object storage integrations can act as storage connectors to storages located in a cloud and accesses can be directed to a cloud if stated in the component setup.

The device management component supports a simple on/offboarding process, *currently without manual approval of the operations*. If explicit on/offboarding is enabled (by default, this is currently not the case to ease development), a device must be explicitly on-boarded or off-boarded. This may

⁶⁸ <https://github.com/thingsboard/thingsboard>

⁶⁹ <https://github.com/minio/minio>

⁷⁰ <https://github.com/adobe/S3Mock>

lead to the exchange/removal of security certificates or encryption keys. On devices, that were not on-boarded, the platform may not execute operations. *Neither exchange of security information nor denial of operations are currently implemented.*

The functionality of the device management has been validated through many fine-grained test cases, see also [Pid21]. There, the performance of the direct execution of individual device management operations using the ThingsBoard device registry and the MinIO S3 connector have been measured and take in average 8-170 ms. If the operations are executed via the device management AAS sub-models, the BaSyx1 operations take in average 11-204 ms.

3.7.3 Monitoring

Service execution shall be monitored, in terms of resources and functionality, e.g., through application specific probes and alerts. Therefore, the platform employs a set of generic built-in monitoring probes in platform (cf. Section 3.6) and application services and allows for application-specific probes. While probing of the individual services or ECS-runtimes/resources happens on the devices, the main task of this component is to aggregate the data on IT level (see also [SSE21]). The aggregation of the received values shall follow existing guidelines, approaches or relevant standards in Industry 4.0.

The basic requirements for the monitoring component in [ESA+21] focus on devices/resources, services and alarming/alerts, in generic or application-specific fashion (e.g., through specific monitoring services hooked into the data processing chain). Akin for the device management, one important general requirement is R7 which requests the use of AAS for the interfaces of all layers/components in the platform. On the one side, the monitoring must take the information in the platform AAS on available resources into account and use the information provided by services and resources through their local monitoring. On the other side, the device management shall provide relevant (aggregated) information and own operations to upper layers such as the user interface of the platform. The functionality of the monitoring component shall rely on the combined information. However, due to potential performance issues of the AAS approach, for urgent alarms/alerts the path via the Transport Component is more adequate. As usual, the monitoring component defines an own AAS submodel, which currently consists of a list of recent alerts.

As first (alternative) monitoring component we decided for an integration of the Prometheus⁷¹ service and resource monitoring approach (open source, Apache License). Prometheus is based on gathering data from exposed HTTP/REST endpoints, allows for configuring evaluation rules on the gathered data, stores the data in a time series data base and exposes the aggregated information again as HTTP/REST service endpoints (including an alert manager). However, direct HTTP access across all resources in a production system may not be permitted, i.e., some intermediary representation might be required.

Instead of scraping AAS (for performance reasons), we rely on the approach illustrated in Figure 27, i.e., to provide the monitoring information through the transport layer as envisioned in [CE21]. For Prometheus, this is similar to monitor resources over network borders, where a firewall or a gateway provides a proxying service⁷², which collects all relevant information in the subnet on behalf of Prometheus and offers the collected information on individual endpoints for scraping. Although this requires an additional server process, it also allows for the flexible integration of other monitoring systems, as the data is provided independently and just must be translated into a format that can be understood by the respective monitoring system, e.g., a transport-to-MQTT translation that feeds information into the monitoring component.

⁷¹ <https://prometheus.io/>

⁷² <https://github.com/pambrose/prometheus-proxy>

In our case, the integration of Prometheus into oktoflow receives the monitoring data of individual resources via transport communication, exposes this information in an own (local) web server and adjusts the Prometheus configuration so that new devices are considered for scraping via the (local) web server. The implementation of such a metrics exporter is prepared in the generic monitoring component, while the Prometheus-specific integration is done in the alternative Prometheus plugin. The Prometheus plugin also contains the Prometheus binaries for Windows and Linux as well as an own lifecycle descriptor that starts and shuts down Prometheus. Within this lifecycle, a bridging metrics exporter as well as an alert monitor are started. We use the Prometheus alertmonitor⁷³ (Apache License) as the Prometheus client library does not provide support for alerts. The alertmonitor scrapes the alert manager HTTP API of Prometheus in regular fashion and turns alerts into alert instances of the Transport. The setup of Prometheus is defined in the configuration model and the setup information is generated during the platform instantiation process. Part of this setup is also the information, whether we rely on an installed Prometheus server or whether the platform manages the execution/lifecycle of the included binaries.

The monitoring of system-level meters and application-level meters (data items received/sent) has been validated through the Prometheus UI. Added resources (ECS-Runtime/Service Manager) when they occur are taken up, system- and application-level meters are categorized according to device id (and service id for application-level) and displayed individually. Aggregated values or rates can be calculated from this information on Prometheus level. Currently, the underlying approach based on micrometer automatically adds several technical system-level meters. Moreover, Spring also adds additional meters. Most of these meters are not relevant on platform level and could be filtered out.




Figure 27: Monitoring

3.8 Storage, Security and Data Protection Layer

The Storage, Security and Data Protection Layer is responsible for managing security aspects of the platform based on the platform configuration, for offering security-enhancing services (such as anonymization or pseudonymization), but also for secure integration of (encrypted) data lakes or clouds. As discussed in Section 3.1, the purpose of this layer is not to realize typical cross-cutting security mechanisms. We do not focus on the configuration aspects (R40a, R40b, R41a, R42, R44, R64a, R65a) here, as we do so later in the discussion of the Configuration Layer in Section 3.10.

⁷³ <https://github.com/matjaz99/alertmonitor>

3.8.1 KODEX platform service

The privacy enhancing service (plugin) in this layer integrates the KODEX privacy and security engineering toolkit⁷⁴ by KIPROTECT into the platform. It is important to emphasize that KODEX (realized in GO) is a generic tool that requires some form of setup to operate on the incoming data in the intended manner. To cope with the genericity of KODEX, some design decisions were made for KODEX that apply analogously (as a blueprint) to other external services:

- The `KodexService` is parameterized over the incoming and outgoing data types. To transfer data instances correctly to KODEX, respective type translators are required. These type translators shall be provided by the utilizing code, which knows the incoming and outgoing types of all service meshes of all applications on a certain platform instance.
- The customization of KODEX happens in terms of certain files that specify the data model. Akin to the type translators, the contents of these files are determined upon integration into a service mesh and are generated from the configuration model. These files are packaged into a ZIP archive (named according to the using node in the service mesh), stored in the service implementation artifact and specified in the respective process part of the service deployment descriptor. When starting the service node, the deployment descriptor is consulted, the artifacts are extracted and the customization files are placed into the temporary home directory of the process implementing the service, here KODEX.
- In the (extracted, temporary) home directory of the process, also the service implementation must be located, i.e., in the KODEX case the operating-system specific binaries. Such implementation files shall be packaged into a “binary” Maven ZIP artifact and deployed along with the service integration code, here `KodexService`. When integrating the generic service code into a service mesh, the Maven identification of the service implementation, here KODEX, is known, and so is the deployed implementation (here binaries for different operating systems). During automated instantiation/integration, the “binary” ZIP is packaged into the service implementation artifact and named respectively so that it can be extracted upon service start along with the customization files as described above.
- Upon code generation of the Spring nodes, further customizations may happen, e.g., service-specific customization files could be created.

Besides an integration via command-line streams, oktoflow contains a form that utilizes the REST API of KODEX. Initial performance results confirm that REST outperforms command line streams on Windows. For example, processing a batch of 1000 tuples, windows with command line streams takes 15 ms per tuple in average, REST on Windows 0.22 ms, command line streams on Linux 1.4 ms and REST 2 ms on Linux.

Regarding licenses, it is important to mention that at the point in time of writing this document, KODEX is licensed under AGPL. However, viral AGPL rules do apply to binary code, i.e., using the KODEX binaries with respective credits does not taint the license limitations of oktoflow. Moreover, KODEX is only integrated, if it is explicitly selected and used as a service in an application and only becomes active when the respective service mesh is started, i.e., on decision of the user.

3.8.2 Influx DB connectors

Although technically belonging to the connectors layer (Section 3.5.2), the Influx DB connectors (for v1/2 and v3) are logically part of the data storage and protection layer. Basically, these connectors are intended to write data points transported as part of oktoflow application streams to an Influx DB. Therefore, incoming data tuples are split into individual Influx field entries of the same measurement id and data point, while the timestamp can either be determined from the data (specialized

⁷⁴ <https://heykodex.com/>, <https://github.com/kiprotect/kodex>

LongIndexType) or by the client when writing the data. Receiving a data stream from an Influx DB requires a connector trigger query, in this case a simple timeseries (start/end time without further filter conditions) or string queries (for v1/v2 in flux query format, for v3 either in SQL or InfluxDB format as specified). While the simple timeseries query is by default turned into a stream of monotonic ascending timestamps, this must be explicitly ensured for the more open string-based queries. When receiving the individual data points, the connector joins them back to data tuples and uses the generated data transformers to ingest corresponding application-specific types (optional oktoflow fields may be helpful if the data in Influx is partially incomplete). The periods between two successive datapoints is obeyed as far as possible, if not overridden by a fixed delay per tuple given by the trigger query. The connectors support Influx v2/v3 authentication via issued tokens as well as Influx v1 support for username/password authentication.

The connector is subject to regression testing based on mocking the Influx client (the usual approach) and further was manually/in internal examples tested with InfluxDb2 version 2.7.6.

3.9 Reusable Intelligent Services Layer

On top of the layers discussed before, the Reusable Intelligent Services Layer provides AI mechanisms in reusable and configurable manner and integrates received/monitored data with additional information such as product order information or floor plans to provide further valuable input to the AI. In this section, we briefly discuss the specific requirements (Section 0), the integration of RapidMiner RTSA as AI platform service (Section 3.9.2), further service candidates ahead (Section 3.9.3) as well as the deferred AI service concept.




Figure 28: Reusable Intelligent Services and data processing function library

3.9.1 Data Processing Function Library

Application of AI methods encompass more than just the AI methods. Usually, also functionalities for data pre-processing etc. are required. As a basis to meet the requirements regarding pre-processing and data transformation functions (in [SSE21], e.g., frequency analysis), we equipped the platform with a library of functions. In this version of the platform, the functions below are

- Image transcoding from/to base64 strings.
- Image processing such as grayscaling, rescaling or thresholding.
- Barcode/QR-code detection based on the Java library zxing⁷⁵ and, as optional fallback, the Python library pyzbar⁷⁶. For the Python fallback, respective packages must be installed.

⁷⁵ <https://zxing.org/w/decode.jspx>

⁷⁶ <https://pypi.org/project/pyzbar/>

- Generic channel-based time-series aggregation. The need for this functionality arised during the realization of the EMO'23 demonstrator where a condition monitoring AI shall be fed with channeled data from a linear drive and, from a different sensor, related energy measurements. The energy measurements were delivered as individual telegrams representing a time-series, while the condition monitoring AI shall be fed with summarized data over different measurement channels for the time span the drive is operating. As two kinds of data, individual data points for all energy channels per point in time and one data point for all energy channels per point in time were available, we realized two generic, re-usable aggregators, which were then instantiated in the respective aggregator service of the application.

3.9.2 RapidMiner RTSA service

Altair/RapidMiner offers pioneering company solutions in data analysis and AI. Their RapidMiner Platform (in Java) and the RapidMiner Studio shows that AI composable from building blocks is not only a vision. One fundamental component in the RapidMiner ecosystem is the Real Time Scoring Agent (RTSA), a REST-based execution environment for deployments created by RapidMiner Studio.

Following the ideas in [SSE21], a separation of data science exploration and design processes from the actual execution/deployment is desirable. RapidMiner is an example for such an approach integrated into oktoflow. While the DataAnalyst can first create a data science process for given data in his/her own environment, the created process (a “deployment”) can later be used and executed by oktoflow under the control of the Service Manager. For the integration, mainly the data input/output formats must match, i.e., the data provided by the oktoflow app (output of a connector/service) becomes the input for the RTSA deployment and, in turn, its output the input for upstream oktoflow app services.

Along these lines, the platform-supplied `RtsaRestService` (see Figure 28) integrates RTSA and links the platform data streams to the input/output of RTSA. The realization is similar to the REST-based integration of KODEX discussed in Section 3.8.1. Through oktoflow's app configuration, a deployment file for the RTSA is specified, which is generated and packaged by the application/platform instantiation along with the RTSA binary into the respective service artifact. Here it must be considered that RTSA is a commercial service, i.e., it requires a license file and cannot be distributed openly. For testing the service integration and the platform instantiation the test part of the component ships with an RTSA mockup (`FakeRtsa`), which acts as a REST server pretending to be an RTSA instance with a deployment. This fake RTSA can be configured in limited form to transform the input data, e.g., by changing fields or adding fields having a random number value. For the platform instantiation, either the real or the mocking fake RTSA and its deployment are given in a specific folder according to the platform naming for binary files from which the instantiation process takes up the binaries.

3.9.3 Flower-based Federated Learning

One AI method that is of particular interest for the industrial production and for the integration within the platform is federated learning. Federated learning clients act as services as they consume data and produce predictions, but they also share information about their AI model with other federated learning services (usually of the same application) so that the other services can learn new knowledge faster, e.g., acceptable anomalies.

oktoflow integrates the Flower⁷⁷ framework written in Python. Services (federated learning client) and the server (usually assigned to one or multiple applications) are specified in the configuration model. Both, client services and server service are generated into the Python part of the application code templates based on the configured settings (including the basic technology code for, e.g., tensorflow and numpy), must be completed during service realization and are executed through the Python

⁷⁷ <https://flower.dev/>

Service Environment. As explained in Section 3.6.3, in contrast to the client services, the server does not consume/produce regular service data streams. To be manageable within the platform, the server needs a Java counterpart which is also a Server, which manages the (hidden) service application lifecycle and relies in most of its functionality directly on `PythonAsyncProcessService`. Akin to the service and server code, respective test code for the Flower services is generated into the Python test part of the application code templates.

3.10 Configuration Layer

It is important to recall that in oktoflow all configuration relevant information is reflected in terms of IVML structures, relations and constraints, while the IVML reasoner validates the platform/application configuration before and at runtime by identifying problems and deviations from validation rules. The Configuration Layer provides functionality to define applications in terms of the platform IVML configuration on top of the (reusable) services.

Figure 29 illustrates the design of the **configuration component**. While the diagram (and the implementation) may appear rather trivial, most of the complexity is in the configuration model, the instantiation process and the underlying framework EASy-Producer [SE15]. The foundation is the `configuration.interface`, which defines the plugin interfaces for configuration technologies. `configuration.easy` is the oktoflow configuration plugin integration the EASy-Producer technology. Finally, `configuration` is the actual configuration component collecting the metamodels the configuration plugins (multiple may exist) into an integratable Maven model artifact. `configuration` also builds up the configuration AAS, which serves as the main backend for the management UI. In addition, `configuration.maven` implements a Maven build plugin that executes the platform/application instantiation via configuration as well as other plugins for testing oktoflow apps as well as for building/testing Angular apps.

The configuration metamodel in `configuration.easy` follows the layered architecture of the platform, i.e., each platform layer is represented by a configuration module. Figure 29 just indicates the technical integration of that model into the platform. We will discuss the model in more details in Section 6. For each platform to be installed, a dedicated platform configuration is created which specifies the AAS settings, the data types, the employed services and the applications etc. The platform instantiation process is defined based on the metamodel, i.e., a configuration is used as input defining the platform and the applications shall be instantiated. The platform instantiation process turns the configured information into source code artifacts, setup information, deployment descriptors and executable build scripts. This process also significantly contributes to the invisible complexity of this component. We will discuss also the instantiation process in more details in Section 6.




Figure 29: Configuration and instantiation of Definition of applications and orchestration of services

configuration.easy contains sub-packages for turning OPC UA XML into IVML (opcua), for turning IDTA AAS into IVML (aas) as well as for reading IVML structures and manipulating IVML (ivml), in particular for graphs such as the serviceMesh.

On top of the models and the instantiation process, the Configuration component just orchestrates the relevant processes. The ConfigurationSetup (read from a Yaml setup file) defines the file system paths where the metamodel, its instance and the instantiation process are defined (metamodel and instantiation process are part of the respective release). The ConfigurationManager ensures the consistency of the operations, e.g. loading, validating and instantiating the model.

The configuration model, the actual settings and in particular the application data flows are reflected in the ConfigurationAas. To ease UI integration, the ConfigurationAas offers via AasIvmMapper operations to read out the data flows in various formats, to write back/delete data flows, to create and delete configuration variables, to change configuration variables and to initiate the application instantiation. For these operations, it is essential that no arbitrary modifications are permitted. The result must always be a valid configuration that can be instantiated. For this purpose, all operations perform a validation through the IVML reasoner and persist the configuration model only if the model is valid. Further, the PlatformInstantiator realizes a command line tool to perform the basic operations of the ConfigurationManager, i.e., to allow a user to instantiate the platform and the defined applications. The PlatformInstantiator offers various modes, from instantiation of interfaces for applications, full instantiation of applications, instantiation of platform components, etc. This tool is integrated in configuration.maven into Maven.

The configuration metamodel is extensible, i.e., it consists of a core model as well as extensions for devices, services and connectors. Already mentioned extensions integrate, e.g., KODEX (Section 3.8.1), Flower (Section 3.9.3) or RapidMiner (Section 3.9.2). One further, particular extension represents an integration of the MIP technology magnetic identification sensor, one of the IIP-Ecosphere dynamic demonstrators. The extension specializes the platform-supplied MQTT connector through specific data types for the MIP sensor technology and, thus, eases the use of this technology in user-defined applications. Besides a connector, also some code to control the sensor is helpful for application projects. This code is supplied by the **configuration default library**, a platform component based on data types defined by the MIP extension of the configuration model. Thus, the configuration default library showcases how to supply code and related dependencies with the configuration model and provides a blueprint for extension libraries. For short, the build process of the default library just instantiates the interfaces of a minimal platform configuration without any IoT application, i.e., only defined types are generated. This code be used in apps in terms of Maven artifacts. It is important to mention that such libraries intentionally ship without the generated types, which are then generated as part of the respective IoT build process.

The configuration metamodel and the platform instantiation are subject to regression testing in the continuous integration. This encompasses simple testing applications with a varying number of interconnected connectors/services as well as configurations for platform and container instantiation. Moreover, the platform examples to be detailed on github⁶⁷, are used as application-level regression tests. Most of those applications are not only built rather than executed and validated against expected output through the build plugins provided by configuration.maven. To demonstrate the setup of the platform, the platform instantiation as well as the creation of example service artifacts is part of the Docker platform containers provided on DockerHub.

3.11 Application Layer

Ultimately, the Application Layer represents individual applications, i.e., it is the actual home of the application generated from the configuration, i.e., the generated artifacts and additional application-specific (handcrafted) services. Currently, this layer exists only virtually, e.g., in terms of the example applications discussed on [github](#)⁶⁷.

3.12 Platform Server(s)

Besides ECS-runtime, service manager and monitoring, there must be one platform service that is responsible of maintaining the overall platform AAS, the application and container artifacts and, if required, also the lifecycle of the AAS server instances (registries, repositories). Moreover, this “platform” component also provides a simple command line interface to operate with the platform if there is no user interface, e.g., to start containers or services. This component serves for:

1. Powering up the servers to run the platform. Therefore, the component defines a lifecycle descriptor (`PlatformLifecycleDescriptor`), which reads information from the `PlatformSetup` representing the YAML setup file. The lifecycle descriptor is loaded via JSL into the `LifecycleHandler`, which, in turn, is called by the `platform` component during its main program. During this startup process, all “installed” lifecycle descriptors (e.g., the descriptor for the network manager; the platform instantiation is responsible for this) are also started up. As part of the startup also the platform AAS is constructed, which contains the platform “nameplate” (`TechnicalInformation` sub-model), further software-specific information (`Platform` sub-model) as well as a listing of all available application `Artifacts` (service artifacts, containers, deployment plans).
2. Observing a heartbeat of the distributed components via their regular monitoring messages, in particular ECS-runtimes and service managers, to detect whether a component instance dropped out and whether the respective AAS submodel still exists and shall be removed.
3. Providing a simple command line interface (CLI) to execute the operations of the platform, from rather low level to high-level commands such as executing an app deployment plan. The command line interface does not rely on the lifecycle mechanism, but on the `PlatformSetup` and, in particular, on the AAS clients of the service and the resources layer to ease executing the operations defined there. Figure 30 illustrates an example interaction with the interactive mode of the command line interface, here turning into the resources commands, showing the commands for resources (`help`), listing the available resources and, finally, ending the CLI. For the single resource shown in Figure 30, in particular the integrated container manager (for Docker) and various initial runtime measurements for disk and memory allocation are shown. It is important to emphasize that the command line performs its operations via the platform AAS and the respective AAS clients for services and the ECS-runtime.

oktoflow interactive platform command line

AAS server: `http://127.0.0.1:9001`

AAS registry: `http://127.0.0.1:9002/registry`

Type “help” for help.

> `resources`

`resources> help`

`list`

`help`

`back`

`resources> list`

- Resource `a005056C00008`

`systemdisktotal: 1023887356`

`systemmemorytotal: 2147483647`

`simplemeterlist: ["system.cpu.count", "system.cpu.usage",`

```

    "system.disk.free", "system.memory.free"...]
containerSystemName: Docker
systemmemoryfree: 2147483647
systemdiskfree: 464061712
systemmemoryused: 2147483647
systemdiskusable: 464061712
systemmemoryusage: 0.5555296172875698
systemdiskused: 559825644
resources> back
> exit

```

Figure 30: Interaction with the preliminary interactive platform command line interface.

Using the platform CLI, we validated the interaction among the components. Therefore, we started platform, ECS-runtime and service manager component as individual programs. Through the CLI, we validated the resource represented by the ECS-runtime and started a simple generated application (cf. Section 6). We also validated the execution of services in a service manager container, starting and stopping of containers via the platform and the ECS-runtime execution in terms of a (Docker-out-of-Docker) container. Please refer to the technical guidelines on github on how to install, instantiate and containerize the platform, i.e., to perform the steps that we also executed for validating the command line interface and the instantiated platform components. The platform CLI also supports creating snapshots of the platform AAS that can be explored with the AASX Package Explorer⁷⁸.

3.13 Platform Management User Interface

The user interface focuses on information display, app management operations well as on editing the configuration model and designing apps through the Platform AAS in terms of an Angular WebApp. The management user interface is separated into three main parts / views, the available and on-boarded resources known through their AAS, the platform and application configuration as well as runtime operations such as starting or stopping applications. The configuration part is structured according to the logical flow of setting up a platform and specifying applications.

Figure 31 illustrates the overview of the available resources allowing for more details from the respective AAS on demand. Each on-boarded resource is displayed with information from its device AAS, e.g., the vendor and the product image. Pressing the button “resource details” leads to an overview of technical information, a combination of the static device AI and the runtime information contributed by the platform monitoring.




Figure 31: Management user interface, available resources.

The next screenshots illustrate the configuration part. The sub-items of the configuration menu indicate the sequence of steps to define an application: Constants, types, software dependencies, AAS nameplates, servers, services, meshes, applications. Most of these steps are displayed as lists with a

⁷⁸ <https://www.plattform-i40.de/IP/Redaktion/DE/Newsletter/2019/Ausgabe21/2019-21-Praxisbeispiel2.html>
required version depends on configured version of AAS/BaSyx used in the platform.

dialog-based editor for creating or editing respective elements. Figure 32 shows the configuration of AAS nameplates to be used in the configuration of individual services. Figure 33 depicts the graphical viewer/editor for service meshes and Figure 34 the configuration of applications, here with actions to obtain application code templates and to integrate service implementations to an application.




Figure 32: Management user interface, configuration of AAS nameplates for service vendors.




Figure 33: Management user interface, configuration of service meshes.




Figure 34: Management user interface, application configuration allowing to obtain implementation templates and integrating implemented templates.

In the last step illustrated in Figure 34, the user can generate application code templates supporting the realization of services. When pressing the “generate template” or the “integrate application” button, a specialized form of the instantiation is executed. If successful, the result of “generate template” is a download offer for interface and template artifacts.

- Without specific Maven repository, we assume that the user is developing on the same computer as the platform is running. Then the downloaded interface artifact shall be unpacked and installed (`mvn install`), the implementation template shall be imported into Eclipse (or VSCode, with less Python path support than Eclipse) and utilized as discussed in the Tutorial Videos. Finally, the service implementation can be uploaded and integrated using the “integrate application” button.
- If a maven repository is specified in the configuration (e.g., SCP or FTP upload to a maven repository, a Sonatype Nexus⁷⁹ or a JFrog Artifactory⁸⁰ server), all created application-specific build processes will use that for deployment so that except for the download of the implementation template, no further up/download is needed.

In the runtime part, the user can access the available deployment plans (Figure 35), upload new deployment plans or start enabled plans. Starting a plan multiple times may (if not prevented by the plan) lead to the execution of multiple instances of the same application. In the instances view, all running application instances are listed and can be undeployed individually. Figure 36 illustrates the currently running services with interactive access to the service logs (Figure 37).




Figure 35: Management user interface, overview of deployment plans and actions to start application instances.

⁷⁹ <https://help.sonatype.com/repomanager3/product-information/download>

⁸⁰ <https://jfrog.com/artifactory>




Figure 36: Management user interface, overview of running services actions to access the runtime service logs.




Figure 37: Management user interface, active runtime service logs.

The management UI is able to configure most aspects of an application including the service meshes, i.e., the data flow graphs. Each modification requires a validation of the underlying configuration model so that consistency issues or problems can be identified before instantiating the results.

The user interface requires some setup, in particular knowledge about the installation location of the platform AAS servers. The platform instantiation turns the Management UI into a compiled, instantiated version, where in particular the settings in the Angular environment/setup JSON are adjusted or respective start scripts, e.g., for an Express webserver are generated.

As the management UI is based on the platform AAS, which is typically running on a different network port and may even run on a different server, access may need to be granted due to Cross-Origin Resource Sharing (CORS)⁸¹. By default, CORS is disabled in the configuration metamodel, but it is

⁸¹ https://de.wikipedia.org/wiki/Cross-Origin_Resource_Sharing

enabled for all accesses in the install package. CORS can be set up through the configuration variable `aasAccessControlAllowOrigin`, e.g., by setting the value to `"*"` (the default, typically in `TechnicalSettings.ivml`). If CORS is not explicitly enabled, usually a browser plugin is required.

Figure 38 depicts the structure of the implementing TypeScript classes, mainly in terms of Angular components (UI elements of different granularity with code, display and style) and Angular Services (re-usable functionality). From top-to-bottom and left-to-right, the implementation consists of the top-level Application module and the Application component that bootstrap the functionality and the dependencies. The FlowchartComponent and its subordinate feedback component integrate drawflow as flowchart editor and customize it for oktoflow service meshes. The ServicesComponent displays running services and via its associated LogsComponent displays runtime logs of individual services in own windows. The DeploymentPlansComponent displays the available deployment plans and allows executing them. Subordinate components allow for file uploads, status updates as the execution of a deployment may take some minutes and a subordinate StatusBoxDetailsComponent to display a sequence of related status messages. The ListComponent is rather generic and responsible for all configuration lists, e.g., for constants, datatypes, nameplates, servers, services, or applications.




Figure 38: Management user interface, implementing components and services.

The most complex component is the EditorComponent, which represents an editor for complex configuration types that are selectable in the aforementioned configuration list. It interprets the configuration AAS to dynamically create editors for all (in particular complex) configuration types, including mandatory and initially hidden optional settings. The EditorComponent further delegates to individual editor components representing individual configuration settings within a complex type, in particular for Boolean values, enumerations, AAS/OPC-UA language strings, references among configuration elements (e.g., applications link to meshes, services to servers) and representations of nested complex values (e.g., datatypes declaring service in/outputs), which, in turn, use a (filtering) instance of the EditorComponent. Similarly, when editing a reference to another configuration element, the EditorComponent for the reference target is opened.

Furthermore, the InstancesComponent displays the list of running application instances and allows terminating individual instances. Finally, the ResourcesComponent displays the overview of the

available resources, delegating a detailed view on a single component to the `ResourceDetailsComponent`.

The briefly explained components utilize sharable services for realization. As Angular services are made available through dependency injection, in principle every service can be used in any other service or component. However, in practice also the services form a usage hierarchy, for which we display in Figure 38 selected dependencies, e.g., `DeploymentPlansComponent` and `InstanceComponent` rely on the `PlanDeployerService` for executing/stopping individual deployment plans, the `ResourceDetailsComponent` in particular on the `SemanticResolutionService` for resolving AAS semantic IDs as well as on the `TechnicalDataRetrieverService` for the actual technical data of an individual resource. Moreover, the `EditorComponent` utilizes the `IvmlFormatterService` which turns user input into IVML values to be passed to operations of the Configuration AAS.

In turn, these services depend on even more basic services, e.g., the `StatusCollectionService` receiving status updates from the Service Environment or the `ApiService` which realizes AAS communication primitives. The `ApiService` is based on the environment configuration service (`EnvConfigService`), which reads a customizable JSON file written during the platform instantiation. Finally, the `UtilsService` provides helpful static and instance utility functions.

Since version 0.7.0, we are equipping the Management UI with regression tests (more than 100 tests). Although we might mock the platform services, in particular the AAS, we decided to test the Management UI against a real platform instance for better consistency. Therefore, the Maven plugins from `configuration.maven` are used to compile Angular as well as to test an application in its lifecycle, i.e., to instantiate a simple application including platform, start the platform and an `ECS-Runtime-ServiceManager`, if requested the simple application, run the Angular regression tests and shut down this temporary platform instance.

3.14 Test support

So far, we focused on the elements to construct the platform, the services as well as the applications. In this section, we provide a cross-cutting overview on the testing support, in particular to answer the question, how the platform supports the user in testing his/her own services and applications.

Besides internal component and service testing, in particular of services supplied with the platform, it is essential to test user-developed services as well as their interactions in an application. At a glance, the answer might be to construct a unit test and to test the supplied code. However, reality is not so trivial, as, e.g., connectors may be based on external devices or their server instances, e.g., MQTT broker or OPC UA server, and these devices may not be available in certain testing situations. Moreover, setting up a test for a single spring-based service involving a Java or, in the more complex case, a Python service (involving the Java integration and the Python service environment) requires much internal knowledge.

However, **testing as early as possible is essential** to avoid time-consuming detours through app instantiation, packaging and execution. Thus, as a general advice, we recommend to apply testing to all levels of an application, ranging from tests of the code that you supply up to entire applications.

Moreover, we recommend to base the implementation of individual services or applications on the generated application code templates. The code templates provide the structure of implementation projects where only explicitly marked parts must be filled with own code, e.g., glue code to own classes or AI methods. Such templates also contain a customized build process for Java and, if demanded by the application configuration, Python. However, it is important to keep in mind that the **code generation never touches hand-written code**, i.e., code templates are re-generated with each run of

the application instantiation while integration of differences due to changes in the application design is currently a manual developer task.

In summary, the platform offers different forms of testing support:

- **Testing the connectivity of an individual connector:** To validate that a certain connector correctly communicates with its counterpart, e.g., a device or a server, the platform generates per connector a simple connectivity test program that uses the configured connection information, creates a connector instance and requests data. If implemented by the connector, it also emits the data structure provided by the connector, e.g., an OPC UA tree.
- **Testing individual services:** For a user-defined service typically self-supplied code is entered into a generated service template. This combination shall be tested individually. However, writing tests against internals of an infrastructure like the platform is not trivial. Thus, the application code templates contain template classes for production as well as for testing for both, Java and Python. The generated build process of an application template performs compilation (Java), syntax testing (Python) and execution of test cases (Java, Python). The generated Python service tests read a JSON data file, feed the contained data points through the generated serializers and data classes into the service and assert the output.
- **Testing services in the service environment:** Testing a service in its real execution environment requires mocking the service execution environment, e.g., unpacking Python and AI models in the right folder, executing the service environment and service lifecycle, e.g., through Spring Cloud Stream. The setup of these tests is, unfortunately, not trivial. For this purpose, the generated app templates contain test cases that set up an appropriate environment. A JSON file⁸² (see generated templates in the app template) determines the input data, which may define a repeated/timed ingestion behavior. This data is fed through the DataWrapper into the service. Resulting data, synchronous or asynchronous, is received by the generated test and emitted. In the generated version, all data received from the service under testing is asserted as true as we currently do not support data correctness specification in the application configuration. Thus, the test only ensures that after feeding a service, it also emits data. Statistics about received data types are collected by default and can be used for asserts. As the generated test must be taken over manually into your service implementation project, you may also extend the test for more realistic assertions.
- **Testing an individual connector as a service:** A connector is a kind of platform-supported service. Usually, the generic connector (type) implementation is tested sufficiently. The instantiation process wraps the connector into a service and adds generated input/output data translators, user-supplied data translators or event handlers based on the application configuration to turn the generic connector type into an application-specific connector. Although the user-supplied parts shall be tested through individual unit tests, it remains unclear whether these parts will work correctly when interacting with the connector. However, this can be tested as discussed above as the generated service (instance) including user-supplied parts is a service that can be tested in its service environment. Akin to services, the application instantiation generates code to test individually connectors in their service environment. As for service testing, data injection happens through JSON files
- **Testing/mocking applications:** Although components may be working after applying the test opportunities discussed above, there is no guarantee that the integrated application will be working. Here, again, mocking may be required as not all devices or even software environments are available. The basis for this form of test is that applications can be executed

⁸² For legacy reasons, the file name endings in some examples or generated implementation templates may look like YAML files. The mocking implementation of a connector actually reads “.yaml” and “.json” files.

even without the service execution environment provided by an installed platform. For this purpose, the generated build process of an application code template contains specific commands in the build process. These commands allow for starting an application through a mockup of the service environment, which, in turn, starts the services in an arbitrary order on the same machine (in contrast to distributed execution and controlled order as provided by the service manager of the platform). Please notice that longer-running tasks must be executed in the start method of the service rather than in its constructors (as detailed in Section 4); otherwise, data injection may already be completed while some services may still not be ready for processing. Connectors can be mocked as described above, e.g., if certain devices or their OPC UA servers are not available. Services can also be mocked, by replacing the service classes defined in the model, e.g., by extensions of the services that disable environment-specific functionality. For example, if no GPUs are at hand, one may replace a TensorFlow-based Python AI script by a simple mock script (requiring a respective modification of the application configuration).

- **Testing the application:** Despite all tests, ultimately also the application with all services and all required devices in place must be tested. This can be achieved by running the application standalone in mocking mode or in a real/simulated environment with the respective devices.
- **Testing the entire platform and distributed apps:** Apps can only finally be validated or evaluated, e.g., for resource consumption, response time or throughput, when running in distributed fashion. As the whole cycle of configuring, installing, deploying starting, evaluating, stopping, and uninstalling applications as well as platform components is not trivial, oktoflow contains the **Platform Evaluation and Testing Environment (PETE)**⁸³ that allows for automated, distributed, monitored, and repetitive tests and evaluations. PETE is based on the concept of an experimentation workbench [SEK21], there implemented in terms of Jupyter Notebook Project⁸⁴, which in case of PETE were used to develop and realize the individual execution steps. After converting Jupyter Notebook to a Python script, PETE can be executed in headless manner. PETE covers the 13 steps shown in Figure 39, while measurement/validation still need to be realized:




Figure 39: The steps executed automatically by PETE

It is important to mention that only testing on application-level may include all resources and service implementations in the final form as it will be deployed. Thus, accidental overlaps of resources, e.g., identity stores may only be detected when running an integrated application.

⁸³ <https://github.com/iip-ecosphere/platform/blob/main/platform/tests/test.environment/README.md>

⁸⁴ <https://jupyter.org/>

The requirements documents [SSE21, ESA+21] even demand in-place pre-deployment tests. Currently, the platform does not offer functionality for these optional (but important) requirements.

4 Architectural Decisions and Constraints

Often, an architecture explicitly or implicitly defines constraints that must be obeyed by an implementation. We summarize and explain the constraints for oktoflow below:

- C1. Higher layers and contained components are allowed to have **dependencies** only to downstream layers and components, if possible, only to the directly adjacent lower layer. This constraint is induced by the basic layered architecture style of the platform.
- C2. As an exception from C1, the **ECS-runtime shall not depend on the Services Layer** so that the services layer can be installed separately (as explained in Section 3). Both, Services Layer and ECS-runtime may depend on certain classes of the services environment.
- C3. **Asset Administration Shells** are the mechanism for the **communication among platform components**, e.g., remote method calls. Each component that may originate from a different vendor shall be equipped with an own AAS linked into the platform AAS [EN23]. AAS shall form a unified backend for the platform and the management UI. Exceptions are monitoring / tracing streams [CE21], soft-realtime transport streams or operation status updates.
- C4. Wrapped **singleton components** or libraries, usually in oktoflow plugins, shall not be called by other components than the wrapper. This applies e.g., to transport and connector protocols, the AAS implementation, the container management. One exception is `support.aas.basyx.server`, which is allowed to access (as the only component) `support.aas.basyx` as it represents the server component with full dependencies.
- C5. There shall be **no direct references into optional or alternative components** except for refining optional/alternative components. Moreover, generic components shall not reference their specialized components. For access to the specialized implementation, plugin techniques such as descriptors, factories, facades and JSL shall be used. In contrast to production code, this constraint may be relaxed for test code (Maven scope “test”), which may declare dependencies to specific alternatives/plugins to enable functional testing, e.g., to explicitly rely on the AMPQ transport protocol.
- C6. **Protocol brokers/servers** are internal and shall not be used for communication with external components. On the one side, these broker/server instances host internal data traffic, i.e., additional channels (accidentally) using the same names may disturb platform operations or overload transport capacities. On the other side, the internal protocol/broker/server may change due to some configuration decision and, thus, may break (accidental) assumptions.
- C7. **Protocol brokers/servers for testing** such as Apache Qpid, HiveMq or Moquette shall be in testing plugins and no other component shall directly use classes from them. These testing servers may be used during platform instantiation to provide a default broker/server for a configured transport protocol.
- C8. **Generated artifacts** shall be separated from manual code (usually an own folder such as `gen`) and generated artifacts shall not be modified as they may/will be re-generated.
- C9. **Implementation of services** shall be separated per service, so that services can be composed/integrated free of other dependencies. For convenience, in test code, we may intentionally invalidate this rule, e.g., `test.configuration.configuration` implements all service artifacts for all tests in `configuration.easy`.
- C10. **Exception handling** demands basic rules on how and where to use/handle exceptions enforcing certain responsibilities. Exceptions indicate abnormal situations in the program execution that shall not be handled by normal program code rather than by stopping the execution at the point of occurrence and tracing back the calls until the exception at hands is handled (or on top-level it terminates the actual thread). While often programmers try to handle an exception at the point where it obviously occurs, we believe that in most cases the caller, i.e., the cause of executing the code that throws the exception shall be informed. This

shall not imply that each exception shall be handled in top-level code. For example, consider some complex data format processing code, e.g., reading an AASX file for an asset administration shell. If we handle an Input-Output exception within the AASX component, the caller does not know that and why the format processing fails. Let us now assume, that reading the AASX file was triggered by the ECS-runtime when building the AAS of the ECS-runtime, e.g., to link device vendor and ECS AAS. Here, the lifecycle handler of the ECS-runtime (more or less top-level code) that starts the creation of the ECS AAS is not interested why an AASX file processing fails. However, the code creating the AAS trying to establish the AAS link is better suited to handle the exception, e.g., to insert an empty link or to log the problem. Thereby, logging (cf. Sections 2) is often not the only answer to an exception, in particular not emitting an exception stack trace to the console (which may not be logged properly). In contrast, the programmer shall think about handling the exception in a manner that processing can continue, e.g., inserting an empty link into the AAS rather than no AAS property at all, which may cause failures in parts relying on the assumption that such a property exists. In particular the type of the used exceptions shall be selected carefully; in Java code, we typically employ only `ExecutionException`, `IOException` and `IllegalArgumentException`, wrapping all more specific exceptions into those.

- C11. Apply **defensive logging**, i.e., carefully think about what is an “error”, a “warning”, an “information”. Errors shall only be emitted if a component fails to operate. If the component can compensate this, e.g., by falling back to some strategy or default plugin, then a “warning” is more adequate. Often, “information” is helpful to better understand failing processing sequences, while too much “information” is also not helpful.
- C12. **Logging setup/filtering is decided during integration, not before.** Basically, logging must happen through oktoflows logging interface and may be handled by an implementing logging plugin (or, initially, oktoflow’s default logger). This allows that “bigger” plugins like BaSyx or Spring set up and integrate their logging in a special way with the platform logging. In the extreme case, the platform instantiation decides about how to set up logging, actually a deferred logging decision. Concrete logging decisions and implementation may only occur in test code/dependencies, preferably using the platform logging implementation plugin.
- C13. Basic libraries like **YAML and JSON shall not be used directly**, only via the platform support interfaces and, thus, indirectly through the implementing plugins, akin to logging.
- C14. **Interfaces, signatures and behaviors**, in particular of services and connectors, must comply with the design of the respective components so that the code generation can properly take them up. In many cases, arbitrary changes or customizations cannot be easily realized as they must conform with all existing components and may require changes to the code generation. E.g., constructor signatures shall declare the same parameters in the same sequence as the (abstract) superclasses. Likewise, generic classes must comply with the given template parameters. However, in individual cases, template parameters of connectors may be extended (new parameters added to the end) or an abstract connector class can be instantiated by generating additional, required methods. However, typically constructor calls cannot easily be modified, if at all, by adding parameters, for which the values must be obtainable from the platform configuration. Akin, the behavior of new connectors and services must comply with the expected behavior of implanting or overriding methods as documented in code.
- C15. **Dependencies** of all kind including Java, Python or artifacts like AI models or resources are managed through Maven and, thus, must either be available through official or own/local/private Maven repositories or packaged during the platform/application build processes into Maven artifacts, e.g., using Maven assembly descriptors. Application code templates include generated Maven build processes with respective assembly descriptors,

e.g., for Python services or specified (AI) artifacts. Due to the plugin concept and the isolated class loading introduced in version 0.8.0, managed Maven dependencies are not allowed anymore as they pollute the effective dependencies (thus, the basic parent POM `platformDependencies` does not contain any dependencies). Thus, platform core layer components like `support` or `transport` are not allowed to declare any external dependency. All implementing components/plugins (e.g., `support.aas.basyx` or `transport.amqp`) may declare dependencies locally, but are encouraged to rely on the managed dependencies of the platform's bill-of-material POM (`platformDependenciesBOM`). In particular, `platformDependenciesBOM` aims for commonly used versions of dependencies as well as for reducing the installation footprint.

- C16. **Spring dependencies** are considered optional so that implementation components can rely independently on individual spring versions (although this may require packaging such components as platform plugin, i.e., loading it into an own classloader). The common spring dependencies for platform implementation components/plugins are in `platformDependenciesSpring`, an extension of `platformDependenciesBOM`. However, plugins may rely on more specific/recent versions of Spring, e.g., BaSyx2, and, thus are allowed to declare own dependencies or to override managed dependencies.
- C17. The **platform core components must be free of dependencies** except for Java and downstream platform components (C1). For functionality that require libraries, either use those provided by oktoflow (C13) or create own platform with respective interfaces in carefully selected layers of the platform.

It would be desirable to check and enforce these dependencies. However, so far tools that we tried, e.g., in the continuous integration, failed for multiple components using a central or even adequately distributed rule set as they require an application rather than a component to be checked.

5 Asset Administration Shells

oktoflow heavily relies on asset administration shells (AAS) to describe the capabilities and interfaces of its components and to interface with external components, in particular using standard structures for AAS/sub-models exist and where we have yet adopted them. Often, during initial development of components, there were no such standards so that we pragmatically defined own submodels aiming for an adoption of standardized formats as soon as possible. Until version 0.8.0 we integrated

- Generic Frame for Technical Data for Industrial Equipment in Manufacturing [IDTA 02003-1-2]
- Handover Documentation [IDTA 02004-1-2]
- Hierarchical Structures enabling Bills of Material [IDTA 02011-1-0]
- Product Carbon Footprint [IDTA 2023-01-24]
- Time Series Data [IDTA 02008-1-1]
- Submodel for Contact Information [IDTA 02002-1-0]
- Nameplate for Software in Manufacturing [IDTA 02007-1-0]

through automated generation of IVML per submodel format and subsequent API code generation based on oktoflows AAS abstraction/plugins [EW24]. Through this approach, oktoflow supports now 26 IDTA submodel specifications and one specification draft. It is important to recall that upgrading crosscutting structures, e.g., the device representation, requires synchronization among several components including the management UI and, thus, may be subject to future releases. With this pragmatic approach in mind, we designed and partially realized the platform AAS structure shown in Figure 40. As already explained in Section 3.1.2, we separate between AAS describing an (external) artifact and internal information (usually in sub-models). We follow the main rule that AAS are created for each type of component that may originate from a different vendor [EN23]. Thus, AAS do exist for

- The platform with its various sub-models like equipment/software nameplate, dynamic network port assignment, transport setup, available artifacts such as apps, containers or deployment plans, installed connector/service types and their utilized data types, the device management and available devices (with their containers, service artifacts, running services).
- Further assets represented in their own AAS like devices, service or composed applications (with vendor information). Device and service AAS are linked from the respective platform submodels to make the information in the AAS available. For each application running on top of the platform, an AAS shall be provided (currently via the `TraceToAasService` discussed in Section 3.6.2.1), which states the creator of the App, lists the utilized services and may provide application specific operations.

For the platform AAS and its sub-models, we distinguish between installed/available descriptors and their active instances in apps, in particular as in many cases only the active instances provide the full information about in/outgoing types. Examples are in particular the connectors, the services and their relations, the containers, the apps etc. These structures are dynamic, i.e., they change due to installed components as well as due to started/terminated instances. Some sub-models are active, in particular those providing operations, e.g., for network or service management. A specific example is the dynamic relation between resources and services. When an ECS-runtime starts, it contributes its AAS to the resources collection of the platform AAS. When a service manager starts, it contributes further operations to the resource it is running on, i.e., both components contribute their specific information/operations to the same, unified AAS sub-model.




Figure 40: oktoflow AAS structure

Figure 41 depicts a screenshot of an excerpt of (an earlier rel of) the platform AAS using AAS metamodel v2 in the AASX Package Explorer, i.e., an excerpt of the full AAS shown in Figure 40.




Figure 41: oktoflow AAS in the AASX Package Explorer showing a running service (SimpleReceiver).

6 Platform Configuration

This section provides an overview on oktoflow's configuration model and the employed concepts as well as insights into a simple example configuration. Section 6.1 dives deeper into the configuration model in terms of modeling patterns, Section 6.2 provides more details on the structure of the configuration metamodel, Section 6.3 treats special configuration topics and Section 6.4.6.3 discusses additional support for standardized protocols or connectors such as OPC UA. Please note that a technical documentation of the configuration elements is now located in github. Next, Section 6.5 details the instantiation process and Section 6.6 the container instantiation. More technical information, e.g., steps how to implement an oktoflow application, project structures and service implementation rules are summarized in the platform documentation on github.

In essence, the configuration model mirrors the component hierarchy of the platform and describes per component the configurable elements, their dependencies and constraints. For modeling, we use the default configuration plugin for EASy-Producer [SE15] and, thus, for modeling, the Integrated Variability Modeling Language [IVML]. The configuration model consists of three parts:

1. The **configuration metamodel** defining the configurable elements, their structure, relations, properties and consistency constraints.
2. A **platform configuration** uses the configuration metamodel to specify the configuration of a certain platform installation. This encompasses the selection of alternative platform components, e.g., the transport protocol, AAS meta model, or container manager, as well as the apps to run on that platform installation including data types, services, and service meshes. A platform configuration may define application/installation specific constraints.
3. A **valid platform configuration** complies with the configuration metamodel and fulfills all constraints. Notable exceptions are the **application templates** introduced in ReGaP, i.e., incomplete IVML application skeletons that can be turned easily into complete applications. These templates are subject to (full) validation when turned into application. A valid platform configuration can be instantiated using a process description delivered with the configuration metamodel consisting of (VIL, variability implementation language) and artifact instantiation templates (VTL, variability template language) [VIL]. The instantiation process assembles the components of the specified platform and generates implementation artifacts (code, tests, data files, build specifications, app code templates) and packages them adequately.
4. VIL and VTL can be used at **runtime to adapt the underlying system** [Eic16].

The configuration model is managed by the configuration component (Section 3.10) at runtime and used during build time for platform instantiation. The configuration submodel represents the configuration metamodel and the actual platform configuration as an interface for the management UI. Through specific AAS operations, the platform configuration can be modified while ensuring that a modification does not invalidate the consistency constraints imposed by the configuration model.

Figure 42 illustrates the structure of the configuration metamodel, mirroring the layers and components of the platform, each given in terms of an IVML project/module. The most basic project (MetaConcepts) introduces even more abstract, generic concepts for adaptive, configurable software systems. These concepts are incrementally refined into platform specific concepts. The first specific model defines the concepts for DataTypes (the types used in services and service meshes), in particular PrimitiveType and RecordType consisting of Field instances. Primitive types are already frozen⁸⁵ on this level.

⁸⁵ Frozen elements cannot be modified outside the defining IVML project. Only frozen elements can be instantiated before runtime, while the remaining elements may be frozen later or remain changeable at




Figure 42: Illustrative structure of the platform configuration metamodel (simplified).

On service level, several refinements of the platform's service term are defined as configurable elements. `ServiceBase` is abstract and contains information common to all services, e.g., name, id, version, description, input types, output types, service kind or operation mode (synchronous/asynchronous). `ServiceBase` already defines constraints prescribing which information must be present for which kind of service. A `Service` is a refinement of `ServiceBase`

runtime. MetaConcepts defines so called binding times to conditionally freeze configurations as well as the CReversibleProperty, which partially remains unfrozen.

and also the parent of specific services like `JavaService` (e.g., detailed by a Java qualified class name denoting the implementation) or `PythonService`. A special kind of Service is the Connector, one type per implemented connector type (only OPC-UA is shown here, similar elements exist for AAS, MQTTv3, file, Influx etc.). A `ServiceFamily` represents multiple, alternative but equivalent services with the same input/output types. Service families steer the selection of alternative services, usually at startup or runtime. As `ServiceFamily` it inherits from `ServiceBase`, it can transparently be used as a `Service`. From the configured services, the code generation derives implementation interfaces (Java, Python) and service stubs (Java) for the integration of non-Java service implementations.

The `Devices` module defines the properties of the ECS-runtime, in particular the container manager to use. Moreover, it defines the `EcsDevice`, which represents an installed/connected device. In the next release we plan that `EcsDevice` instances steer the automated creation of Docker containers as well as the automated and optimized assignment of containers to resources.

The `Applications` module allows specifying one or multiple applications consisting of one or multiple `ServiceMesh` instances. A `ServiceMesh` is a directed graph (as introduced in Section 3.1.2) rooted by sources, linked by connectors/relations possibly leading to sinks. Graph edges are represented by `MeshConnector`. Each node in such a graph has an implementation in terms of a `ServiceBase`, which is refined to application-specific Java or services as well as platform-supplied services like connectors or pre-integrated services like the KODEX, the RapidMiner RTSA or the Trace2AAS service. Services declare their input and output data types, typically for forward or backward data flows (cf. Section 3.1.2). During model validation, service properties are pulled up from service level to mesh level and allow for checking whether a flow graph is valid, i.e., input/output types of the services correctly sequenced. `SubServiceMesh` represents a part of a mesh wrapped into a service (refines `ServiceBase`) and, thus, is intended as a reusable sub-mesh that can be placed like a service in any mesh where the incoming/outgoing data types of the sub-mesh fit the integrating mesh. During code generation, selected/all applications are processed, i.e., the service meshes are traversed and for each node glue code for the configured service management/ stream engine is generated. In the default case, Java classes with Spring Cloud Stream annotations are created and bound to the respective service interfaces. Based on the given plugin ids or implementation class names, implementing services are dynamically instantiated. At startup time of the app, these instances are mapped into the respective AAS (via the `ServiceMapper` from the service environment) and made available for monitoring and service management. For easing the creation of an application AAS, the configuration model also allows specifying reusable (optional) vendor information. This information is mapped into the individual AAS of the services or app, as each component that potentially is created by a different vendor is equipped with its own AAS (cf. Section 5).

Besides production code also build specifications (Maven), assembly specifications, Spring application specifications, oktoflow deployment descriptors, log setting files, JSL specifications and test classes (for validating generated descriptors) are created automatically. For the major platform services, i.e., platform service, ECS-runtime, service manager, central monitoring and management UI, the technical configuration (including AAS network, Maven repository or container setup) is instantiated into respective Yaml application settings. Finally, the generated build specifications are executed so that for a complete instantiation, the binary artifacts of the major platform services as well as a combined Java/Python artifact per application is generated/deployed to the configured Maven repository.

We now provide a brief insight into the platform/application configuration using a simple model that we also use for configuration plugin and management UI regression tests. Most⁸⁶ of the test/example

⁸⁶ Essentially, a managed platform configuration is semantically the same as joining all aspects into a single IVML project/module, but the separated modules are easier to handle by the platform.

projects follow the rules of a **managed platform configuration** (cf. Section 6.2) as it would be maintained by the management UI through the configuration AAS based on the operations implemented in the configuration component/plugin. A managed platform configuration separates the following aspects into individual IVML files/modules particularly including technical setup, (all) types, (all) services and individual files per application and mesh. These modules are tied together into a single top-level platform configuration module through IVML imports.

```
project TechnicalSetup {

  import IIPEcosphere; // still the original name

  annotate BindingTime bindingTime = BindingTime::compile to .;

  // ----- component setup -----

  serializer = Serializer::Json;
  // serviceManager, containerManager are already defined

  aasServer = {
    schema = AasSchema::HTTP,
    port = 9001,
    host = "127.0.0.1"
  };

  // further AAS omitted, looks similar

  // ----- transport -----

  transportProtocol = TransportProtocolAMQP {
    port = 8883,
    security = { // -> identityStore.yml
      authenticationKey = "amqp"
    },
    gatewayPort = -1
  };

  serviceProtocol = ServiceProtocolAMQP {};

  // ----- monitoring -----

  // current default: no monitoring configured

  // ----- UI -----

  // current default: no UI configured

  // ----- freezing -----

  freeze { // only those mentioned/relevant here
    serializer;
    aasServer;
    transportProtocol;
    serviceProtocol;
    .; // "." means every variable declared in this project
    // but: freeze stated variables except for those marked for runtime change
  } but (f|f.bindingTime >= BindingTime::runtimeMon);

}
```

Figure 43: Simplified Technical platform configuration.

Figure 43 depicts the technical part specifying the setup of the platform. IVML is textual DSL for variability modeling. Each model is surrounded by a project declaration forming a namespace, here for `TechnicalSetup`. Within that namespace, first model imports are stated, here an import of the metamodel (still named `IIPEcosphere`) and the declaration of annotations, orthogonal variables

added to all variables declared in the namespace, here the so-called binding time, the latest point in time when decisions for configuration values must be made. After this header, the configured values are stated, typically as value assignments to typed variables (a typed variable indicates configuration option/decision in IVML). Essentially, typed variables can be primitive (e.g., Boolean, Integer, String, Real), of type enum or of type compound, a complex type consisting of typed fields. Here, the wire format data serializer (variable is defined in the metamodel) is set to Json, an enumeration literal for serializers defined in the metamodel. Then the global `aasServer` (a compound variable) receives its schema, port number and host name (similarly but not shown for AAS registries and local AAS asset implementation server). Similarly, the transport protocol (with an authentication key pointing to the identity store) and, accordingly, the service transport protocol is specified, all using the wire format configured before. Device management looks similar but is omitted in the example. As no monitoring and management UI are desired, which is the default setup in the metamodel, we can skip these parts here.

```
// ----- data types (in AllTypes) -----

RecordType rec1 = {
  name = "Rec1",
  fields = {
    Field {
      name = "intField",
      type = refBy(IntegerType)
    }, Field {
      name = "stringField",
      type = refBy(StringType)
    }
  }
};

// ...

// ----- services (in AllServices) -----

Service mySourceService = JavaService {
  id = "SimpleSource",
  name = "Simple Data Source",
  description = "",
  ver = "0.1.0",
  deployable = true,
  asynchronous = true,
  class =
    "de.iip_ecosphere.platform.test.apps.serviceImpl.SimpleSourceImpl",
  artifact = "de.iip-ecosphere.platform:apps.ServiceImpl:" + iipVer,
  kind = ServiceKind::SOURCE_SERVICE,
  output = {{type=refBy(rec1)}}
};
```

Figure 44: Data type and service definition from `AllTypes.ivml` and `AllServices.ivml`, respectively.

Data types and services to be used in an app are defined in a managed platform configuration in further IVML modules. Fragments without project namespace are illustrated in Figure 43. In contrast to the technical setup, data types and services are defined as own variables, i.e., variable declarations. `rec1` is a variable of type `RecordType`, which, in turn, is defined as compound type in the metamodel. A `RecordType` has a name (turned e.g., into a class name during instantiation) and a set of fields, each with at least a name and a type (further properties are whether the value is nullable, optional, the description of the value, an optional semantic identifier denoting the value unit as well as display settings for automated dashboard creation). Types, here the type of a field, are stated in terms of a

reference to the type definition (also a variable in IVML, the reference is stated by `refBy`). Here we define an integer and a String field using pre-defined types from the metamodel⁸⁷.

Based on the definition of the RecordType `Rec1`, i.e., the IVML variable `rec1`, we introduce a Java service as a data source (that will create arbitrary data of type `rec1`). The source is described by its `id`, its `name`, an empty `description`, a `version`, whether it is `deployable`, whether it is a `synchronous` or `asynchronous` service and its implementation `class` located in the given `Maven artifact`. Please note that we use here the version of the platform defined by the metamodel in the variable `iipVer`. The service is a `source service` (one of the main service kinds) and its output is constituted by one record of the type represented by `rec1`, i.e., the RecordType `Rec1`. In fact, multiple types can be given as sequence (the outer brackets) of the compound `IOType` (the default element type⁸⁸ of that sequence) given as compound value holding the type (thus the inner brackets).

```
// ----- application and service nets -----

Application myApp = {
    id = "SimpleMeshApp",
    name = "Simple Mesh Testing App",
    ver = "0.1.0",
    description = "",
    services = {refBy(myMesh)}
};

ServiceMesh myMesh = {
    description = "initial service net",
    sources = {refBy(mySource)}
};

MeshSource mySource = {
    impl = refBy(mySourceService),
    next = {refBy(myConnMySourceMyReceiver)}
};

MeshConnector myConnMySourceMyReceiver = {
    name = "Source->Receiver",
    next = refBy(myReceiver)
};

MeshSink myReceiver = {
    impl = refBy(myReceiverService)
};
```

Figure 45: Application and service mesh part of a simple platform configuration.

The second part of the example in Figure 45 defines an application with a simple service mesh, also depicted without project namespace. First an application is defined, with `id`, `name`, `version` and empty `description`. In `services` the service meshes are stated, here a single reference to `myMesh`. `myMesh` may consist of multiple sources, here just `mySource` of type `MeshSource`. `mySource` uses the previously defined `mySourceService` as implementation and states the connection to next mesh element along the data flow graph in terms of a `MeshConnector`, which, in turn, points to the next

⁸⁷ We do not use IVML types here as IVML is a domain-independent language. We aim for typical Industry 4.0 types, such as `int32`. Thus, the metamodel defines an own symbolic type hierarchy that is turned into implementation types by the code generation of the instantiation process.

⁸⁸ `IOTypeWithPath` is an `IOType` with more fields for certain connectors, e.g., as value `IOTypeWithPath{type=refBy(rec1), path="abc"}`.

data processor, here a MeshSink. The implementation of the receiver sink is myReceiverService, which is defined similar to mySourceService but not shown here.

Besides structural aspects, the metamodel also defines validation constraints, e.g., that types must have unique names, services are configured correctly or services meshes fit together. By default, we perform a validation before instantiation and in the management UI after each change. Validation is important, as an invalid model typically leads to invalid artifacts that, e.g., cannot be compiled.

The code generation creates more than 14 different types of artifacts (Maven XML, assembly XML, Java source, Python source, application Yaml, logging XML, Java test code, Windows batch/Linux shell startup scripts, Linux/system service descriptors, README files, Broker setup specifications, Docker files, Type script files, Angular environment setups), which leads to different types of artifact structures, e.g., various forms of Java code. The number of generated artifacts varies with the number of services/mesh elements defined per application/platform configuration.

6.1 Modeling Patterns

As already illustrated in Figure 42, oktoflow's configuration model consists of several layers reflecting the architectural layers of the platform. Each layer of the model defines the decisions to be made that are related to that layer. This section dives a bit deeper into employed (meta) modeling patterns.

```
decision = X::Alternative2;

enum X {Alternative1, Alternative2, ...};

X decision = X::Alternative1;
```

Figure 46: IVML model pattern for simple alternatives without detailing properties.

Figure 46 shows the IVML model pattern to represent **simple alternatives** that do not need to be detailed further, e.g., the transport layer serializer format. The lower box illustrates the meta model, the upper box a specific configuration. Alternatives are defined (in the lower box) as an enumeration type *X* listing all potential alternatives as well as a variable representing the *decision*. The metamodel may assign a default value to ease creating a configuration. The configuration (upper box) overwrites the value to indicate that a different alternative shall be used in the platform instance. However, this pattern does not allow for openness as IVML enums are fixed and cannot be extended.

```
decision = Alternative2 {
  // assign properties as needed
};

abstract compound X {
  Type p = default;
}

compound Alternative1 refines X {
  //optional further properties, constraints
}

// further alternatives, constraints

X decision = Alternative1 {
  p = default1
  // assign further properties
};
```

Figure 47: IVML model pattern for alternatives with detailing properties.

Many alternatives demand further information, e.g., the transport protocol, or the AAS server settings. In this case, we model **extensible alternatives with detailing information**: an abstract base compound defines a common type for all alternatives and refining compounds allow for individual configuration of the alternatives. In Figure 47, the base type is compound *X*, which defines properties that are common for all alternatives, e.g., a server port, usually with default values. The individual alternatives such as *Alternative1* refine the base type *X* may add further properties. Each type representing a specific alternative can define constraints that become active only if that specific alternative/type is used. The alternatives may override the default values by re-declaring the properties with the same name and the same or a refining type. Also here, the metamodel defines a variable representing the respective decision and assigns a default instance of *Alternative1* including specific values for the properties. The configuration in the upper box may then demand a different, here *Alternative2* with respective properties. To increase consistency, property values can be derived from, e.g., common global variables. In contrast to enums, modeling of alternatives via compound refinements allows for openness, as further refining alternatives can be defined on any upstream model level.

In several situations, the configuration model must remain open for extensions by the user or by third parties. Typical **model extensions** are specific service or connector definitions as well as service meshes or app definitions in a managed platform configuration. Akin to the configuration model, also the instantiation process must be extensible, e.g., to perform service-specific generations when an externally provided generic service shall be integrated. The EASy-Producer languages IVML, VIL and VTL offer mechanisms to support such requirements.

- One basic mechanism is **dynamic dispatch** of constraint operations as well as for code generation. In dynamic dispatch, the operation to be executed is dynamically selected based on the actual types of all parameters [EQS+16]. This allows to consider (unknown) extensions of the actual model through type refinement.
- A second mechanism allows for a dynamic model structure through **wildcard imports**, i.e., model parts that may even not be known at modeling time are dynamically added to the model structure. However, imports just make model elements known to other model elements and do not influence the dynamic dispatch mechanism. Thus, the EASy-Producer languages provide a special import statement, that allows to extend the dynamic dispatch in the model at hand to dynamically loaded models.




Figure 48: Model structure for openness and extensibility.

To exploit these mechanisms, we employ the model import structure as illustrated in Figure 48, also illustrating oktoflow's naming convention for such situations. Let *X* be a model module for that we want to equip with openness, e.g., assume *X* to be *Devices* for the properties of edge, server and cloud devices. Instead of placing all related model elements into *X*, we split it up into a base model for the aspect of *X*, extensions that contribute specialized aspects and, finally *X* composing base and extension into a usable model of the platform.

- The base module *XBase* (i.e., *DevicesBase* for *Devices*) defines the basic types, constraints or operations the dynamic extensions shall hook into,
- The extensions import the *XBase* and add refined/own types, constraints and operations. The names of the extensions shall be composed from a prefix that makes the extensions uniquely importable through wild cards and an extension specific suffix that is matched by the wildcard.

In oktoflow, the prefix consists of the original module name *X*, the infix “Part” and the individual name of the extension, e.g., `DevicePartPhoenixContact`⁸⁹ for specific device types contributing Device properties of the AXC PLC/edge series of Phoenix Contact.

- Finally, some modules must use the extensions. For convenience, we let *X* do this job and let using modules import *X*. For this purpose, *X* must be an extension of *XBase* that dynamically imports *XPart** so that the extensions become available and hook into the dynamic dispatch of *XBase* and, thus, *X*. Often *X* is empty except for basic and wildcard imports as well as the model extension statement.




Figure 49: Metamodel concepts for defining services and alternatives.

The configuration of user-defined applications on top of the platform, the involved services, their data flows and the resources to execute the services on requires more information and, thus, is more complex than the three IVML model patterns discussed before. The most relevant configuration concepts for applications are illustrated in terms of the UML-like class diagram in Figure 49:

1. Configure **re-usable services** (the *Services* module in Figure 49) including platform-supplied and user-defined services, their properties including families of equivalent services that can be flexibly exchanged, e.g., alternative AI services.
2. Represent **data transformation and mappings** to reduce the effort of manual coding in standard situations, e.g., to already transform data in a connector (illustrated in Figure 50) or to specify conditions or actions for behavioral state machines. For a connector we specify an external and a platform-side app I/O data format, usually instances of `RecordType`. Fields with same (nested) field names are mapped onto each other in both directions, extern-to-app and app-to-extern. Fields that cannot be mapped are either ignored, i.e., filtered out, or left uninitialized. To fill such fields, assignment expressions between both sides can be stated, allowing for simple data transformations, e.g., unit calculations. For a model connector (cf. Section 3.5.2), the given data formats can be turned by the generator into connector-specific paths and data can be obtained and transformed automatically. For a channel connector (cf.

⁸⁹ Depending on the naming, the prefix may be adjusted deliberately for a more “speaking” name.

Section 3.5.2), the input is always binary and requires further mechanisms. A configurable parser turns the data into named/indexed fields that are then mapped by the connector to the external I/O format, which is further processed as for model connectors. In the opposite direction, a configurable output formatter takes the pre-processed data in App I/O format and maps it back to the external I/O format. If no platform-supplied parsers/formatters fit to the data at hands, own Java components supplied as Maven artifact can be specified. Similarly, the entire mapping process can be treated manually by own serializers or model adapters.




Figure 50: Overview of low-code data mapping for connectors.

3. Configure **physical and logical compute resources** the services are executed on (the *Devices* module in Figure 49). While the device management cares for the actual instances, the configuration model focuses more on kinds/categories of devices with same properties, e.g., to steer the container generation.
4. **Compose connectors and services to applications** (the *Applications* module in Figure 49) so that one service can occur in multiple applications, that the data flows within an application can be defined and can be instantiated automatically.

As already mentioned in the overview, the most basic module in Figure 49 is *MetaConcepts*, which aims at generic concepts of configurable runtime-adaptive systems. Thus, *MetaConcepts* introduces basic notions of resources (CResource), components (CComponent), families of components (CFamily) and connectors among components (CConnector). As these concepts define properties using these types, which are re-defined by oktoflow's modules, e.g., in the *Services* module.

From the generic *MetaConcepts* perspective, we now turn to the specific configuration concepts. The *Services* module currently just introduces the notion of a device, where additional properties will be added in the future. The *DataTypes* module introduces the ability to express types that are reflected in Industry 4.0 or relevant programming languages, such as String or Integer or int32, but also more complex, composed types (RecordType). These types are used in the *Services* module to specify the inputs and outputs of individual services. Specific service types represent platform implementation concepts, e.g., Connector for the connectors in Section 3.5.2, JavaService for services implemented in Java and PythonService for services implemented in Python as well as service/connector implementations shipped with the platform (e.g., AAS, OPC UA and MQTT connectors from Section 3.5.2).

On top of these configuration layers, the *Applications* module defines data flow graph among services, the service meshes. An Application consists of one or multiple ServiceMesh instances, and, in turn, a service mesh starts at one or multiple sources (of type MeshSource). Sources are linked via MeshConnector instances to processor or, ultimately, sink nodes.




Figure 51: Instance view on a platform application.

As illustration of these concepts, Figure 51 shows how instances of the aforementioned types can be linked together (broadcasting backward flows are not shown in Figure 51). The Application consists of one ServiceMesh, which, in turn, consists of a chain of three services, a source, a processor and a sink, all linked by instances of MeshConnector. The source is implemented by an OPC UA connector, the processor by some Python implementation and the sink by some Java implementation. Each service has its own input/output types, which must comply with the predecessor/successor services along the data flow of the service mesh. Further, each service is (at latest at startup time of the application) deployed to a certain resource.

6.2 Configuration Model Structure

The overall structure of the configuration metamodel as shown in Figure 52 (an abstraction of Figure 49) follows the layering of the platform architecture. As stated above, `MetaConcepts` capture the configurable basics of an adaptive software system. `DataType` defines primitive and extensible data types used for specifying input/output types of services and connectors. On the next level, `AAS` server/implementation properties, `AAS Nameplates` with default instances, `Devices` including capabilities and requirements, and `Resources` (including resource management and monitoring) are configured. Then, the transport protocol (including authentication and transport layer security) is specified. `Services` (including Java and Python services), reuse all these configuration types and form a basis for the more specific `Connectors`. Applications and their service meshes are defined on top, imported by `UI` (settings) and finally the top-level `IIPEcosphere` module for global technical installation settings such as installation paths. Module names suffixed with a `*` are extensible through the import mechanism and modeling patterns explained in Section 6.1

Besides this main structure, `DataOperations` defines the structure of operator-expression trees for specifying conditions or state actions. The `Connectors` use these expressions to allow configuring data transfer actions between external data sources/sinks and apps in low-code manner. The `ServiceStateMachines` (since version 0.8.0⁹⁰) allow for modeling service behavior in terms of state machines, where the expressions are used for transition conditions, time conditions and state actions on local state machine variables.

⁹⁰ We will cite here the MSc thesis of Monika Staciwa as soon as it is published and available.

An actual copy of the metamodel is included in an instantiated platform as well as in each implementation/example project after downloading the metamodel as Maven artifact as part of the respective build process. In some of these projects, the actual platform/app configuration is typically defined as a free structure using one or two IVML modules, in particular to focus on the specific properties of the example and to ease gaining an overview.




Figure 52: Basic structure of the IVML configuration metamodel.

When the platform takes over control of its own model, fixed structures are needed so that configuration.easy can store settings, service instances and application meshes in distinct places. This **managed platform configuration** structure required then is depicted in Figure 51. In this setup, the topmost module is the PlatformConfiguration storing settings that override global non-frozen configuration options. Service instances are stored in AllServices, related type definitions for input/output specifications in AllTypes. In turn, AllTypes relies on AllConstants, containing e.g., server host names or commonly used port numbers. The application instances (ApplicationPartX with X being the application name) are stored in individual extensions/folders pointing to their linked service mesh parts (ServiceMeshPartX with X being the mesh name). The top-level metamodel module IIPEcosphere and transitively imported modules are linked through AllConstants into the managed configuration model. Further, TechnicalSetup defines the fundamental technical abilities of the platform, e.g., transport protocol, monitoring or management UI setup. In turn, TechnicalSetup may rely on the constants in AllConstants. In addition, the templates folder contains the application templates (TemplatePartX) introduced in version 0.8 for the ReGaP innovation community.




Figure 53: Managed platform configuration structure.

Besides dynamic imports, there are also intentionally unused IVML modules in the metamodel, in particular for not regularly used types. As the IVML models for OPC UA companion specs or IDTA AAS specs can become rather large, they, if imported without care, may heavily impact the speed of model operations. However, individual applications may need these files, through model plugins even add further files. While adding such imports to PlatformConfiguration is straightforward, this is not the case for the management UI. There, the configuration (meta-)model is used as a consistent data basis for the management user interface. Besides AAS operations, e.g., to add variables, or entire service meshes, we also need operations to obtain unimported IVML modules as well as to add or to remove them from a managed platform configuration.

Due to technical reasons, the transformation of an IVML configuration into an AAS may, dependent on the used AAS/implementation version, impose some limitations:

- IVML variables shall not be named `value`, as this `idShort` may be a reserved name. We use name prefixes, e.g., `aValue`, which, at a glance, may be a bit confusing. As an approach, affected IVML variables can be annotated with a `displayName`, overriding the displayed name on the UI. In further cases, the type name may cause an implicit renaming, e.g., for fields that cannot be named as intended to to conflicts with IVML keywords, e.g., `ver` of type `OktoVersion` are displayed then on the UI as `version`.
- IVML variables of type `Any` in `MetaConcepts` are intended to be redefined with a specialized type. However, if this does not happen, we assume them to be superfluous and skip them when creating the configuration AAS, so that they do not occur in the UI.
- IVML variables can be annotated with the integer value `uiGroup`. Elements that shall not be visible shall have the value 0, positive numbers denote mandatory groups in decreasing order of priority, similarly negative numbers associated optional groups.
- Further meta-properties steer the interpretation of the model by the UI, e.g., `size` for collections, `abstract` for abstract types, `refines` for parent types, `required` for required fields, `constant` for constant elements, etc.

6.3 Special configuration topics

Besides the main configuration concepts there are some specific topics. We will detail these steps incrementally, here as well as in the github online documentation.

6.3.1 Deployment plans

Deployment plans are needed to conveniently distribute, start and stop applications. A deployment plan states for each service the target device, on which ECS-runtime and service manager, e.g., in terms of containers must be installed. Deployment plans can directly be stated as YAML files and uploaded through the management UI or specified using the same concepts in the platform/application configuration. If stated in the configuration, the deployment plans are instantiated and made available as if uploaded.

6.3.2 App Templates

In certain cases, it might be difficult to start off application modeling. App templates try to address this issue by allowing "applications with intended gaps". This templating concept is new and part of the ongoing work in the ReGaP innovation community.

- An app template is an oktoflow application for which not all required fields may be filled. Thus, an app template cannot be instantiated as an usual application. In contrast, the management UI must identify the gaps and explicitly ask the user how to fill these gaps prior to turning the template into an app on behalf of the user. The resulting app must then comply with all configuration rules of oktoflow's application metamodel.

- Following the managed configuration structure, apps reside as single IVML projects (name prefix `TemplatePart`) in the optional `templates` folder. Templates may contain constants, data types (including dashboard information), services, connectors, service meshes and application instances.
- In particular, data types may overlap between different templates; if two data types have the same name, only one is taken over into the platform configuration.
- All services used in a template app shall be platform services, i.e., shipped with the platform, preferably generic platform services that can be operated with arbitrary data (possibly detailed by additional configuration information).
- Simple gaps are fields that are not filled out or configured in a conflicting way. Alternative services can be represented by service/connector families.
- If specialized data types are needed, applications created from a template can define data type substitutions which are taken into account during app instantiation. To be effective, service implementations must then use the generic ways to create instances of data types or to transfer data among data type instances.

As initial validation/example, we equipped the `SimpleMesh` testing app in `configuration.easy` with a simple template. Further concepts or support operations may be needed; their design and realization depend on the evaluation of the app template concept.

6.3.3 Identity store

Primarily, the identity store aims at encapsulating complex secure identity and keystore mechanisms. Basically, the keystore is just a YAML file that details accounts, passwords, keystores or, more generally, identity tokens.

If an app is derived from an app template, the required identity information must be either made accessible by default or specifiable by the user, preferably using the gap mechanism if the app template instantiation. Thus, in particular for this purpose, in contrast to the usual practice of not having security information in the platform/application configuration, the configuration metamodel allows for specifying the contents of an app's identity store, which is turned into identity store files by the platform/application instantiation.

6.3.4 State machines

For Java services in general and for app templates in particular, the oktoflow metamodel now, as part of the ReGaP innovation community and through the MSc thesis of Monika Staciwa, supports modelling of behavior in terms of state machines with local variables. State machines represent state-based lifecycles with conditional or timed transitions, transition actions as well as timed constraints and actions in states. Actions, expressed in terms of (extended) `DataOperations` may access and modify local variables. A future version of the management UI shall include for this purpose an extended version of the expression editor `TransEx` [Sal25]. The platform/application instantiation recognizes the optional reference from a `JavaService` to a state machine and, if present, generates a refined abstract base class for that service, into which user code may need to hook into.

6.3.5 Dashboard support

Visualization of application results is an important step in Industry 4.0 applications. As part of the work in the ReGaP innovation community, the metamodel was extended so that `RecordType` fields can express their value unit in terms of a semantic id (e.g., `ECLASS`), their display name, their target display row and their target display panel. For more reuse of information, also datatypes as well as alias types can specify their value unit as semantic id, whereby the default semantic id of the field is taken from the employed data type. Such a value type semantic id may be known to the semantic id resolution mechanism in the support layer to be shown with the respective value unit name on the management

UI. More importantly, the mapping of semantic ids to display technologies shall be specified in the translation of the display information to an AAS-JSON describing the desired dashboard. Currently, the translation is a Java program, but shall be, along with the AAS plugins from the support layer, become an EASy-Producer VIL instantiator. From the viewpoint of data flows, the relevant results of an oktoflow app shall be stored in a timeseries database using the Influx connector. An additional Python program (by Bitmotec) takes up the AAS-JSON, turns it into a Grafana⁹¹ dashboard and links the dashboard with the Influx DB used in the oktoflow app.

6.4 Support for Standardized Connectors/Protocols

Defining the input/output data types for complex, nested data structures can be a non-trivial process. For standardized protocols/information models such as OPC UA or AAS, more and more data structures with standardized fields and semantics are created. For a standard-based platform like the oktoflow it is important to take up such approaches and to ease the use of standardized data structures, thus, supporting the platform user in creating applications.

As an example, the **OPC UA Companion Specifications**, is a set of standardized models for OPC UA. Currently, more than 50 such models have been specified and further models are in preparation. Over time, also the defined model structures are evolving. Thus, manually translating OPC UA Companion Specifications into IVML for further processing in the platform is not sustainable. Fortunately, the OPC Companion Specifications are available in a machine-readable XML format, which can automatically be translated into IVML and then used further in application configurations. We demonstrate this approach by an automated model translator [Cep23].

The model translator reads an OPC Companion Specification XML file and translates it into IVML using a base metamodel, which extends the configuration metamodel, in particular the `DataTypes` module. The created IVML files, can be imported into an application model when needed. Besides the main types representing a Companion Specification, also declared subtypes can be used in custom applications, in particular as input/output to generated OPC connectors.

The approach was successfully validated with all 55 available OPC UA companion specs. Valid IVML models were produced for all companion specs and valid connector code was generated. The largest specs, e.g., the Tobacco Machine Communication (TMC)⁹² or the IEC61850-7-4 spec triggered a restructuring of the generated connector code to cope with the more than 120 KLOC XML specification. As far as available in the VDW UMATI OPC test server, 3 of the generated connectors that also can operate with the NodeIds from the specification were successfully functionally validated against an OPC reference implementation [Cep23].

Similarly, we approached in [EW24] the IDTA AAS submodel specifications. The optional IVML module `AASDataTypes` extends `RecordType` and `Field` from oktoflow's metamodel to represent AAS structures used in the submodel specifications. A specialized parser can read IDTA PDFs (transferred into XML for reading) or AASX files into the types of `AASDataTypes`. Besides directly using the structures akin to OPC UA companion models, and a special variant of oktoflow's code generator turns the IDTA models into Java APIs, specializing the AAS abstraction from support in a way, that the platform can easily create standard-compliant AAS submodels. However, when publishing [EW24], we also found several inconsistencies in the PDF specifications as well as in their accompanying AASX files.

6.5 Platform Instantiation Process

After successfully configuring a platform and the related apps, the configuration must be instantiated. This happens through further languages of EASy-Producer [VIL], namely the Variability Instantiation

⁹¹ <https://grafana.com/>

⁹² <https://reference.opcfoundation.org/TMC/v200/docs/8.1>

Language (VIL) for the instantiation process and the Variability Template Language (VTL) to modify or create artifacts of a certain type, e.g., Java or Python code files or XML for Maven build specifications.

The VIL model defines several **entry points**. Figure 54 illustrates selected entry points and selected steps that are executed during the instantiation. The entry points are

- `generateInterfaces` generates the interfaces of the applications required for the realization of user-defined services as well as default basic implementations of the service interfaces, e.g., to ease implementation of parameters or data ingestors. These classes are deployed as Maven artifacts to be used by application code templates, which are also generated when executing this entry point.
- `generateAppsNoDeps`⁹³ instantiates the applications, e.g., the glue code binding the application interfaces with the service execution engine, while intentionally skipping implementation dependencies. Although these applications would not run, the user can figure out if they would build at all. Also generates the application code templates.
- `generateApps` for the instantiation of the configured apps for a platform including all required dependencies. Additional parameters can limit the generation to certain apps (via their id). The resulting app artifacts are packaged to be executable. If enabled, creates respective application containers.
- `generateBroker` generates a default broker/service instance for the configured transport protocol, e.g., for regression testing, app development testing etc. These instances may not be used for production, but they are helpful for a first setup or for executing examples.
- `generatePlatform` for the instantiation of the major platform services, i.e., customized versions of service manager, ECS-runtime, monitoring, management UI and platform service. If enabled, creates respective platform containers.
- `generateApi`, generates API classes, e.g., for IDTA AAS submodel specifications.
- `main` (not shown in Figure 54) which executes all aforementioned entry points, in particular for testing.

For the instantiation of the **application interfaces**, we first iterate over all data types declared in a platform configuration and create their Java and Python realization (`JavaType`, `PythonType`). Moreover, we create the related serializers based on the declared types, for Java in order to realize the platform transport wire format (e.g., JSON, XML, protobuf) and for Python a JSON-String Serializer for the fixed data communication with the Python Service Environment. For all services in the platform configuration, we generate the service interfaces (`JavaServiceInterface`, `PythonServiceInterface`) and where feasible a basic implementation for service parameter and ingestor handling (`JavaServiceBaseImpl`)⁹⁴.

Please note that, as discussed in Section 3.6.2, although Java and Python Service Environments are similar, they also differ in the required level of programming, which is reflected in the different instantiation steps. Thus, Java services demand a more complex code generation due to the direct integration into the service execution engine than Python services that are “just” executed through the Python Service Environment, which, in turn, is executed by generic Java services. When completing the generation of application interfaces, we create a Maven assembly descriptor for the Python

⁹³ Since version 0.8 considered as legacy/deprecated due to the plugin-based creation of service and connector instances. May be removed in future versions.

⁹⁴ To allow for service implementations that are not based on the basic implementation, the default values of the service parameters are set during app startup through the reconfiguration operation, i.e., the parameter values may not be available during constructor execution but when the first data arrives for processing.

interfaces and base classes, a Maven build specification that creates the deployable artifacts as well as an Ant script to execute the deployment in the continuous integration.

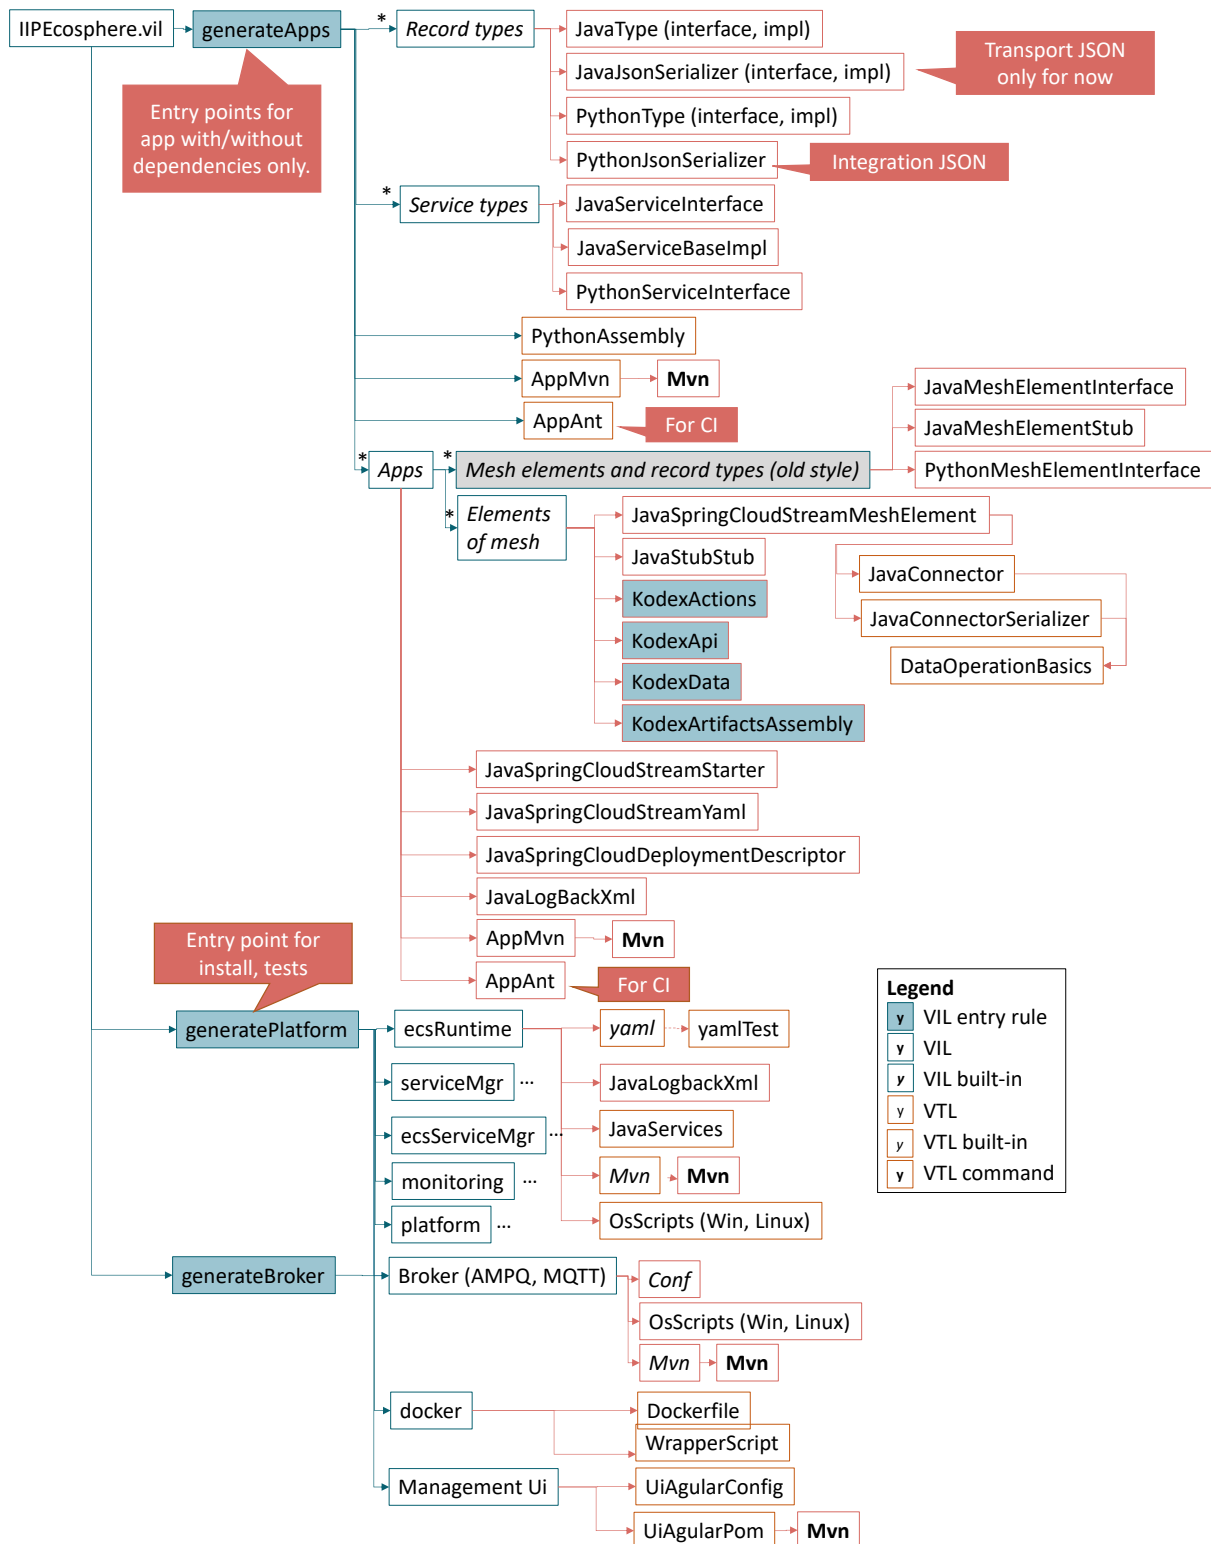


Figure 54: Overview of the platform instantiation process.

Further, we generate **application code templates** (not indicated in Figure 54). For each Java or Python services to be implemented, the application code template contains a code skeleton, a unit test, respective input data template files and a service level test. For each connector used in the app, a connector connectivity test is generated. If required, specialized assembly descriptors e.g., for Python

services are created. Finally, a Maven build specification, bash and shell scripts to ease the Maven execution as well as specialized files for the use of the template in an Integrated Development Environment (IDE)⁹⁵ are generated. It is important to mention that

- As the code generation is not allowed to overwrite user code, the template must be obtained, extracted and loaded into an IDE. When the configuration of the respective app is updated, a careful comparison of the potentially re-generated template files and the own code based on the original templates is required, as, e.g., data types, field access, service methods and test input data files may have changed.
- Templates become particularly usable after the first execution of the build process as then oktoflow's meta model, the respective Java dependencies, the generated Python interface code and the Python service environment are made available by downloading/unpacking them from the Maven repository where the generated application interfaces have been deployed.
- The meta model and, similarly, the Python interface/service environment code is not updated automatically rather than on user request. This behavior is intended to leave the decision to the developer when to cope with potentially changed code.

For obtaining the **applications** (integrating handcrafted service code), we iterate over all application configurations and all their linked service meshes. Here, we create in particular code to bind the respective service into Spring Cloud Stream (`JavaSpringCloudStreamMeshElement`). Depending on the actual service, also further artifacts and assembly descriptors are generated, e.g., as shown in Figure 54 for KIPROTECT KODEX. Furthermore, for each app, the deployment descriptors, a starter class (for registering the mesh elements in the service framework and for registering the serializers), and the Maven POM file (`AppMvn`) are created. The POM file is executed, ultimately creating application artifacts, including the packaged application service artifact. Further, for purposes of continuous integration, an ANT build file is created. If enabled, application containers are created, i.e., Docker containers that contain all required dependencies to execute the configured applications.

On platform level, the instantiation process creates packaged artifacts of the **major platform services**, i.e., the ECS-Runtime, the service manager, a combined version of ECS-Runtime and service manager as well as the central platform (AAS) services. For each of these components, first the application setup file (`yaml`) and a test class to validate the Yaml file are created. Then the logging configuration, the selected JSL service descriptors and ultimately a Maven POM with the respective components selected in the platform configuration are created. The Maven POM is executed, obtaining the respective artifacts, i.e., folders containing all required binary dependencies. Moreover, for starting the components, operating-system specific scripts for Windows and Linux (also descriptors for operating system services) are generated. Similarly, the configured transport protocol leads to the instantiation of a corresponding default broker. Further, the platform management UI is instantiated to comply with the network and AAS setup and, if enabled, Docker containers are for the major platform services.

6.6 Container Instantiation

Virtualization of platform services or applications intertwines technical requirements with ease-of-use. On the one side, devices, in particular edge devices, may not allow for extensive software installation outside their own ecosystems. In such environments, Java/Python in general as well as certain versions of dependencies may not be available. Here virtualization, a recent trend in IIoT, may be key [ESA+21, AE24]. On the other side, complex services require non-trivial dependencies, which, in particular for Python, may require a physical installation and, in turn, may demand installation of native libraries, the execution installation procedures or code compilers. Such operations may not be permitted on a

⁹⁵ Primary target is Eclipse, for which also a Pythonpath tooling configuration is generated. However, so far we were not successful creating a similar setup for VSCode.

target system. Further, separating dependencies among different (versions) of applications and enabling removal of installations may end in a nightmare without virtualization.

Creating adequate containers for an IIoT application requires technical knowledge and may involve non-trivial tradeoff decisions, e.g., pre-installed vs. dynamically installed dependencies, layering of target containers respect to desired system properties (transfer time, upgradability, adaptability), etc. As support, the platform enables automatic creation of containers based on a set of strategies. As ideally no human is involved in this process, this allows for application and device-specific containers (e.g., considering vendor builds of dependencies) instead of one-size-fits-it-all containers that are often used in cloud settings where resources do not matter so much [EPR+22, EPN22, AE24].

Container creation, as part of platform/application instantiation, is the process of creating container images that include an application or major platform services with all required dependencies. We primarily focus on the Docker container technology, but are open to other approaches, e.g., LXC. Container creation must be enabled explicitly (IVML variable `createContainer`). Then, based on the technical setup of `ContainerType` in the configuration model, three different types of container images are created. Successfully created container images are stored in a (public, private, local) container registry as also defined in the configuration model.

Ecs_Svc_App	EcsSvc_App	C1Ecs_C2Svc_App
- Local broker - ECS-Runtime - Service manager <i>Application dependencies</i>	- Local broker - ECS-Runtime + Service Manager <i>Application dependencies</i>	ECS-Runtime - Local broker - Service manager <i>Application dependencies</i>
Ecs_Svc_AllApps	EcsSvc_AllApps	C1Ecs_C2Svc_AllApps
- Local broker - ECS-Runtime - Service manager <i>All dependencies for all applications</i>	- Local broker - ECS-Runtime + Service Manager <i>All dependencies for all applications</i>	ECS-Runtime - Local broker - Service manager <i>All dependencies for all applications</i>

Figure 55: Container types and contained services/parts in oktoflow.

Currently, as discussed in more details in [Sta22], the platform supports the automated creation of six **types of container images**. The container types, as illustrated in Figure 55, are:

1. `Ecs_Svc_App` creates a container for each application including an ECS-runtime and a service manager as separate processes as well as a local communication broker for intra-device communication among services.
2. `EcsSvc_App` instantiates a container for each application including ECS-runtime and a service manager as a single process. Although running the central services as own processes increases resilience, a single process allows for saving resources. As for `Ecs_Svc_App`, the container includes all service dependencies for the configured application as well as a local broker.

3. C1Ecs_C2Svc_App generates two separate containers for each application: a container with an ECS-runtime and a second container with service manager, application dependencies and a local broker.
4. Ecs_Svc_AllApps creates a container for all applications containing an ECS-runtime and a service manager as separate processes. The container also includes all service dependencies for the configured applications and a local broker.
5. EcsSvc_AllApps leads to a container for all applications including ECS-runtime and service manager as single process, all service dependencies for the configured applications as well as a local communication broker.
6. C1Ecs_C2Svc_AllApps creates two containers, one container with an ECS-runtime and a second container with service manager, dependencies for all applications and local broker.

After creating the images and pushing them to the registry, a descriptor with container image properties is generated. This descriptor also states the container startup options (environment variables, exposed network ports, utilized volumes, Docker-out-of-Docker settings for C1Ecs_C2Svc_App or C1Ecs_C2Svc_AllApps as well as the location from where to obtain the container. These descriptors are stored in the artifacts folder of the platform as specified in the configuration model, used by the management user interface to display the available containers as well as by the container manager installed on a device to obtain and to start a container on demand. If multiple device types are configured, instantiation process creates one additional container image per device type, e.g., considering vendor-local installations of certain libraries.

Further, there we support creating **shared container base image** (containerBaseImageMethod). Typically, this method reduces container creation time and storage size of the containers. The approach is illustrated in Figure 56.




Figure 56: Container base image creation

The steps for creating and using a container base image are:

1. Create a container for common dependencies including the dependencies that are shared among all devices.
2. Create an all-dependencies container image: This image uses the common dependencies container as a base image and adds all non-common dependencies.
3. Application container image or all-applications container image:
 - a. If the device provides installed dependencies, then the common dependencies container image will be used as base image. Further, potentially missing dependencies as well as device-specific files are added as layers.
 - b. If the device does not provide any dependencies, the all-dependencies container image will be used as base image and device-specific files are added as layer.

The instantiation process can also create a container image for the central platform service including management UI and monitoring if enabled (IVML variable platformContainerGeneration).

To further speed up the container instantiation process, the platform employs **semantic fingerprinting** for all the files used in containers to identify whether a container shall be created at all. Docker just compares files, which typically leads to a container creation, when a JAR file or archive has been re-created by previous build steps, i.e., the file time stamps have changed, but not actual content of the files. For this purpose, we record MD5 checksums per file (for archive per contained file) and consider them when the container shall be re-created. However, this behavior can be disabled and container creation can be enforced via the IVML variable `forceContainersCreation`.

If individual Python services specify a conda or a **virtual environment** (venv), e.g., to use different versions of the same dependency, the container instantiation will install the required Python dependencies into that conda environment. However, virtual environments also have an impact on the resource usage of a container as dependencies may be installed multiple times, i.e., there is no sharing among the supported Python virtual.

7 Implementation Aspects


In this section, we briefly discuss relevant implementation aspects of the oktoflow platform. While particular technical information is given in github, this section just focuses on overarching aspects.

7.1 New components

For easing the implementation of new components or examples, we provide several templates that already have built in conventions and development setup, e.g., the `basicMaven` project for platform components or the `impl.model` template for applications. Please refer to the online documentation.

7.2 Compiling the Platform

Due to the various optional and alternative components in the platform, manually customing or compiling the platform is not trivial. Usually, the binaries of the individual components are either available via Maven central (releases) or SSE Maven repository (snapshot, releases). These builds are created and deployed by the SSE Continuous Integration (CI) server as illustrated in Figure 57, which knows about the build dependencies among the components and builds the parts and pieces along the dependency tree when the code of a component changes. As part of building, it executes the respective component tests, assembles the documentation and, if successful, deploys the respective snapshots to the SSE Maven repository or the stable releases from Maven (or related repositories).



✓	☀	IIP_configuration.defaultLib	5 Stunden 2 Minuten #953	Unbekannt	2 Minuten 35 Sekunden
✓	☀	IIP_configuration.easy	5 Stunden 19 Minuten #128	Unbekannt	13 Minuten
✓	☀	IIP_configuration.interface	4 Tage 10 Stunden #78	Unbekannt	54 Sekunden
✓	☀	IIP_configuration.maven	5 Stunden 4 Minuten #1586	Unbekannt	1 Minute 54 Sekunden
✓	☀	IIP_connectors	4 Tage 11 Stunden #1100	Unbekannt	1 Minute 46 Sekunden
✓	☀	IIP_connectors.file	4 Tage 11 Stunden #204	Unbekannt	1 Minute 14 Sekunden
✓	☀	IIP_connectors_aas	4 Tage 11 Stunden #1123	Unbekannt	1 Minute 24 Sekunden
✓	☀	IIP_connectors_ads	4 Tage 11 Stunden #350	Unbekannt	55 Sekunden

Figure 57: Screenshot of the SSE Continuous Integration server (cropped)

For completeness, we discuss below the dependencies among the individual components of the platform (as illustrated in Figure 58). As indicated in Figure 58, some components need specific settings for successful testing, e.g., the RTSA components need to know which JDK to use for RTSA execution (strict Java 8 for the original RTSA).

The `platformDependencies` project defines the common build process steps as well as administrative information such as authors or source code management required for Maven Central. Since version 0.8, the `platformDependencies` only declares the versions for dependencies that are used by multiple platform components, but no dependencies (cf. constraint C15). Thus, `platformDependencies` serves as parent POM for all platform core components, which must not have any external dependencies except for the JDK Java class library. These version properties are turned into managed dependencies in the platform bill-of-material (`platformDependenciesBOM`), which is used as parent POM for all implementing components (cf. Section 4).




Figure 58: Dependencies among the components (folder names as in GitHub)

The bill-of-material POM is helpful to restrict the diversity of dependency versions, while, if required, individual implementing components can override the version. In an ideal case, if all dependencies are homogenized after an entire update of the dependencies, such overrides are not needed. Moreover, components and application services that rely on the application/service execution version of Spring (Cloud Streams) are based on `platformDependenciesSpring`, which extends `platformDependenciesBOM`.

Building the platform starts with compiling and packaging the fundamental library `tools.lib` as basis for the the platform's Maven plugins for dependency management (obtaining, unpacking and updating the configuration metamodel), Maven invocation (executing sequences of goals in different profiles on the same POM, shortcutting build processes for app artifact updates) and Python (syntax check and testing). As these plugins are used by the build process defined in the platform dependencies, they must be built before without upstream dependencies.

The Support Components (also based on `tools.lib`) consisting of the basic plugin mechanisms and plugin interfaces in `support.boot` (along with implementing plugins), the basic utility classes and interfaces (with their implementing plugins) in `support` and the AAS abstraction `support.aas` are the most basic components. As `support.boot` and `support` provide some basic functionality that is also used in components that strictly depend on Java 8, `support.boot` (and its implementing plugins) as well as `support` are also based on Java 8. However, plugins for which interfaces are defined in `support` but which are not directly used by `support` may be based on more recent Java versions. In particular, `support` integrates a Python helper class identifying the actual Python binary from the platform's Maven Python plugin, which, in turn, forces the Maven Plugin to Java 8. The BaSyx plugins (consisting of sub-projects for server functionality and, for BaSyx2 also of sub-projects for common and plugin testing) and the `iip-aas` support functions are built after `support.aas`.

The Transport Component (`transport`) is built after the Support Layer. Subsequently, the transport connectors (`transport.*`), the basic (optional) Spring integration (`transport.spring`) and the Spring binders (`transport.spring.*`) utilizing the transport connectors are built. As the Connectors Component (`connectors`) relies on the type translation and serialization mechanisms of the Transport Component, then the Connectors Component as well as the individual connector types (`connectors.*`) can be built, in particular as the MQTT connectors are based on the corresponding transport connectors.

The components of the service layer (`services.*`) consist of the generic service environment (`services.environment`) and the service manager interface (`services`), which are built before processing the Spring-specific service environment (`services.environment.spring`), the `service.spring.loader` for isolated class loading and the specific plugin for Spring Cloud Streams (`services.spring`). Along with `services.spring`, the CI also builds its subprojects `test.simpleStream.spring`, a testing artifact realizing a simple stream processor chain as well as `services.spring.pluginTests` for testing the Spring service manager in a class-loading isolated clean-root environment.

Next follows the resource/deployment components (`ecsRuntime.*`), i.e., the ECS-runtime and the container managers plugins for Docker, LXC and the experimental manager for Kubernetes. Subsequently, components for the device management (`deviceMgt.*`), including the plugins for Minio or S3 object storage, the `basicRegistry` as well as `ThingsBoard` for device management. Next, the components for security and data protection services (`securityDataProtection.*`) are built. Within the reusable intelligent services, the RapidMiner RTSA (version 14) requires Java 8 for execution. As RTSA is an IPR protected component, the regression testing is based on a mock version of RTSA, which needs to be built with Java 8 to mimic the prerequisites of RTSA. Then, the `platform server(s)` component hosting the server processes can be processed by the CI.

On the next layer, the platform configuration interfaces (`configuration.interface`), the default configuration technology plugin (`configuration.easy`) based on EASy-Producer (stand-alone, Maven-based integration, JDK driving force through Eclipse/xText) and the configuration component (`configuration.configuration`) providing the configuration AAS are built. `configuration.easy` contains a sub-project for defining the implementations of test application services (`test.configuration.configuration`) for testing the configuration model. The platform's configuration Maven plugin `configuration.maven` executes the platform instantiation via `configuration.configuration` and provides goals for all instantiation tasks, an encompassing testing task, which can start an entire platform/apps, as well as a goal to manipulate text files.

On top of the configuration model, there is the management UI. As UI is an Angular Web app based on the platform AAS as backend, it requires a different build process. The TypeScript code of the UI is compiled using Angular, packaged, archived by the CI server and then, using a pseudo Maven POM, deployed as binary component into the Maven repository. Testing the platform management interface heavily relies on the configuration Maven plugin, which instantiates a platform and runs it together with the Angular tests in coordinated manner. The platform instantiation takes this binary up akin to other (binary) platform components, unpacks and customizes the UI.

The Test Components (`test.*`) are a side track, build after the basic Maven plugins, and required for testing. These protocol related test components/plugins encapsulate embedded protocol brokers, such as Apache Qpid, HiveMq or Moquette, which are either used as test dependencies or plugins.

The platform examples (`example.*`) act as platform regression tests, i.e., the respective configuration model is instantiated and the application is executed in mock/testing mode using the test goal of the configuration maven plugin. Usually, the platform examples are all-in-one projects, for which the build process heavily relies on the orchestration by the platform's invoker Maven plugin (running the instantiation of the broker, the interfaces, the application build and the application instantiation/integration in coordinated manner as sub-maven executions).

If a local build is required, we provide a multi-module Maven POM on the top-level of the platform code repository. Given that all required software (Java, Maven, Python and dependencies due to building the examples) is installed, the Maven command `mvn install` builds the full platform. To speed up the build (otherwise it may take around two three), tests may be disabled (`-DskipTests`).

7.3 Installing and Using the Platform

As discussed above, the platform must be configured and instantiated before it can be installed or executed. Thus, the continuous integration or the released binaries do not provide or represent complete platform executables, respectively. Please refer to the online documentation in [github](https://github.com/iip-ecosphere/platform/blob/main/platform/documentation/INSTALL.md)⁹⁶ on how to install/start oktoflow.

⁹⁶ <https://github.com/iip-ecosphere/platform/blob/main/platform/documentation/INSTALL.md> and <https://github.com/iip-ecosphere/platform/tree/main/platform/tools/Install>

8 Summary & Conclusions

Realizing an open (experimental) IIoT/I4.0 platform is a significant amount of work. A solid foundation was performed in IIP-Ecosphere and is taken on in the oktoflow platform and in further projects such as the DATIpilot innovation community ReGaP. This handbook provides insights into the ideas, concepts, rationales and designs of the current release of the oktoflow platform. The rationale behind this document is to enable interested parties to better understand the internals of the platform. To strengthen this focus, technical documentation on installing, configuring or building apps has been migrated into the (more agile) online documentation on github. However, as the platform is evolving, this document is just a snapshot in time and will further evolve with oktoflow.

We discussed the technical basis for architecture modeling, the overview of the layered architecture, the individual layers and the components they contain. We discussed architectural constraints, the use of Asset Administration Shells (AAS), the approach to platform configuration and instantiation as well as selected implementation aspects.

In summary, the release accompanied by this handbook significantly improves oktoflow for future work, case studies and, in particular, as technical core of the DATIpilot innovation community ReGaP. For the next release, we plan for the following topics:

- More and improved modeling concepts for App templates
- Improved plugin handling and updated/evolved plugin dependencies
- Management UI, e.g., for configuring data transfer expressions [Sal25], app templates and behavioral state machines

9 References

- [AE24] A. Alamoush, H. Eichelberger, Open source container orchestration for Industry 4.0 – requirements and systematic feature analysis, *International Journal of Software Tools for Technology Transfer*, pp. 527-550, 2024
- [BBB+20] S. Bader, H. Bedenbecker, M. Billmann, A. Bondza, B. Boss, S. Erler, K. Garrels, T. Hadlich, M. Hankel, O. Hillermeier, M. Hoffmeister, M. Kiele-Dunsche, J. Neidig, A. Orselzki, S. Pollmeier, B. Rauscher, W. Rieder, S. Stein, B. Waser, Generic Frame for Technical Data for Industrial Equipment in Manufacturing (Version 1.1), *Plattform Industrie 4.0*, 2020, https://www.zvei.org/fileadmin/user_upload/Presse_und_Medien/Publikationen/2020/Dezember/Submodel_Templates_of_the_Asset_Administration_Shell/201117_I40_ZVEI_SG2_Submodel_Spec_ZVEI_Technical_Data_Version_1_1.pdf
- [Cas21] M. G. Casado, Service and device monitoring on devices in IIP-Ecosphere, IT-Studienprojekt, Universität Hildesheim, 2021
- [CE21] M. G. Casado, H. Eichelberger, Industry 4.0 Resource Monitoring - Experiences with Micrometer and Asset Administration Shells, *Symposium on Software Performance 2021, EUR Workshop proceedings*
- [Cep23] J.-H. Cepok, Projektarbeit, Uni Hildesheim, 2023
- [Eic16] H. Eichelberger, A Matter of the Mix: Integration of Compile and Runtime Variability, *Workshop on Dynamic Software Product Lines, FAS'16*, 2016.
- [IVML] H. Eichelberger, S. El-Sharkawy, C. Kröher, K. Schmid, IVML Language specification, http://projects.sse.uni-hildesheim.de/easy/docs-git/docRelease/ivml_spec.pdf
- [EW24] H. Eichelberger, A. Weber, Model-driven realization of IDTA submodel specs - The good, the bad, the incompatible? *ETFA'24, IEEE*, 2024
- [ESA+25] H. Eichelberger, C. Sauer, A. S. Ahmadian, C. Kröher, Industry 4.0/IIoT Platforms for manufacturing systems — A systematic review contrasting the scientific and the industrial side, *Journal of Information and Software Technology (IST)*, volume 179, 107650, <https://doi.org/10.1016/j.infsof.2024.107650>, 2025
- [ESS22] H. Eichelberger, H. Stichweh, C. Sauer, Requirements for an AI-enabled Industry 4.0 Platform – Integrating Industrial and Scientific Views, *SOFTENG'22*, pp. 7-14, 2022
- [EN23] H. Eichelberger, C. Niederée, Asset Administration Shells, Configuration, Code Generation: A power trio for Industry 4.0 Platforms, *ETFA'23*, pp. 1-8, IEEE, 2023
- [EPR+22] H. Eichelberger, G. Palmer, S. Reimer, T. Trong Vu, H. Do, S. Laridi, A. Weber, C. Niederée, T. Hildebrandt, Developing an AI-enabled IIoT platform - Lessons learned from early use case validation, *SASI4'22 @ ECSA'22*, 2022
- [EPN22] H. Eichelberger, G. Palmer, C. Niederée, Developing an AI-enabled Industry 4.0 platform - Performance experiences on deploying AI onto an industrial edge device, *Symposium on Software Performance (SSP'22)*, 2022
- [ESA+21] H. Eichelberger, C. Sauer, A. S. Ahmadian, M. Schicktanz, A. Dewes, G. Palmer, C. Niederée, IIP-Ecosphere Platform – Requirements (Functional and Quality View), Version 1.0, March 2021, IIP-2021/02-en, DOI: 10.5281/zenodo.4485774, 2021
- [EQS+16] H. Eichelberger, C. Qin, R. Sizonenko, K. Schmid, Using IVML to Model the Topology of Big Data Processing Pipelines In *Proceedings of the International Systems and Software Product Line Conference SPLC'16*, p. 204 – 208, 2016.

- [EW25] H. Eichelberger, A. Weber, J. Hildebrand, ADS Performance Revisited, Softwaretechnik-Trends Band 45, Heft 1, 2025
- [VIL] H. Eichelberger, K. Schmid, EASy Variability Instantiation Language: Language Specification, http://projects.sse.uni-hildesheim.de/easy/docs-git/docRelease/vil_spec.pdf
- [IDTA 02003-1-2] IDTA 02003-1-2 Generic Frame for Technical Data for Industrial Equipment in Manufacturing (https://industrialdigitaltwin.org/wp-content/uploads/2022/10/IDTA-02003-1-2_Submodel_TechnicalData.pdf)
- [IDTA 02004-1-2] IDTA 02004-1-2 Handover Documentation (https://industrialdigitaltwin.org/wp-content/uploads/2023/03/IDTA-02004-1-2_Submodel_Handover-Documentation.pdf)
- [IDTA 02011-1-0] IDTA 02011-1-0 Hierarchical Structures enabling Bills of Material (https://industrialdigitaltwin.org/wp-content/uploads/2023/04/IDTA-02011-1-0_Submodel_HierarchicalStructuresEnablingBoM.pdf)
- [IDTA 2023-01-24] IDTA 2023-01-24 Product Carbon Footprint (https://github.com/admin-shell-io/submodel-templates/blob/main/published/Carbon%20Footprint/0/9/IDTA%202023-0-9%20_Submodel_CarbonFootprint.pdf)
- [IDTA 02008-1-1] IDTA 02008-1-1 Time Series Data (https://industrialdigitaltwin.org/en/wp-content/uploads/sites/2/2023/03/IDTA-02008-1-1_Submodel_TimeSeriesData.pdf)
- [IDTA 02002-1-0] IDTA 02002-1-0 Submodel for Contact Information (https://industrialdigitaltwin.org/wp-content/uploads/2022/10/IDTA-02002-1-0_Submodel_ContactInformation.pdf)
- [IDTA 02007-1-0] IDTA 02007-1-0 Nameplate for Software in Manufacturing (https://industrialdigitaltwin.org/wp-content/uploads/2023/08/IDTA-02007-1-0_Submodel_Software-Nameplate.pdf)
- [IDS] International Data Spaces, IDS reference architecture model version 3.0, <https://internationaldataspaces.org/22m-3-0/>
- [IIRA] The Industrial Internet Reference Architecture Technical Report, <https://www.iiconsortium.org/pdf/IIRA-v1.9.pdf>
- [KGR20] H. Koziol, S. Grüner, J. Rückert, A Comparison of MQTT Brokers for Distributed IoT Edge Computing, ECSA, 2020
- [LNI40] LNI 4.0 Testbed Edge Configuration – Usage View, https://www.plattform-i40.de/PI40/Redaktion/EN/Downloads/Publikation/LNI4.0-Testbed-Edge-Configuration_UsageViewEN.pdf
- [MBB+18] E. Maleki, F. Belkadi, N. Boli, J. van der B. Zwaag, K. Alexopoulos, S. Koukas, M. Marin-Perianu, A. Bernard, D. Mourtzis, Ontology-Based Framework Enabling Smart Product-Service Systems: Application of Sensing Systems for Machine Health Monitoring, IEEE Internet of Things Journal, 5 (6), pp. 4496-4505, 2018
- [NE25] C. Nikolajew, H. Eichelberger, Industry 4.0 Connectors - A Performance Experiment with Modbus/TCP, Softwaretechnik-Trends Band 45, Heft 1, 2025
- [UML] OMG, Unified Modeling Language, Version 2.5.1, <https://www.omg.org/spec/UML/About-UML/>
- [Pid21] D. Pidun, Geräteverwaltung von IoT-Geräten für die IIP-Ecosphere Plattform, BSc-Abschlussarbeit, Universität Hildesheim, 2021

- [RAMI] Reference Architecture Model Industrie 4.0, <https://www.plattform-i40.de/PI40/Redaktion/EN/Downloads/Publikation/rami40-an-introduction.html>
- [Sal25] Ahmed Gamal Attia Mansour Salem, Design, Implementation and Evaluation of an Angular-based Configurator for oktoflow connectors, MSc thesis, University of Hildesheim, 2025
- [SEA+20] C. Sauer, H. Eichelberger, A. Ahmadian, A. Dewes, J. Jürjens, Aktuelle Industrie 4.0 Plattformen – Eine Übersicht, IIP-Ecosphere Whitepaper IIP-2020/001, 2020, DOI: 10.5281/zenodo.4485756, 2020
- [SEK21] K. Schmid, S. El-Sharkawy, C. Kröher, Improving Software Engineering Research through Experimentation Workbenches. arXiv e-prints, arXiv-2110, 2021
- [SE15] K. Schmid, H. Eichelberger, EASy-Producer: From Product Lines to Variability-rich Software Ecosystems, SPLC' 15, 2015
- [Sch23] L. Schulz, Container-Virtualisierung mit LXC in der IIP-Ecosphere-Plattform, BSc Arbeit, University of Hildesheim, 2023
- [Sta20] M. Staciwa, Experimentelles Container-Deployment auf Industrie 4.0 Geräte, Projektarbeit, University of Hildesheim, 2020
- [Sta22] M. Staciwa, Modell-basierte Erstellung containervirtualisierter Industrie 4.0 Anwendungen am Beispiel der IIP-Ecosphere-Plattform, BSc thesis University of Hildesheim, 2022
- [SSE21] H. Stichweh, C. Sauer, H. Eichelberger, IIP-Ecosphere Platform Requirements (Usage View), Version 1.0, Januar 2021, IIP-2021/001, DOI: 10.5281/zenodo.4485801, 2021
- [ZVEI-N] ZVEI, Specification Submodel Templates of the Asset Administration Shell – ZVEI Digital Nameplate for industrial equipment (Version 1.0), https://www.plattform-i40.de/IP/Redaktion/DE/Downloads/Publikation/Submodel_Templates-Asset_Administration_Shell-digital_nameplate.html